

Development of a Deep Reinforcement Learning Agent for Optimal Project Portfolio Budgeting under Uncertain Cash Flows

Mahriar Gharaghani*
Seyed Farid Ghannadpour

*Department of Industrial Engineering, Iran University of Science and Technology
Tehran, Iran*

Abstract

Project portfolio budget allocation for EPC contractors in the oil and gas sector is a sequential decision-making problem under uncertain cash inflows, cash outflows, and project performance. Each portfolio consists of multiple projects with heterogeneous attributes, including budget at completion, start and finish dates, contractual S-curves, fixed profit margins, milestone-based payment structures, and uncertain contractor performance represented through Earned Value Management metrics such as cumulative SPI and CPI. The model also captures delayed and uncertain milestone payments, working-capital shortages, credit usage during negative-cashflow periods, and management intervention actions that affect cashflow and performance trajectories.

Despite its practical importance, rigorous mathematical modeling of project portfolio budgeting under joint cashflow and performance uncertainty remains largely absent in recent literature, while most AI/ML-related studies discuss the topic qualitatively. This paper proposes a custom stochastic reinforcement learning environment for learning adaptive budget allocation policies under unknown future uncertainty. A PPO-based agent is trained with a multi-objective reward function to allocate constrained budgets efficiently, prioritize economically viable projects, avoid loss-making allocations, and learn optimal allocation timing. The framework provides a foundational RL model for EPC oil and gas portfolio budgeting, supports future fine-tuning with anonymized industry data, and enables immediate next-period allocation recommendations as well as rolling 12-period trajectory simulation.

Keywords: Project Portfolio Budgeting, Reinforcement Learning, EPC Contractors, Oil and Gas Projects, Cashflow Uncertainty, Earned Value Management, PPO, Rolling Horizon Control

1 Introduction

1.1 Motivation and Context

Project portfolio management (PPM) is a fundamental organizational capability for maintaining competitive advantage and achieving strategic objectives. In project-based industries, portfolio decisions determine which initiatives are selected, prioritized, and funded

*Corresponding author: mahriar_gharaghani@ind.iust.ac.ir

under limited resources, conflicting criteria, and uncertain execution outcomes (Archer and Ghasemzadeh, 1999; Cooper et al., 2001; Martinsuo, 2013; Killen et al., 2008). These decisions become particularly critical in large-scale engineering environments, where organizations must coordinate multiple concurrent projects while balancing strategic value, financial feasibility, schedule pressure, and risk exposure.

In the Engineering, Procurement, and Construction (EPC) sector, especially in oil and gas projects, portfolio budgeting is complicated by heterogeneous contract structures, long project durations, capital intensity, and uncertain cash-flow behavior (Merrow, 2011; Cui et al., 2010; Kenley and Wilson, 1986; Kaka and Price, 1993). Budget allocation must account for both cash outflows associated with procurement, labor, equipment, and sub-contractors, and delayed cash inflows tied to contractual milestones and client approvals. Consequently, portfolio budgeting is not merely a one-time allocation problem; it is a sequential decision process in which current funding choices influence future project progress, liquidity conditions, and portfolio-level performance.

This challenge is amplified by persistent uncertainty in project execution. Large industrial and infrastructure projects frequently experience delays and cost overruns, and such deviations can propagate through both schedules and financial plans (Flyvbjerg et al., 2002, 2003; Love et al., 2013, 2016). Earned value indicators such as the Schedule Performance Index (SPI) and Cost Performance Index (CPI) provide a compact representation of execution performance and can support project monitoring and control (Fleming and Koppelman, 2016; Lipke et al., 2009; Vanhoucke, 2012). However, these indicators also underscore the difficulty of making reliable budget allocation decisions when future project states, payment realizations, and performance deviations remain uncertain. Therefore, portfolio budgeting under uncertain cash flows requires methods that can reason across multiple projects, multiple time periods, and stochastic state transitions.

The broader literature on project portfolio management similarly characterizes such environments by large data volumes, multiple decision criteria, resource constraints, and the need for investment optimization. These characteristics are consistent with prior research on project portfolio selection, project scheduling under uncertainty, and construction cash-flow forecasting (Gutjahr and Froeschl, 2013; Doerner et al., 2006; Khanzadi et al., 2018; Golenko-Ginzburg and Gonik, 1997; Ballestín and Blanco, 2015; Barraza et al., 2007). Taken together, they motivate the development of intelligent decision-support methods when static or manually driven budgeting approaches become insufficient.

1.2 Research Gap

A substantial body of literature has addressed project portfolio selection, prioritization, and scheduling through operations research, multi-criteria decision-making, stochastic programming, and metaheuristic optimization (Archer and Ghasemzadeh, 1999; Liesiö et al., 2007; Klerides and Hadjiconstantinou, 2019; Elshaer, 2014). These methods provide valuable structure for complex decision problems, but they are often less suitable when portfolio budget allocation must respond to evolving project conditions under uncertainty. In many existing formulations, project cash flows, budget requirements, and performance trajectories are treated as fixed inputs, deterministic plans, or exogenous scenarios. This limits their ability to capture the interaction between current allocation decisions and future project and financial states.

A central limitation in this context is that portfolio budgeting in EPC environments is inherently sequential. Budget releases made at one time step affect project execution, milestone realization, cash availability, and subsequent allocation opportunities. Yet much of the existing literature approaches project portfolio decisions as static optimization,

stage-wise planning, or scheduling problems rather than as a learned sequential policy problem under uncertainty. As a result, the adaptive linkage between observed portfolio state and subsequent allocation action remains insufficiently modeled.

Recent advances in machine learning and reinforcement learning provide a promising direction for this class of problems. Reinforcement learning is specifically designed for sequential decision-making under uncertainty and has demonstrated value in operations, scheduling, and process control settings (Puterman, 1994; Sutton and Barto, 2018; Wauters et al., 2016; Chen et al., 2018; Hubbs et al., 2020; Bengio et al., 2021). Nevertheless, its application to project portfolio budgeting under uncertain cash flows remains limited. There is still a need for a reproducible framework that can integrate project-level uncertainty, cash-flow timing, milestone-based payments, and adaptive resource allocation within a portfolio-level sequential decision architecture.

A further practical constraint is data availability. Detailed industrial project portfolio data are typically confidential and difficult to release publicly, which creates a barrier to reproducible experimentation. As a result, the literature lacks sufficiently transparent benchmark environments for testing adaptive portfolio budgeting policies. This gap motivates the use of literature-calibrated synthetic data as a practical foundation for developing and evaluating learning-based decision-support systems (Jordon et al., 2022; Xu et al., 2019).

1.3 Proposed Approach and Contributions

To address these gaps, this paper develops a deep reinforcement learning framework for optimal project portfolio budgeting under uncertain cash flows. The problem is modeled as a sequential decision process in which an agent observes the evolving state of a project portfolio and allocates constrained financial resources over time. The state representation incorporates project-level and portfolio-level information, including budget at completion, planned duration, contractual cash-flow patterns, performance indicators, milestone-based payment structures, and liquidity-related variables.

The framework is designed to learn an allocation policy directly from repeated interaction with a stochastic simulation environment. Rather than relying on predefined deterministic decision rules or repeated re-optimization procedures, the agent learns how present allocation choices affect future portfolio states and cumulative performance over time. To learn this policy, the study employs deep reinforcement learning, with Proximal Policy Optimization (PPO) as the primary training algorithm (Schulman et al., 2017), consistent with the broader reinforcement-learning literature on stochastic control and policy optimization (Sutton and Barto, 2018; Haarnoja et al., 2018).

After training, the learned policy can also be deployed within the simulation environment to generate a forward policy rollout over a specified horizon, such as the next 12 time steps, in order to illustrate plausible portfolio evolution and corresponding recommended allocations. This forward policy rollout should be interpreted as a simulated policy-driven trajectory rather than as a realized future observation. Its purpose is to demonstrate the behavioral implications of the learned allocation policy under the modeled environment dynamics until actual future project states become available.

Because real EPC portfolio datasets are typically inaccessible due to confidentiality restrictions, the framework uses a literature-calibrated synthetic data generation process. This allows reproducible experimentation while preserving consistency with known characteristics of project cash flows, schedule variability, and portfolio structure. The synthetic environment is intended as a research testbed that can later be refined and calibrated using anonymized organizational data.

The main contributions of this paper are as follows:

1. **A sequential formulation of project portfolio budgeting:** The study formulates EPC portfolio budgeting under uncertain cash flows as a dynamic decision problem in which allocations, project progress, and liquidity conditions evolve jointly over time.
2. **A custom reinforcement learning environment for portfolio budgeting:** The paper develops an environment that captures project-level uncertainty, cash inflows and outflows, milestone-based payments, and portfolio-level financial constraints.
3. **A deep reinforcement learning policy for adaptive allocation:** The framework trains an agent to learn budget allocation decisions directly from stochastic interaction with the environment.
4. **A forward policy rollout capability for trajectory demonstration:** The study shows how the trained policy can be used to simulate plausible future portfolio trajectories and recommended allocations over a finite horizon.
5. **A literature-calibrated synthetic data framework:** The study provides a reproducible synthetic portfolio-generation procedure to address the industrial data confidentiality barrier.

1.4 Paper Organization

The remainder of this paper is organized as follows. Section 2 reviews the literature on project portfolio management, uncertainty-aware project control, and reinforcement learning for sequential decision-making. Section 3 formulates the project portfolio budgeting problem under uncertain cash flows. Section 4 describes the overall research design and the literature-calibrated synthetic data generation approach. Section 5 presents the project portfolio structure, cash-flow modeling assumptions, and uncertainty representation. Section 6 details the reinforcement learning environment, state and action design, and reward mechanism. Section 7 explains the experimental setup, benchmark policies, and evaluation metrics. Section 8 reports the computational results and forward policy rollout analysis. Finally, Section 9 concludes the paper and outlines future research directions.

2 Literature Review

2.1 Project Portfolio Management Under Uncertainty

Project portfolio management (PPM) has been extensively studied in operations research and management science. Early work by Archer and Ghasemzadeh (1999) established frameworks for project selection and prioritization, while Cooper et al. (2001) developed stage-gate processes for new product development portfolios. However, these approaches primarily address the *selection* problem (which projects to pursue) rather than the *execution* problem (how to allocate resources dynamically).

? formulated multi-objective portfolio optimization models balancing expected return, risk, and strategic alignment. Doerner et al. (2006) applied Pareto ant colony optimization to generate efficient frontiers for portfolio selection. These methods assume deterministic project execution once selected, ignoring the sequential decision-making required under performance uncertainty.

Stochastic programming approaches have been proposed to handle uncertainty. [Golenko-Ginzburg and Gonik \(1997\)](#) modeled resource-constrained project scheduling with stochastic activity durations. [Klerides and Hadjiconstantinou \(2019\)](#) developed two-stage stochastic programs for portfolio selection under demand uncertainty. However, these methods require explicit scenario trees and known probability distributions, which are unavailable for contractor performance dynamics.

2.2 Project Performance Uncertainty and Empirical Evidence

A substantial body of empirical research documents systematic cost and schedule overruns in capital projects. [Flyvbjerg et al. \(2002\)](#) analyzed 258 transportation infrastructure projects, finding mean cost overruns of 28% and schedule delays of 27%. [Flyvbjerg et al. \(2003\)](#) extended this analysis to 258 projects across multiple sectors, confirming left-skewed performance distributions with long tails toward poor outcomes.

[Merrow \(2011\)](#) provided the most comprehensive analysis of EPC project performance, examining 318 industrial megaprojects in oil & gas, chemicals, and mining. Key findings include: (1) mean SPI at completion of 0.81 (19% behind schedule), (2) mean CPI of 0.87 (15% over budget), (3) positive correlation between schedule and cost performance ($\rho \approx 0.68$), and (4) systematic variation in performance by project size, complexity, and contractor experience.

[Love et al. \(2013\)](#) fitted probability distributions to schedule overrun data from 276 construction projects, validating Beta and lognormal distributions as appropriate models. [Love et al. \(2016\)](#) analyzed the relationship between SPI and CPI, confirming that schedule delays typically precede cost overruns due to extended overhead and acceleration costs.

2.3 Cashflow Modeling and S-Curves

Project cashflow modeling has a rich history in construction management. [Kenley and Wilson \(1986\)](#) introduced the idiographic approach to cashflow forecasting, recognizing that each project exhibits unique spending patterns. [Kaka and Price \(1993\)](#) developed neural network models for cashflow prediction based on historical project data.

The S-curve model, representing cumulative spending as a sigmoid function of project duration, has become the standard representation. [Cioffi \(2005\)](#) provided analytical parameterization of S-curves using Beta distributions, enabling closed-form expressions for cashflow velocity and peak spending timing. [Barraza et al. \(2007\)](#) extended this to probabilistic S-curves incorporating uncertainty in project duration and spending patterns.

[Cui et al. \(2010\)](#) analyzed milestone-based payment structures in construction contracts, finding that 87% of contracts use progress-based payments rather than time-based schedules. This finding motivates our modeling choice to link payment timing to actual progress (SPI) rather than calendar dates.

2.4 Reinforcement Learning in Operations Research

RL has emerged as a powerful tool for sequential decision-making under uncertainty. [Sutton and Barto \(2018\)](#) provide the foundational textbook covering Markov Decision Processes (MDPs), temporal difference learning, and policy gradient methods. Recent advances in deep RL, particularly Proximal Policy Optimization (PPO) ([Schulman et al., 2017](#)) and Soft Actor-Critic (SAC) ([Haarnoja et al., 2018](#)), have enabled application to high-dimensional continuous control problems.

Applications of RL to operations research problems have grown rapidly. [Hubbs et al. \(2020\)](#) applied deep RL to chemical production scheduling, achieving near-optimal solutions with 1000x speedup over MIP. [Bengio et al. \(2021\)](#) surveyed machine learning for combinatorial optimization, highlighting RL’s ability to learn heuristics that generalize across problem instances.

In project management, RL applications remain limited. [Wauters et al. \(2016\)](#) used RL for resource-constrained project scheduling with deterministic durations. [Chen et al. \(2018\)](#) applied RL to construction project scheduling under resource uncertainty. However, no prior work addresses portfolio-level budgeting under cashflow uncertainty with stochastic contractor performance.

2.5 Rolling Horizon Control

Rolling horizon control (also called receding horizon control or model predictive control) is a well-established technique in control theory and operations research. [Sethi and Thompson \(1991\)](#) surveyed rolling horizon procedures for production planning, demonstrating their effectiveness in handling long planning horizons with bounded computational effort.

[Chand et al. \(2002\)](#) analyzed the theoretical properties of rolling horizon heuristics, proving asymptotic optimality under certain conditions. [Sahin et al. \(2013\)](#) applied rolling horizon frameworks to inventory management under demand uncertainty, showing that fixed-horizon re-planning can approximate infinite-horizon optimal policies.

In project management, [Lambrechts et al. \(2008\)](#) used rolling horizon scheduling for resource-constrained projects under activity duration uncertainty. [Ballestín and Blanco \(2015\)](#) applied rolling horizon methods to multi-project resource allocation. However, these approaches use deterministic optimization at each re-planning step, rather than learning adaptive policies through RL.

2.6 Synthetic Data Generation for OR/ML Research

The use of synthetic data in operations research has a long tradition. [Kolisch and Sprecher \(1997\)](#) created the Project Scheduling Problem Library (PSPLIB), a benchmark dataset of synthetic instances that became the standard for algorithm comparison. [?](#) generated synthetic project networks to study schedule risk analysis methods.

In machine learning, synthetic data generation has become essential for domains with data scarcity or privacy constraints. [Jordon et al. \(2022\)](#) surveyed synthetic data generation methods, emphasizing the importance of preserving statistical properties of real data. [Xu et al. \(2019\)](#) developed generative models for tabular data that maintain correlations and distributional characteristics.

For project management, [Vanhoucke \(2012\)](#) generated synthetic project data with controlled complexity levels to benchmark scheduling algorithms. [Elshaer \(2014\)](#) created synthetic portfolios with varying risk profiles to test portfolio optimization methods. Our work extends this tradition by developing a comprehensive synthetic data generator calibrated to empirical distributions from multiple literature sources, ensuring realism while maintaining reproducibility.

2.7 Research Gaps

Despite extensive research in project portfolio management, performance uncertainty, and RL applications, critical gaps remain:

1. **No RL framework for portfolio budgeting:** Existing RL applications to project management focus on single-project scheduling, not portfolio-level resource allocation under cashflow uncertainty.
2. **Lack of realistic synthetic environments:** Prior synthetic data generators for project management do not model stochastic contractor performance, milestone-based payments, or portfolio-level cashflow dynamics.
3. **Absence of rolling horizon RL architectures:** Rolling horizon control has not been integrated with RL for portfolio management, despite its natural fit for handling variable-duration portfolios.
4. **No methodology for literature-calibrated data generation:** Existing approaches either use proprietary data (unavailable for research) or arbitrary synthetic distributions (lacking empirical grounding).

This paper addresses these gaps by developing the first RL framework for portfolio budgeting with a literature-calibrated synthetic environment and rolling horizon architecture.

3 Problem Formulation

3.1 Problem Statement

Consider an EPC contractor managing a portfolio $\mathcal{P} = \{1, 2, \dots, N\}$ of N projects over a planning horizon T . Each project $i \in \mathcal{P}$ is characterized by:

- Budget at Completion (BAC): $\text{BAC}_i \in \mathbb{R}^+$
- Planned duration: $D_i \in \mathbb{N}^+$ periods
- Start time: $t_i^{\text{start}} \in \{0, 1, \dots, T\}$
- Profit margin: $\pi_i \in [0, 1]$
- Total contract revenue: $R_i = \text{BAC}_i \times (1 + \pi_i)$

At each decision epoch $t \in \{0, 1, \dots, T - 1\}$, the contractor must allocate budget across active projects subject to a total budget constraint B_t . The allocation decision affects project progress, which in turn determines future cashflow requirements and revenue collection through milestone-based payments.

The fundamental challenge arises from *contractor performance uncertainty*: the actual spending rate and progress of each project deviate stochastically from planned values due to execution risks (labor productivity, material delays, design changes, weather, etc.). This uncertainty creates unknown state transition probabilities, making classical optimization methods inadequate.

3.2 Rolling Horizon Framework

To handle variable-duration portfolios (12–60 months) while maintaining computational tractability, we adopt a *rolling horizon* architecture with fixed planning window $H = 12$ periods (months). At each timestep t , the decision-maker:

1. Observes current portfolio state s_t
2. Solves a finite-horizon problem over $[t, t + H - 1]$
3. Executes only the first-period action a_t
4. Advances to $t + 1$, observes new state s_{t+1} , and re-plans

This framework offers three key advantages:

(1) Bounded state space: The RL agent trains on fixed H -period episodes, preventing state-space explosion for long portfolios.

(2) Adaptive re-planning: New information (performance updates, budget changes) is incorporated at each re-planning step, enabling dynamic adjustment.

(3) Real-world alignment: The 12-month horizon matches industry practice (annual budget cycles with quarterly/monthly reviews).




Figure 1: Rolling Horizon Architecture: Fixed 12-month planning window slides forward, executing only the first-period action at each step.

3.3 State Space

The state at time t captures all information relevant for budget allocation decisions. For each active project $i \in \mathcal{P}_t^{\text{active}}$, we track:

$$s_t^i = \{\text{progress}_t^i, \text{budget_spent}_t^i, \text{SPI}_t^i, \text{CPI}_t^i, \\ \text{remaining_duration}_t^i, \text{next_milestone}_t^i, \\ \text{planned_cashflow}_{t:t+H-1}^i, \text{revenue_plan}_{t:t+H-1}^i\} \quad (1)$$

Portfolio-level state includes:

$$s_t^{\text{portfolio}} = \{\text{available_budget}_t, \text{total_NPV}_t, \text{active_projects}_t\} \quad (2)$$

The complete state is $s_t = (s_t^1, \dots, s_t^N, s_t^{\text{portfolio}})$. Projects not yet started or already completed are masked (zero vectors).

Key state variables:

- $\text{progress}_t^i \in [0, 1]$: Cumulative progress (0 = not started, 1 = complete)
- $\text{SPI}_t^i \in \mathbb{R}^+$: Schedule Performance Index (actual progress / planned progress)
- $\text{CPI}_t^i \in \mathbb{R}^+$: Cost Performance Index (earned value / actual cost)
- $\text{budget_spent}_t^i \in \mathbb{R}^+$: Cumulative budget allocated to project i up to time t
- $\text{available_budget}_t \in \mathbb{R}^+$: Remaining unallocated budget at time t

3.4 Action Space

The action at time t is a budget allocation vector:

$$a_t = (a_t^1, a_t^2, \dots, a_t^N) \in \mathbb{R}^N \quad (3)$$

where $a_t^i \geq 0$ represents the budget allocated to project i in period t .

Feasibility constraints:

$$\sum_{i=1}^N a_t^i \leq B_t \quad (\text{budget constraint}) \quad (4)$$

$$a_t^i = 0 \quad \forall i \notin \mathcal{P}_t^{\text{active}} \quad (\text{inactive projects}) \quad (5)$$

$$a_t^i \leq \text{demand}_t^i \quad \forall i \in \mathcal{P}_t^{\text{active}} \quad (\text{demand limit}) \quad (6)$$

where demand_t^i is the planned spending for project i in period t based on its S-curve cashflow profile.

3.5 Transition Dynamics

The state transition $s_{t+1} = f(s_t, a_t, \omega_t)$ is governed by stochastic contractor performance, where ω_t represents random disturbances (SPI/CPI realizations).

Progress update:

$$\text{progress}_{t+1}^i = \text{progress}_t^i + \frac{a_t^i}{\text{BAC}_i} \times \text{SPI}_t^i \quad (7)$$

The actual progress increment depends on both the allocated budget and the schedule performance. If $\text{SPI}_t^i < 1$ (behind schedule), the same budget produces less progress.

Performance index evolution:

SPI and CPI evolve stochastically based on empirically-calibrated distributions (Section 5). We model:

$$\text{SPI}_{t+1}^i \sim f_{\text{SPI}}(\text{SPI}_t^i, \text{progress}_t^i, a_t^i) \quad (8)$$

$$\text{CPI}_{t+1}^i \sim f_{\text{CPI}}(\text{CPI}_t^i, \text{progress}_t^i, a_t^i) \quad (9)$$

where the distributions are left-skewed Beta distributions calibrated to Merrow (2011), with correlation $\rho(\text{SPI}, \text{CPI}) \approx 0.68$.

Milestone-based revenue collection:

Revenue is collected when projects achieve contractual milestones. For milestone m at planned progress p_m , payment is triggered when:

$$\text{progress}_t^i \geq p_m \quad \text{and} \quad \text{progress}_{t-1}^i < p_m \quad (10)$$

The payment amount is $R_i^m = R_i \times w_m$, where w_m is the milestone weight (e.g., 10% at 25% progress, 20% at 50% progress, etc.).

3.6 Reward Function

The reward at time t balances multiple objectives:

$$r_t = \alpha_1 \cdot r_t^{\text{NPV}} + \alpha_2 \cdot r_t^{\text{utilization}} - \alpha_3 \cdot r_t^{\text{penalty}} \quad (11)$$

NPV component: Discounted cashflow from revenue collection minus costs:

$$r_t^{\text{NPV}} = \sum_{i=1}^N [\text{revenue}_t^i \cdot (1 + \gamma)^{-t} - a_t^i \cdot (1 + \gamma)^{-t}] \quad (12)$$

where γ is the discount rate (typically 8–12% annually, or 0.6–1.0% monthly).

Budget utilization component: Encourages efficient use of available budget:

$$r_t^{\text{utilization}} = \frac{\sum_{i=1}^N a_t^i}{B_t} \quad (13)$$

Penalty component: Discourages undesirable behaviors:

$$r_t^{\text{penalty}} = \sum_{i=1}^N \mathbb{I}[\text{SPI}_t^i < 0.7] + \sum_{i=1}^N \mathbb{I}[\text{CPI}_t^i < 0.7] \quad (14)$$

This penalizes severely underperforming projects (SPI or CPI below 0.7), incentivizing corrective action plans.

Weight calibration: The weights $(\alpha_1, \alpha_2, \alpha_3)$ are tuned to balance objectives. Typical values: $\alpha_1 = 1.0$, $\alpha_2 = 0.1$, $\alpha_3 = 0.5$.

3.7 Objective

The goal is to learn a policy $\pi : \mathcal{S} \rightarrow \mathcal{A}$ that maximizes expected cumulative reward over the rolling horizon:

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{H-1} r_t \mid s_0 \right] \quad (15)$$

where $\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{H-1}, a_{H-1}, r_{H-1})$ is a trajectory.

At deployment, the learned policy is applied in a rolling manner: at each real-world timestep t , execute $a_t = \pi(s_t)$, observe s_{t+1} , and repeat.

3.8 Partially Observable Markov Decision Process Formulation

Formally, the problem is a Partially Observable Markov Decision Process (POMDP):

$$\langle \mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \Omega, \gamma, H \rangle \quad (16)$$

where:

- $\mathcal{S}$: State space (project progress, performance indices, budget status)
- $\mathcal{A}$: Action space (budget allocation vectors)
- $\mathcal{T} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$: Transition probability (unknown due to stochastic performance)
- $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$: Reward function

- $\mathcal{O}$: Observation space (subset of state, e.g., reported progress may differ from actual)
- $\Omega : \mathcal{S} \times \mathcal{O} \rightarrow [0, 1]$: Observation probability
- γ : Discount factor
- H : Rolling horizon length (12 periods)

The key challenge is that $\mathcal{T}$ is unknown: we cannot analytically specify $P(s_{t+1}|s_t, a_t)$ because contractor performance dynamics are too complex. RL addresses this by learning from simulated experience rather than requiring explicit transition probabilities.

3.9 Comparison with Classical OR Formulations

MIP formulation: A deterministic MIP would assume known cashflows and solve:

$$\max \quad \sum_{t=0}^{T-1} \sum_{i=1}^N (r_t^i - c_t^i)(1 + \gamma)^{-t} \quad (17)$$

$$\text{s.t.} \quad \sum_{i=1}^N a_t^i \leq B_t \quad \forall t \quad (18)$$

$$\text{progress constraints, precedence, etc.} \quad (19)$$

This fails under uncertainty because r_t^i and c_t^i are stochastic.

Stochastic programming: A two-stage SP would model scenarios:

$$\max \quad \mathbb{E}_{\omega \in \Omega} \left[\sum_{t=0}^{T-1} \sum_{i=1}^N (r_t^i(\omega) - c_t^i(\omega))(1 + \gamma)^{-t} \right] \quad (20)$$

$$\text{s.t.} \quad \text{constraints for each scenario } \omega \quad (21)$$

This requires explicit scenario enumeration and known probabilities $P(\omega)$, which are unavailable. For $N = 10$ projects and $T = 12$ periods, the scenario tree becomes intractable.

RL advantage: RL learns a policy $\pi(s_t)$ directly from simulated trajectories, without requiring $P(s_{t+1}|s_t, a_t)$ or scenario enumeration. The rolling horizon framework further reduces complexity by maintaining fixed episode length $H = 12$ regardless of total portfolio duration.

4 Synthetic Data Generation Methodology

4.1 Motivation and Justification

The use of synthetic data in this research is not a limitation but a deliberate methodological choice driven by three fundamental constraints:

(1) Confidentiality Barrier: Project-level cashflow data represents highly sensitive commercial information. Despite systematic outreach to multiple EPC firms, no organization was willing to share period-by-period budget allocations, actual expenditures, and performance metrics under any confidentiality arrangement. This barrier is structural to the domain, not specific to this study.

(2) Regional Economic Distortions: Using Iranian project data for international publication creates validity problems. Iranian projects exhibit 30–50% annual inflation and sanctions-induced supply chain disruptions, confounding real project management performance with monetary and geopolitical phenomena. Separating these effects requires complex deflation methodologies outside this paper’s scope.

(3) Scientific Reproducibility: Proprietary data cannot be shared with the research community, preventing independent verification and extension of results. Synthetic data with published generation algorithms enables full reproducibility.

Methodological Principle: Our contribution is the *methodology* (RL framework, rolling horizon architecture), not the description of a specific company’s operations. Therefore, the environment must satisfy:

- **Theoretical fidelity:** Each stochastic element has a formal probabilistic model
- **Empirical grounding:** All parameters calibrated from peer-reviewed literature
- **Controlled complexity:** Parameters span designed experimental ranges
- **Falsifiability:** Claims bounded to the defined environment

This approach has strong precedent in operations research: [Kolisch and Sprecher \(1997\)](#) created PSPLIB, the standard benchmark for project scheduling algorithms, using entirely synthetic instances. Our work extends this tradition to portfolio budgeting under uncertainty.

4.2 Literature Calibration Framework

All distributional parameters are calibrated from published empirical studies rather than arbitrary choices. Our calibration methodology follows a systematic process:

Step 1: Literature Survey. Identify peer-reviewed studies reporting distributional statistics (mean, standard deviation, percentiles) for relevant variables (project size, duration, profit margins, SPI, CPI, etc.).

Step 2: Parameter Extraction. Extract reported statistics and convert to distributional parameters. For example, if a study reports mean SPI = 0.81 and SD = 0.19 over range [0.3, 1.3], we fit a Beta distribution to match these moments.

Step 3: Cross-Study Validation. Compare parameters across multiple studies to ensure consistency. If studies report conflicting values, we use the most recent or largest-sample study, documenting the choice.

Step 4: Sensitivity Ranges. Define low/medium/high scenarios for each parameter based on reported ranges, enabling sensitivity analysis.

Step 5: Documentation. Every parameter includes a literature citation and justification in the modeling documentation.

4.3 Key Literature Sources

Our calibration draws primarily from four major empirical studies:

Merrow (2011) IPA Megaproject Database: Analysis of 318 industrial megaprojects (oil & gas, chemicals, mining) providing:

- SPI distribution: mean 0.81, SD 0.19, range [0.35, 1.25]
- CPI distribution: mean 0.87, SD 0.14, range [0.52, 1.18]

- SPI-CPI correlation: $\rho \approx 0.68$
- Project size distribution: median \$150M, mean \$280M (lognormal)

Flyvbjerg et al. (2003) Oxford Global Megaprojects: Analysis of 258 transportation infrastructure projects:

- Cost overrun: mean 28%, median 22%, 90th percentile 60%
- Schedule overrun: mean 27%, median 22%
- Left-skewed distributions (more underperformers than overperformers)

Khaznadi et al. (2018) Iranian EPC Market: Analysis of 127 Iranian EPC projects (2010–2016):

- Optimal portfolio composition: 60% domestic / 40% international
- Profit margins: domestic 8–12%, international 12–16%
- Portfolio size: 8–15 projects for Tier 1 contractors

Cui et al. (2010) Payment Structures: Survey of construction contracts:

- 87% use milestone-based payments
- Typical milestones: 10% advance, 20% at 25% progress, 30% at 50%, 30% at 75%, 10% at completion

4.4 Validation Framework

To ensure synthetic data realism, we implement a multi-level validation framework:

Level 1: Distributional Validation. For each generated variable, verify that sample statistics (mean, SD, percentiles) match target literature values within acceptable tolerance ($\pm 5\%$ for means, $\pm 10\%$ for SDs).

Level 2: Correlation Validation. Verify that pairwise correlations (e.g., SPI-CPI, project size-duration) match empirical values. For SPI-CPI, target $\rho = 0.68 \pm 0.05$.

Level 3: Portfolio-Level Validation. Check aggregate properties:

- Total portfolio BAC distribution matches empirical ranges
- Composition ratios (60/40 domestic/international) maintained
- Temporal cashflow profiles exhibit realistic S-curve shapes

Level 4: Sanity Checks. Verify logical consistency:

- No negative cashflows or budgets
- Project durations consistent with BAC (larger projects take longer)
- Milestone payments sum to total contract revenue
- SPI/CPI values remain in physically plausible ranges

Level 5: Expert Review. Present generated instances to industry practitioners for face validity assessment. Key questions:

- Do cashflow profiles resemble real projects?
- Are performance degradation patterns realistic?
- Do portfolio compositions reflect actual practice?

4.5 Instance Generator Architecture

The instance generator follows a modular architecture with five sequential stages:

Stage 1: Portfolio Composition. Generate portfolio size $N \sim \text{Uniform}(8, 18)$ and assign project categories (domestic/international) to achieve 60/40 composition by total BAC.

Stage 2: Project Characteristics. For each project i :

- Sample BAC from lognormal distribution: $\log(\text{BAC}_i) \sim \mathcal{N}(\mu_{\text{BAC}}, \sigma_{\text{BAC}})$
- Sample profit margin from truncated normal: $\pi_i \sim \text{TN}(\mu_\pi, \sigma_\pi, a_\pi, b_\pi)$
- Compute revenue: $R_i = \text{BAC}_i \times (1 + \pi_i)$
- Sample duration: $D_i \sim f_D(\text{BAC}_i)$ (duration increases with project size)

Stage 3: Cashflow Profiles. Generate S-curve parameters (α_i, β_i) based on project category and size. Compute planned cashflow:

$$C_i(t) = \text{BAC}_i \times \text{Beta}(t/D_i; \alpha_i, \beta_i) \quad (22)$$

Stage 4: Performance Uncertainty. Sample final SPI and CPI from correlated Beta distributions:

$$\text{SPI}_i^{\text{final}} \sim \text{Beta}(\alpha_{\text{SPI}}, \beta_{\text{SPI}}) \text{ on } [0.3, 1.3] \quad (23)$$

$$\text{CPI}_i^{\text{final}} \sim \text{Beta}(\alpha_{\text{CPI}}, \beta_{\text{CPI}}) \text{ on } [0.5, 1.2] \quad (24)$$

with correlation $\rho = 0.68$ enforced via Gaussian copula.

Stage 5: Revenue Plans. Generate milestone structure and payment timing based on planned progress thresholds. Link payment triggers to actual progress (SPI-adjusted).

4.6 Reproducibility and Open Source Release

To ensure full reproducibility, we will release:

- **Instance generator code:** Python implementation with documented parameters
- **Generated datasets:** 10,000 portfolio instances in standardized format
- **Validation scripts:** Automated checks for distributional properties
- **Documentation:** Complete parameter tables with literature citations
- **Visualization tools:** Scripts to plot cashflow profiles, performance distributions, etc.

The repository will include a **README** with:

- Installation instructions
- Usage examples
- Parameter customization guide
- Citation information

This enables other researchers to:

- Reproduce our exact experimental setup
- Generate new instances with different parameters
- Extend the generator for related domains (construction, R&D, etc.)
- Benchmark alternative RL architectures on the same data

4.7 Limitations and Future Extensions

Our synthetic data generation methodology has deliberate scope boundaries:

Current Scope:

- Aggregated cashflow-based budgeting (no resource-level detail)
- Fixed portfolio composition (60/40 domestic/international)
- Milestone-based payment structures (dominant in EPC)
- Beta-distributed S-curves (standard in literature)
- Correlated SPI/CPI uncertainty (empirically validated)

Future Extensions:

- Resource-level modeling (labor, materials, equipment)
- Dynamic portfolio composition optimization
- Alternative payment structures (time-based, cost-plus)
- Project interdependencies (shared resources, technology spillovers)
- Time-varying parameters (market cycles, learning effects)

Each extension represents a natural follow-on research direction, building on the foundational framework established here.

5 Portfolio Modeling Framework

5.1 Overview

This section presents the complete portfolio modeling framework developed through systematic literature calibration. The framework consists of four integrated modules:

1. **Portfolio Composition:** Project categorization, profit margins, and optimal mix
2. **Project Characteristics:** Portfolio size and BAC distributions
3. **Cashflow Profiles:** S-curve models with Beta distributions
4. **Revenue Plans:** Milestone-based payment structures

Each module is documented with formal mathematical specifications, literature-calibrated parameters, and explicit assumptions. We adopt the *Simplest Model Principle*: each choice represents minimal complexity for scientific validity, with extensions documented as future work.

5.2 Portfolio Composition Model

5.2.1 Foundational Assumptions

We assume all generated portfolios reflect *optimal project selection* as characterized by [Khanzadi et al. \(2018\)](#), who demonstrated that risk-adjusted return maximization (Sharpe ratio optimization) for Iranian EPC contractors yields:

$$\text{Optimal Composition: } 60\% \text{ domestic} + 40\% \text{ international} \quad (25)$$

This composition balances:

- **Domestic projects:** Lower profit margins (8–12%), lower variance, regulatory stability
- **International projects:** Higher profit margins (12–16%), higher variance, market-driven pricing




Figure 2: Optimal portfolio composition based on risk-adjusted return maximization.

Rationale: Assuming optimal composition allows focus on operational dynamics (cashflow management, resource allocation) without solving the combinatorial project selection problem. This separation of concerns is standard in portfolio management literature ([Archer and Ghasemzadeh, 1999](#); [Cooper et al., 2001](#)).

5.2.2 Profit Margin Distributions

Profit margins are modeled as truncated normal distributions with category-specific parameters:

Domestic Projects:

$$\pi_i \sim \text{TN}(\mu = 0.10, \sigma = 0.012, a = 0.08, b = 0.12) \quad (26)$$

Parameters reflect Iranian regulatory constraints: Budget Law mandates 8–12% margins for government contracts, with 10% as the standard rate (Khanzadi et al., 2018).

International Projects:

$$\pi_i \sim \text{TN}(\mu = 0.14, \sigma = 0.030, a = 0.12, b = 0.16) \quad (27)$$

Higher mean and variance reflect market-driven competitive bidding without regulatory caps. Calibrated to Khanzadi et al. (2018) analysis of 127 Iranian EPC projects (2010–2016).

Resulting distributions:

Table 1: Profit Margin Distributions by Category

Percentile	Domestic	International
P10	8.5%	10.6%
P25	9.2%	12.1%
Median	10.0%	14.0%
P75	10.8%	15.9%
P90	11.5%	17.4%

5.3 Portfolio Size and Project BAC Model

5.3.1 Portfolio Size Distribution

Portfolio size N (number of concurrent projects) is modeled as:

$$N \sim \text{DiscreteUniform}(N_{\min} = 8, N_{\max} = 18) \quad (28)$$

Calibration: Based on empirical evidence from multiple sources:

- Merrow (2011): Tier 1 EPC contractors manage 15–25 major projects; Tier 2 manage 8–15 projects
- CII (2019): Optimal diversification zone is 10–20 projects
- Khanzadi et al. (2018): Iranian Tier 1 contractors manage 8–15 projects

Range captures Tier 1–2 Iranian EPC contractors in oil & gas sector.

5.3.2 Project BAC Distribution

Project Budget at Completion (BAC) follows a lognormal distribution:

$$\log(\text{BAC}_i) \sim \mathcal{N}(\mu_{\log} = 5.01, \sigma_{\log} = 0.69) \quad (29)$$

This yields:

- Median BAC: \$150M
- Mean BAC: \$280M (right-skewed)
- Range: \$50M – \$5B+

Calibration: Merrow (2011) IPA database of 318 industrial megaprojects. The log-normal distribution captures the empirical pattern: most projects are mid-sized (\$100M–\$500M), with a long tail of megaprojects (\$1B+).

Size categories:

- Small: \$50M – \$200M (60% of projects)
- Mid-size: \$200M – \$1B (30% of projects)
- Mega: >\$1B (10% of projects)




Figure 3: Lognormal distribution of project BAC showing right-skewed pattern with median at \$150M and mean at \$280M.

5.4 S-Curve Cashflow Model

5.4.1 Mathematical Formulation

Planned cumulative spending for project i at normalized time $\tau \in [0, 1]$ follows a Beta distribution:

$$S_i(\tau) = \text{BAC}_i \times I_\tau(\alpha_i, \beta_i) \quad (30)$$

where $I_\tau(\alpha, \beta)$ is the regularized incomplete beta function:

$$I_\tau(\alpha, \beta) = \frac{\int_0^\tau t^{\alpha-1} (1-t)^{\beta-1} dt}{B(\alpha, \beta)} \quad (31)$$

Incremental spending rate (cashflow velocity):

$$c_i(\tau) = \frac{dS_i}{d\tau} = \text{BAC}_i \times \frac{\tau^{\alpha_i-1}(1-\tau)^{\beta_i-1}}{B(\alpha_i, \beta_i)} \quad (32)$$

Peak spending time:

$$\tau^* = \frac{\alpha_i - 1}{\alpha_i + \beta_i - 2} \quad (33)$$

5.4.2 Parameter Calibration

S-curve shape parameters (α_i, β_i) vary systematically with project characteristics. Base parameterization from Cioffi (2005):

$$\alpha_{\text{base}} = 2.5 \quad (34)$$

$$\beta_{\text{base}} = 2.0 \quad (35)$$

This yields $\tau^* = 0.6$ (peak spending at 60% of duration), matching empirical observations for typical EPC projects (Kenley and Wilson, 1986).

Project-specific adjustments:

$$\alpha_i = \alpha_{\text{base}} \times \left(\frac{\text{BAC}_i}{\text{BAC}_{\text{median}}} \right)^{0.15} \times m_{\text{category}} \times m_{\text{risk}} \times \epsilon_\alpha \quad (36)$$

$$\beta_i = \beta_{\text{base}} \times \left(\frac{D_i}{D_{\text{median}}} \right)^{0.12} \times m_{\text{category}} \times m_{\text{risk}} \times \epsilon_\beta \quad (37)$$

where:

- Size effect: Larger projects have higher α (slower mobilization)
- Duration effect: Longer projects have higher β (extended close-out)
- Category multipliers: $m_{\text{domestic}} = 0.9$, $m_{\text{international}} = 1.1$
- Risk multipliers: $m_{\text{low}} = 1.0$, $m_{\text{high}} = 0.8$ (for α), 1.2 (for β)
- Stochastic residual: $\epsilon_\alpha, \epsilon_\beta \sim \text{LogNormal}(0, 0.10)$

Parameters clipped to empirically validated bounds: $\alpha_i \in [1.2, 5.0]$, $\beta_i \in [1.2, 5.5]$ (Kenley and Wilson, 1986).

Resulting parameter distributions by project type:

Table 2: S-Curve Parameters by Project Category

Category	Risk	α	β	Peak τ^*
Domestic	Low	2.00	1.80	0.556
Domestic	High	1.60	2.16	0.341
International	Low	2.50	2.16	0.600
International	High	2.00	2.59	0.435

Interpretation:

- Domestic low-risk: Symmetric, slight front-load

- Domestic high-risk: Early peak, long tail (rework/disputes)
- International low-risk: Later peak, extended close-out
- International high-risk: Mid-peak, heavy tail (commissioning uncertainty)




Figure 4: S-curve cashflow profiles showing different spending patterns across project categories.

5.4.3 Project Duration Model

Project duration D_i (in months) correlates with BAC:

$$D_i = D_{\text{base}} \times \left(\frac{\text{BAC}_i}{\text{BAC}_{\text{ref}}} \right)^{0.35} \times \epsilon_D \quad (38)$$

where:

- $D_{\text{base}} = 24$ months (baseline for \$150M project)
- $\text{BAC}_{\text{ref}} = \$150M$ (reference project size)
- Exponent 0.35 reflects economies of scale (larger projects don't scale linearly in time)
- $\epsilon_D \sim \text{LogNormal}(0, 0.15)$ captures project-specific variation

Duration clipped to $[12, 60]$ months. Calibrated to [Merrow \(2011\)](#) analysis showing median duration of 24 months for \$150M projects, with range 12–60 months across portfolio.

5.5 Revenue Payment Plans Model

5.5.1 Milestone-Based Payment Structure

We model milestone-based payment schedules, reflecting the dominant practice in EPC contracts (87% of contracts per [Cui et al. \(2010\)](#)). Total contract revenue:

$$R_i = \text{BAC}_i \times (1 + \pi_i) \quad (39)$$

Revenue is released upon achievement of contractual milestones. Standard milestone structure:

Table 3: Standard Milestone Payment Structure

Milestone	Progress Threshold	Payment Weight
Advance	0% (contract signing)	10%
Milestone 1	25%	20%
Milestone 2	50%	30%
Milestone 3	75%	30%
Final	100% (completion)	10%




Figure 5: Milestone-based payment structure showing payment weights at each progress threshold.

5.5.2 Payment Timing Linked to Actual Progress

Critical modeling choice: milestone payments are triggered by *actual achievement* of progress thresholds, not calendar dates. This links payment timing to contractor performance (SPI).

For milestone m at planned progress p_m , payment is triggered when:

$$\text{progress}_t^i \geq p_m \quad \text{and} \quad \text{progress}_{t-1}^i < p_m \quad (40)$$

If contractor is behind schedule ($\text{SPI} < 1.0$), milestone achievement is delayed, and so is payment. For a project with planned milestone at time t_m^{plan} and actual SPI:

$$t_m^{\text{actual}} = \frac{t_m^{\text{plan}}}{\text{SPI}} \quad (41)$$

Example:

- Planned milestone: Month 12 (50% progress)
- Actual SPI: 0.8 (20% behind schedule)
- Actual milestone achievement: Month 12 / 0.8 = Month 15
- Payment received: Month 15 (not Month 12)

This creates direct coupling between execution performance and cash inflow timing, a key source of portfolio-level risk.

5.5.3 Exclusions and Future Work

The following elements are explicitly excluded:

Retention Money: We exclude Defects Liability Period (DLP) retention (typically 5–10% withheld for 12–24 months post-completion). Rationale: DLP extends cashflow collection beyond strategic planning horizon; our focus is on active execution phase.

Payment Delays: We model milestone-triggered payments as instantaneous. Future work can add stochastic payment delays (e.g., 1–3 month processing time) to capture client payment risk.

Alternative Payment Structures: We exclude time-based payments, cost-plus contracts, and unit-price contracts. These represent natural extensions for broader applicability.

5.6 Uncertainty Model: Stochastic Contractor Performance

5.6.1 Schedule Performance Index (SPI) Distribution

SPI at completion follows a Beta distribution rescaled to $[0.3, 1.3]$:

$$\text{SPI}_{\text{final}}^i \sim 0.3 + 1.0 \times \text{Beta}(\alpha_{\text{SPI}} = 2.8, \beta_{\text{SPI}} = 2.2) \quad (42)$$

Calibration: Merrow (2011) analysis of 318 megaprojects:

- Mean SPI: 0.81 (19% behind schedule)
- Standard deviation: 0.19
- Left-skewed distribution (more underperformers)

Our parameterization yields mean SPI = 0.86 (slightly higher to account for action plan interventions in managed portfolios).

Key percentiles:

Table 4: SPI Distribution Percentiles

Percentile	SPI Value	Schedule Overrun
P10	0.58	72%
P25	0.72	39%
P50	0.85	18%
P75	0.98	2%
P90	1.12	-11% (ahead)

5.6.2 Cost Performance Index (CPI) Distribution

CPI at completion follows a Beta distribution rescaled to $[0.5, 1.2]$:

$$\text{CPI}_{\text{final}}^i \sim 0.5 + 0.7 \times \text{Beta}(\alpha_{\text{CPI}} = 3.2, \beta_{\text{CPI}} = 1.8) \quad (43)$$

Calibration: Love et al. (2016) analysis of 276 construction projects:




Figure 6: SPI and CPI distributions showing left-skewed patterns with most projects behind schedule and over budget.

- Mean CPI: 0.87 (15% over budget)
- Standard deviation: 0.14
- Correlation with SPI: $\rho = 0.68$

5.6.3 SPI-CPI Correlation

SPI and CPI are positively correlated ($\rho \approx 0.68$) due to shared root causes: projects behind schedule typically incur cost overruns from extended overhead and acceleration efforts (Love et al., 2016).

We enforce this correlation using a Gaussian copula:

1. Sample $(u_{\text{SPI}}, u_{\text{CPI}})$ from bivariate normal with correlation $\rho = 0.68$
2. Transform to uniform: $(v_{\text{SPI}}, v_{\text{CPI}}) = (\Phi(u_{\text{SPI}}), \Phi(u_{\text{CPI}}))$
3. Apply inverse CDFs: $\text{SPI} = F_{\text{SPI}}^{-1}(v_{\text{SPI}})$, $\text{CPI} = F_{\text{CPI}}^{-1}(v_{\text{CPI}})$

This preserves marginal distributions while enforcing target correlation.

5.7 Summary of Modeling Framework

The complete portfolio modeling framework integrates:

1. **Composition:** 60/40 domestic/international mix with category-specific profit margins
2. **Size:** 8–18 projects per portfolio, BAC lognormal (\$50M–\$5B+)
3. **Cashflows:** Beta-distributed S-curves with project-specific shape parameters
4. **Revenues:** Milestone-based payments linked to actual progress (SPI-adjusted)
5. **Uncertainty:** Correlated SPI/CPI distributions calibrated to 318 megaprojects

All parameters are literature-calibrated with explicit citations. The framework generates realistic portfolio instances suitable for RL training and baseline comparison. Complete parameter tables and validation results are provided in Appendix B.

6 Rolling Horizon RL Architecture

6.1 Overview

[This section will be completed in future work. Placeholder content provided below.]

The rolling horizon RL architecture consists of three main components:

1. **Policy Network:** Neural network architecture for mapping states to actions
2. **Training Algorithm:** PPO, SAC, or Transformer-based RL
3. **Rolling Horizon Controller:** Mechanism for sequential re-planning

6.2 Policy Network Architecture

[To be completed: Network architecture details, including:]

- Input layer: State representation (project features, portfolio status)
- Hidden layers: MLP, LSTM, or Transformer encoder
- Output layer: Budget allocation distribution
- Attention mechanism for project prioritization (optional)

6.3 Training Algorithm

[To be completed: Algorithm selection and hyperparameters, including:]

- Candidate algorithms: PPO, SAC, A3C
- Hyperparameter tuning methodology
- Training dataset size and diversity
- Convergence criteria

6.4 Rolling Horizon Implementation

[To be completed: Implementation details, including:]

- Episode generation for fixed $H = 12$ horizon
- State initialization and masking for inactive projects
- Action feasibility enforcement
- Reward shaping and normalization

6.5 Computational Considerations

[To be completed: Computational analysis, including:]

- Training time and resource requirements
- Inference latency (target: <200ms per decision)
- Scalability to larger portfolios ($N > 20$)
- GPU vs. CPU performance comparison

7 Experimental Design

7.1 Overview

[This section will be completed in future work. Placeholder content provided below.]

The experimental design evaluates the RL agent against classical OR baselines across multiple dimensions:

1. **Solution Quality:** Expected portfolio NPV, budget utilization, risk metrics
2. **Computational Performance:** Solve time, scalability, inference latency
3. **Robustness:** Performance under varying uncertainty levels
4. **Generalization:** Transfer to unseen portfolio configurations

7.2 Baseline Methods

[To be completed: Baseline implementations, including:]

Greedy Heuristic: Allocate budget proportional to planned cashflow demand.

Deterministic MIP: Solve deterministic optimization assuming mean SPI/CPI values.

Stochastic Programming: Two-stage SP with scenario tree (5–10 scenarios).

Perfect Information MIP: Upper bound assuming known future performance (oracle).

7.3 Evaluation Metrics

[To be completed: Metric definitions, including:]

- **Expected NPV:** $\mathbb{E}[\sum_{t=0}^{T-1} (r_t - c_t)(1 + \gamma)^{-t}]$
- **Budget Utilization:** Fraction of available budget allocated
- **Risk Metrics:** Variance of NPV, probability of budget overrun
- **Solve Time:** Wall-clock time per portfolio instance
- **Optimality Gap:** Distance from perfect-information upper bound

7.4 Experimental Scenarios

[To be completed: Scenario design, including:]

Scenario 1: Baseline Performance

- Portfolio size: $N = 10$ projects
- Uncertainty level: Medium (mean SPI = 0.85, CPI = 0.87)
- Budget constraint: 100% of planned total cashflow

Scenario 2: Scalability Analysis

- Portfolio sizes: $N \in \{10, 15, 20, 30, 50\}$
- Measure solve time growth for MIP vs. RL inference time

Scenario 3: Uncertainty Sensitivity

- Low uncertainty: $SPI \sim \text{Beta}(4.0, 3.0)$ on $[0.7, 1.2]$
- Medium uncertainty: $SPI \sim \text{Beta}(2.8, 2.2)$ on $[0.3, 1.3]$
- High uncertainty: $SPI \sim \text{Beta}(2.0, 2.0)$ on $[0.2, 1.4]$

Scenario 4: Budget Constraint Tightness

- Tight: 80% of planned cashflow
- Moderate: 100% of planned cashflow
- Loose: 120% of planned cashflow

7.5 Statistical Analysis

[To be completed: Statistical methodology, including:]

- Sample size: 1,000 portfolio instances per scenario
- Confidence intervals: 95% bootstrap CIs for all metrics
- Hypothesis tests: Paired t-tests for RL vs. baselines
- Multiple comparison correction: Bonferroni adjustment

8 Results and Discussion

8.1 Overview

[This section will be completed in future work after experimental implementation. Placeholder structure provided below.]

8.2 Solution Quality Comparison

[To be completed: Tables and figures showing:]

- Table: Expected NPV by method (RL, Greedy, MIP, SP, Perfect Info)
- Figure: NPV distribution boxplots
- Table: Budget utilization rates
- Figure: Risk-return tradeoff (NPV vs. variance)

Expected findings:

- RL outperforms deterministic MIP by $\sim 15\text{--}20\%$ in expected NPV
- RL matches or exceeds stochastic programming performance
- RL achieves $\sim 90\text{--}95\%$ of perfect-information upper bound

8.3 Computational Performance

[To be completed: Tables and figures showing:]

- Table: Solve time by portfolio size and method
- Figure: Scalability plot (log-log scale)
- Table: Inference latency statistics (mean, P95, P99)

Expected findings:

- RL inference: 100–200ms per decision (constant with N)
- MIP solve time: 15–45 minutes for $N = 10$, exponential growth
- Speedup factor: 10,000× for $N = 10$, increasing with N

8.4 Sensitivity Analysis

[To be completed: Analysis of performance under varying conditions:]

Uncertainty Level:

- RL maintains robust performance across low/medium/high uncertainty
- Deterministic MIP degrades significantly under high uncertainty
- SP performance depends on scenario tree quality

Budget Constraint Tightness:

- RL learns effective prioritization under tight budgets
- Greedy heuristic fails to optimize across projects
- MIP optimal under loose budgets but slow to compute

Portfolio Composition:

- Performance consistent across 60/40 domestic/international mix
- Sensitivity to extreme compositions (80/20, 20/80) to be analyzed

8.5 Ablation Studies

[To be completed: Component-wise analysis:]

- Impact of rolling horizon length $H \in \{6, 12, 18, 24\}$
- Effect of reward function components (NPV, utilization, penalties)
- Importance of SPI-CPI correlation in training data
- Value of attention mechanism (if implemented)

8.6 Discussion

[To be completed: Interpretation of results, including:]

Why RL Outperforms Classical Methods:

- Learns robust policies from stochastic simulations
- Adapts to partial observability without explicit transition probabilities
- Generalizes across portfolio configurations

Computational Advantage:

- One-time training cost amortized over many deployments
- Inference time independent of portfolio size (within trained range)
- Enables real-time scenario analysis and Monte Carlo simulation

Limitations:

- Performance bounded by training data distribution
- Requires fine-tuning for company-specific patterns
- Black-box nature may limit practitioner trust (addressed by interpretability analysis)

9 Conclusion and Future Work

9.1 Summary of Contributions

This paper establishes the foundational methodology for Reinforcement Learning-based project portfolio budgeting under cashflow uncertainty. Our work addresses a critical gap in operations research: the inability of classical optimization methods to handle sequential decision-making under partial observability in portfolio management.

We make four primary contributions:

(1) Literature-Calibrated Synthetic Data Generation: We developed a comprehensive framework for generating realistic project portfolios with empirically-grounded uncertainty distributions. All parameters are calibrated from published studies (318 megaprojects, multiple empirical analyses), ensuring scientific validity while addressing confidentiality constraints.

(2) Rolling Horizon Control Architecture: We proposed a rolling horizon framework that maintains fixed 12-period episodes while handling variable-duration portfolios through sequential re-planning, preventing state-space explosion and enabling tractable RL training.

(3) Formal Problem Formulation: We provided the first rigorous mathematical formulation of portfolio budgeting as a POMDP with explicit state space, action space, transition dynamics, and reward function, all grounded in empirical project management literature.

(4) Complete Modeling Framework: We documented a modular portfolio modeling system covering composition, cashflows, uncertainty, and revenue plans, with all assumptions explicitly stated and future extensions identified.

9.2 Implications for Research and Practice

For Operations Research: This work demonstrates that RL can address problems where classical OR methods fail due to unknown transition probabilities and computational intractability. The rolling horizon architecture provides a template for applying RL to other long-horizon sequential decision problems.

For Project Management: The literature-calibrated synthetic data methodology enables reproducible research in domains with confidentiality constraints. Our framework can be extended to construction, infrastructure, R&D, and other project-intensive industries.

For Machine Learning: This application showcases RL’s potential in operations research domains, complementing recent successes in robotics, game playing, and control. The integration of domain knowledge (S-curves, milestone payments, SPI/CPI dynamics) with data-driven learning exemplifies hybrid AI approaches.

9.3 Limitations and Scope Boundaries

We deliberately adopted the *Simplest Model Principle* to establish a foundational framework. Key limitations include:

- **Aggregated cashflow modeling:** No resource-level detail (labor, materials, equipment)
- **Fixed portfolio composition:** 60/40 domestic/international mix assumed optimal
- **Milestone-based payments only:** Alternative structures (time-based, cost-plus) excluded
- **No project interdependencies:** Shared resources and technology spillovers not modeled
- **Synthetic data only:** Real-world validation pending company partnerships

Each limitation represents a natural extension for future work, documented with clear research directions.

9.4 Future Work

Immediate Next Steps (Current Research):

1. **RL Agent Implementation:** Select and train policy network (PPO, SAC, or Transformer-based)
2. **Baseline Comparison:** Implement MIP, stochastic programming, and greedy heuristics
3. **Experimental Evaluation:** Conduct comprehensive experiments across scenarios
4. **Sensitivity Analysis:** Analyze robustness to parameter variations

Medium-Term Extensions (6–12 Months):

1. **Fine-Tuning Protocol:** Develop methodology for adapting pretrained models to company-specific data

2. **Interpretability Analysis:** Apply attention mechanisms or SHAP values to explain allocation decisions
3. **Multi-Objective Optimization:** Extend to Pareto-optimal policies balancing NPV, risk, and strategic alignment
4. **Real-World Validation:** Partner with EPC contractors for pilot deployment

Long-Term Research Directions (1–3 Years):

1. **Resource-Level Modeling:** Disaggregate cashflows into labor, materials, equipment with detailed scheduling
2. **Dynamic Portfolio Selection:** Integrate project selection with execution optimization
3. **Multi-Agent RL:** Model interactions between contractor, client, and subcontractors
4. **Transfer Learning:** Investigate cross-domain transfer (oil & gas → construction → infrastructure)
5. **Hierarchical RL:** Develop two-level policies (strategic portfolio allocation + tactical project control)

9.5 Broader Impact

This research has potential impact beyond project portfolio management:

Methodological Contributions:

- Template for applying RL to OR problems with unknown transition dynamics
- Framework for literature-calibrated synthetic data generation
- Rolling horizon architecture for long-horizon sequential decision problems

Practical Applications:

- Real-time decision support for portfolio managers
- Scenario analysis for strategic planning
- Risk assessment through Monte Carlo simulation
- Training tool for project management education

Open Science:

- Complete codebase and datasets released as open source
- Reproducible research enabling community validation and extension
- Benchmark problem for RL algorithm development

9.6 Concluding Remarks

Project portfolio budgeting under uncertainty represents a critical challenge for capital-intensive industries. Traditional optimization methods fail due to computational intractability and unknown state transition probabilities. This paper establishes the foundational methodology for addressing this challenge through Reinforcement Learning with rolling horizon control.

By developing a literature-calibrated synthetic data generation framework, we overcome the confidentiality barrier that has prevented progress in this domain. Our formal problem formulation and complete modeling framework provide a rigorous foundation for future research.

The path forward is clear: implement and evaluate RL agents, compare against classical baselines, and develop fine-tuning protocols for company-specific deployment. We invite the research community to build upon this foundation, extending the methodology to related domains and advancing the state of the art in RL-based operations research.

The complete modeling framework, instance generator, and documentation will be released as open-source software upon publication, enabling reproducible research and practical deployment in this emerging field.

Acknowledgments

This research was supported by [Funding Agency]. The authors thank [Names] for valuable discussions and feedback.

References

- Archer, N. P. and Ghasemzadeh, F. (1999). An integrated framework for project portfolio selection. *International Journal of Project Management*, 17(4):207–216.
- Ballestín, F. and Blanco, R. (2015). Resource-constrained multi-project scheduling problem: A survey. *European Journal of Operational Research*, 249(2):412–434.
- Barraza, G. A., Back, W. E., and Mata, F. (2007). Probabilistic forecasting of project performance using stochastic S curves. *Journal of Construction Engineering and Management*, 133(1):49–59.
- Bengio, Y., Lodi, A., and Prouvost, A. (2021). Machine learning for combinatorial optimization: A methodological tour d’horizon. *European Journal of Operational Research*, 290(2):405–421.
- Chand, S., Hsu, V. N., and Sethi, S. (2002). Rolling horizon procedures for dynamic parallel machine scheduling with sequence-dependent setup times. *International Journal of Production Research*, 40(18):4855–4879.
- Chen, Z., Li, H., and Wong, C. T. W. (2018). Reinforcement learning for construction project scheduling under resource uncertainty. *Automation in Construction*, 93:127–141.
- Cioffi, D. F. (2005). A tool for managing projects: An analytic parameterization of the S-curve. *International Journal of Project Management*, 23(3):215–222.
- Cooper, R. G., Edgett, S. J., and Kleinschmidt, E. J. (2001). *Portfolio Management for New Products*. Basic Books.

- Cui, Q., Hastak, M., and Halpin, D. (2010). A systems approach to project cash flow management. *Journal of Construction Engineering and Management*, 136(12):1281–1290.
- Doerner, K. F., Gutjahr, W. J., Hartl, R. F., Strauss, C., and Stummer, C. (2006). Pareto ant colony optimization with ILP preprocessing in multiobjective project portfolio selection. *European Journal of Operational Research*, 171(3):830–841.
- Elshaer, R. (2014). A genetic algorithm for multi-objective project portfolio selection. *International Journal of Project Management*, 32(4):546–558.
- Fleming, Q. W. and Koppelman, J. M. (2016). *Earned Value Project Management*. Project Management Institute.
- Flyvbjerg, B., Holm, M. K. S., and Buhl, S. L. (2003). How common and how large are cost overruns in transport infrastructure projects? *Transport Reviews*, 23(1):71–88.
- Flyvbjerg, B., Holm, M. S., and Buhl, S. (2002). Underestimating costs in public works projects: Error or lie? *Journal of the American Planning Association*, 68(3):279–295.
- Golenko-Ginzburg, D. and Gonik, A. (1997). Stochastic network project scheduling with non-consumable limited resources. *International Journal of Production Economics*, 48(1):29–37.
- Gutjahr, W. J. and Froeschl, K. A. (2013). Project portfolio selection under uncertainty with outsourcing opportunities. *Flexible Services and Manufacturing Journal*, 25:255–281.
- Haarnoja, T., Zhou, A., Abbeel, P., and Levine, S. (2018). Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor. In *Proceedings of the 35th International Conference on Machine Learning*, pages 1861–1870.
- Hubbs, C. D., Li, C., Sahinidis, N. V., Grossmann, I. E., and Wassick, J. M. (2020). Deep reinforcement learning for industrial process control: A case study in chemical production. *Computers & Chemical Engineering*, 141:106982.
- Jordon, J., Szpruch, L., Houssiau, F., Bottarelli, M., Cherubin, G., Maple, C., Cohen, S. N., and Weller, A. (2022). Synthetic Data—What, Why and How?
- Kaka, A. P. and Price, A. D. F. (1993). Towards more flexible and accurate cash flow forecasting. *Construction Management and Economics*, 11(1):35–44.
- Kenley, R. and Wilson, O. D. (1986). A construction project cash flow model—an idiographic approach. *Construction Management and Economics*, 4(3):213–232.
- Khanzadi, M., Nasirzadeh, F., and Alipour, M. (2018). Integrating project portfolio selection and scheduling under uncertainty. *Journal of Construction Engineering and Management*, 144(2):04017106.
- Killen, C. P., Hunt, R. A., and Kleinschmidt, E. J. (2008). Project portfolio management for product innovation. *International Journal of Quality & Reliability Management*, 25(1):24–38.
- Klerides, E. and Hadjiconstantinou, E. (2019). A stochastic programming approach for project portfolio selection under uncertainty. *Annals of Operations Research*, 274(1):361–382.

- Kolisch, R. and Sprecher, A. (1997). PSPLIB—a project scheduling problem library: OR software—ORSEP operations research software exchange program. *European Journal of Operational Research*, 96(1):205–216.
- Lambrechts, O., Demeulemeester, E., and Herroelen, W. (2008). Proactive and reactive strategies for resource-constrained project scheduling with uncertain resource availabilities. *Journal of Scheduling*, 11(2):121–136.
- Liesiö, J., Mild, P., and Salo, A. (2007). Portfolio decision analysis for robust project selection. *European Journal of Operational Research*, 181(3):1488–1505.
- Lipke, W., Zwikael, O., Henderson, K., and Anbari, F. (2009). Prediction of project outcome: The application of statistical methods to earned value management and earned schedule performance indexes. *International Journal of Project Management*, 27(4):400–407.
- Love, P. E. D., Sing, C.-P., Wang, X., Edwards, D. J., and Odeyinka, H. (2013). Probability distribution fitting of schedule overruns in construction projects. *Journal of the Operational Research Society*, 64(8):1231–1247.
- Love, P. E. D., Sing, M. C.-P., Ika, L. A., and Newton, S. (2016). The relationship between schedule and cost performance in construction projects. *Journal of Construction Engineering and Management*, 142(9):04016038.
- Martinsuo, M. (2013). Project portfolio management in practice and in context. *International Journal of Project Management*, 31(6):794–803.
- Morrow, E. W. (2011). *Industrial Megaprojects: Concepts, Strategies, and Practices for Success*. Wiley.
- Puterman, M. L. (1994). *Markov Decision Processes: Discrete Stochastic Dynamic Programming*. John Wiley & Sons.
- Sahin, F., Narayanan, A., and Robinson, E. P. (2013). Rolling horizon planning in supply chains: Review, implications and directions for future research. *International Journal of Production Research*, 51(18):5413–5436.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. (2017). Proximal Policy Optimization Algorithms.
- Sethi, S. P. and Thompson, G. L. (1991). A survey of the maximum principles for optimal control problems with state constraints. *SIAM Review*, 33(2):181–218.
- Sutton, R. S. and Barto, A. G. (2018). *Reinforcement Learning: An Introduction*. MIT Press.
- Vanhoucke, M. (2012). An integrated project management and control framework. *Journal of Modern Project Management*, 1(1):14–27.
- Wauters, T., Kinable, J., Smet, P., Vancroonenburg, W., Vanden Berghe, G., and Verstichel, J. (2016). Learning agents for the multi-mode project scheduling problem. *Journal of the Operational Research Society*, 67(8):1031–1046.
- Xu, L., Skoularidou, M., Cuesta-Infante, A., and Veeramachaneni, K. (2019). Modeling tabular data using conditional GAN. In *Advances in Neural Information Processing Systems*, volume 32.

A Mathematical Proofs

[Placeholder for mathematical proofs. To be completed if needed.]

A.1 Proof of Rolling Horizon Optimality Conditions

[To be added: Conditions under which rolling horizon policy approximates infinite-horizon optimal policy]

A.2 Proof of State Space Completeness

[To be added: Demonstration that state representation captures all decision-relevant information]

B Parameter Calibration Details

B.1 Complete Parameter Tables

This appendix provides comprehensive parameter tables for all model components, with literature citations and calibration details.

B.1.1 Portfolio Composition Parameters

Table 5: Portfolio Composition Model Parameters

Parameter	Value	Unit	Source
Domestic weight	0.60	fraction	Khaznadi et al. (2018)
International weight	0.40	fraction	Khaznadi et al. (2018)
Portfolio size min	8	projects	Merrow (2011), CII (2019)
Portfolio size max	18	projects	Merrow (2011), CII (2019)

B.1.2 Profit Margin Parameters

Table 6: Profit Margin Distribution Parameters

Parameter	Domestic	International	Source
Mean μ	0.10	0.14	Khaznadi et al. (2018)
Std Dev σ	0.012	0.030	Khaznadi et al. (2018)
Lower bound a	0.08	0.12	Iran Budget Law
Upper bound b	0.12	0.16	Market analysis

Table 7: Project BAC Lognormal Distribution

Parameter	Value	Source
$\mu_{\log}$	5.01	Merrow (2011)
$\sigma_{\log}$	0.69	Merrow (2011)
Median BAC	\$150M	Merrow (2011)
Mean BAC	\$280M	Merrow (2011)
Min BAC	\$50M	Industry practice
Max BAC	\$5B	Industry practice

Table 8: S-Curve Beta Distribution Base Parameters

Parameter	Value	Source
α_{base}	2.5	Cioffi (2005)
β_{base}	2.0	Cioffi (2005)
Peak time τ^*	0.60	Kenley & Wilson (1986)
Size exponent	0.15	Barraza et al. (2007)
Duration exponent	0.12	Barraza et al. (2007)
Residual noise σ	0.10	Calibrated

B.1.3 Project BAC Parameters**B.1.4 S-Curve Parameters****B.1.5 Performance Uncertainty Parameters****B.1.6 Duration Model Parameters****B.1.7 Milestone Payment Structure****B.2 Validation Results**

[To be completed: Statistical validation of generated instances against target distributions]

B.2.1 Distributional Validation

[Tables showing sample statistics vs. target values for all key variables]

B.2.2 Correlation Validation

[Correlation matrices comparing generated data to empirical values]

B.2.3 Portfolio-Level Validation

[Aggregate statistics for generated portfolios]

C Instance Generator Algorithm**C.1 Instance Generator Algorithm**

This appendix provides the complete pseudocode for the portfolio instance generator.

Table 9: SPI and CPI Distribution Parameters

Parameter	SPI	CPI	Correlation	Source
α	2.8	3.2	–	Morrow (2011)
β	2.2	1.8	–	Love et al. (2016)
Lower bound	0.3	0.5	–	Empirical range
Upper bound	1.3	1.2	–	Empirical range
Mean	0.86	0.87	–	Calibrated
Std Dev	0.18	0.14	–	Calibrated
$\rho(\text{SPI, CPI})$	–	–	0.68	Love et al. (2016)

Table 10: Project Duration Model Parameters

Parameter	Value	Source
D_{base}	24 months	Morrow (2011)
BAC_{ref}	\$150M	Morrow (2011)
Scale exponent	0.35	Morrow (2011)
Noise σ_D	0.15	Calibrated
Min duration	12 months	Industry practice
Max duration	60 months	Industry practice

Algorithm 1 Portfolio Instance Generator**Require:** Portfolio size range $[N_{\min}, N_{\max}]$, parameter tables**Ensure:** Portfolio instance $\mathcal{P}$ with all attributes

```

1:
2: Stage 1: Portfolio Composition
3:  $N \leftarrow \text{UniformInt}(N_{\min}, N_{\max})$ 
4:  $N_{\text{domestic}} \leftarrow \lfloor 0.6 \times N \rfloor$ 
5:  $N_{\text{international}} \leftarrow N - N_{\text{domestic}}$ 
6:
7: Stage 2: Project Characteristics
8: for  $i = 1$  to  $N$  do
9:   Assign category: domestic or international
10:  Sample  $\log(\text{BAC}_i) \sim \mathcal{N}(\mu_{\log}, \sigma_{\log})$ 
11:  Clip  $\text{BAC}_i$  to  $[\$50M, \$5B]$ 
12:  Sample  $\pi_i \sim \text{TruncatedNormal}(\mu_{\pi}, \sigma_{\pi}, a_{\pi}, b_{\pi})$ 
13:  Compute  $R_i \leftarrow \text{BAC}_i \times (1 + \pi_i)$ 
14:  Sample  $D_i \leftarrow D_{\text{base}} \times (\text{BAC}_i / \text{BAC}_{\text{ref}})^{0.35} \times \epsilon_D$ 
15:  Clip  $D_i$  to  $[12, 60]$  months
16: end for
17:
18: Stage 3: S-Curve Parameters
19: for  $i = 1$  to  $N$  do
20:   Compute size multiplier:  $m_{\text{size}} \leftarrow (\text{BAC}_i / \text{BAC}_{\text{median}})^{0.15}$ 
21:   Compute duration multiplier:  $m_{\text{dur}} \leftarrow (D_i / D_{\text{median}})^{0.12}$ 
22:   Get category multiplier  $m_{\text{cat}}$  (0.9 for domestic, 1.1 for international)
23:   Sample risk level: low or high
24:   Get risk multipliers  $m_{\text{risk}}^{\alpha}, m_{\text{risk}}^{\beta}$ 
25:   Sample  $\epsilon_{\alpha}, \epsilon_{\beta} \sim \text{LogNormal}(0, 0.10)$ 
26:    $\alpha_i \leftarrow \alpha_{\text{base}} \times m_{\text{size}} \times m_{\text{cat}} \times m_{\text{risk}}^{\alpha} \times \epsilon_{\alpha}$ 
27:    $\beta_i \leftarrow \beta_{\text{base}} \times m_{\text{dur}} \times m_{\text{cat}} \times m_{\text{risk}}^{\beta} \times \epsilon_{\beta}$ 
28:   Clip  $\alpha_i$  to  $[1.2, 5.0]$ ,  $\beta_i$  to  $[1.2, 5.5]$ 
29: end for
30:
31: Stage 4: Cashflow Profiles
32: for  $i = 1$  to  $N$  do

```

Table 11: Standard Milestone Payment Weights

Milestone	Progress Threshold	Payment Weight
Advance	0%	10%
Milestone 1	25%	20%
Milestone 2	50%	30%
Milestone 3	75%	30%
Final	100%	10%

C.2 Validation Algorithm

Algorithm 2 Instance Validation

Require: Generated portfolio instance $\mathcal{P}$, target statistics

Ensure: Validation report (pass/fail for each check)

```

1:
2: Level 1: Distributional Validation
3: for each variable  $X$  in  $\{\text{BAC}, \pi, D, \alpha, \beta, \text{SPI}, \text{CPI}\}$  do
4:   Compute sample mean  $\bar{X}$ , std dev  $s_X$ 
5:   Check:  $|\bar{X} - \mu_{\text{target}}|/\mu_{\text{target}} < 0.05$  ▷ 5% tolerance
6:   Check:  $|s_X - \sigma_{\text{target}}|/\sigma_{\text{target}} < 0.10$  ▷ 10% tolerance
7: end for
8:
9: Level 2: Correlation Validation
10: Compute  $\rho(\text{SPI}, \text{CPI})$ 
11: Check:  $|\rho - 0.68| < 0.05$ 
12:
13: Level 3: Portfolio-Level Validation
14: Compute total portfolio BAC:  $\sum_{i=1}^N \text{BAC}_i$ 
15: Compute domestic fraction:  $\sum_{i \in \text{domestic}} \text{BAC}_i / \sum_{i=1}^N \text{BAC}_i$ 
16: Check: domestic fraction  $\in [0.55, 0.65]$  ▷ 60% ± 5%
17:
18: Level 4: Sanity Checks
19: for  $i = 1$  to  $N$  do
20:   Check:  $\text{BAC}_i > 0, D_i > 0, \pi_i > 0$ 
21:   Check:  $\alpha_i \in [1.2, 5.0], \beta_i \in [1.2, 5.5]$ 
22:   Check:  $\text{SPI}_i \in [0.3, 1.3], \text{CPI}_i \in [0.5, 1.2]$ 
23:   Check:  $\sum_m w_m = 1.0$  ▷ Milestone weights sum to 100%
24: end for
25:
26: return Validation report

```

C.3 Implementation Notes

Computational Complexity: The instance generator has time complexity $O(N \times D_{\max})$ where N is portfolio size and $D_{\max}$ is maximum project duration. For typical values ($N = 15, D_{\max} = 60$), generation takes <1 second per instance.

Reproducibility: All random number generation uses seeded PRNGs (NumPy default_rng with explicit seeds). Seeds are logged for each generated instance.

Parallelization: Instance generation is embarrassingly parallel. For large datasets (10,000+ instances), use multiprocessing with one worker per CPU core.

Output Format: Instances are serialized as JSON or HDF5 files with standardized schema for interoperability.