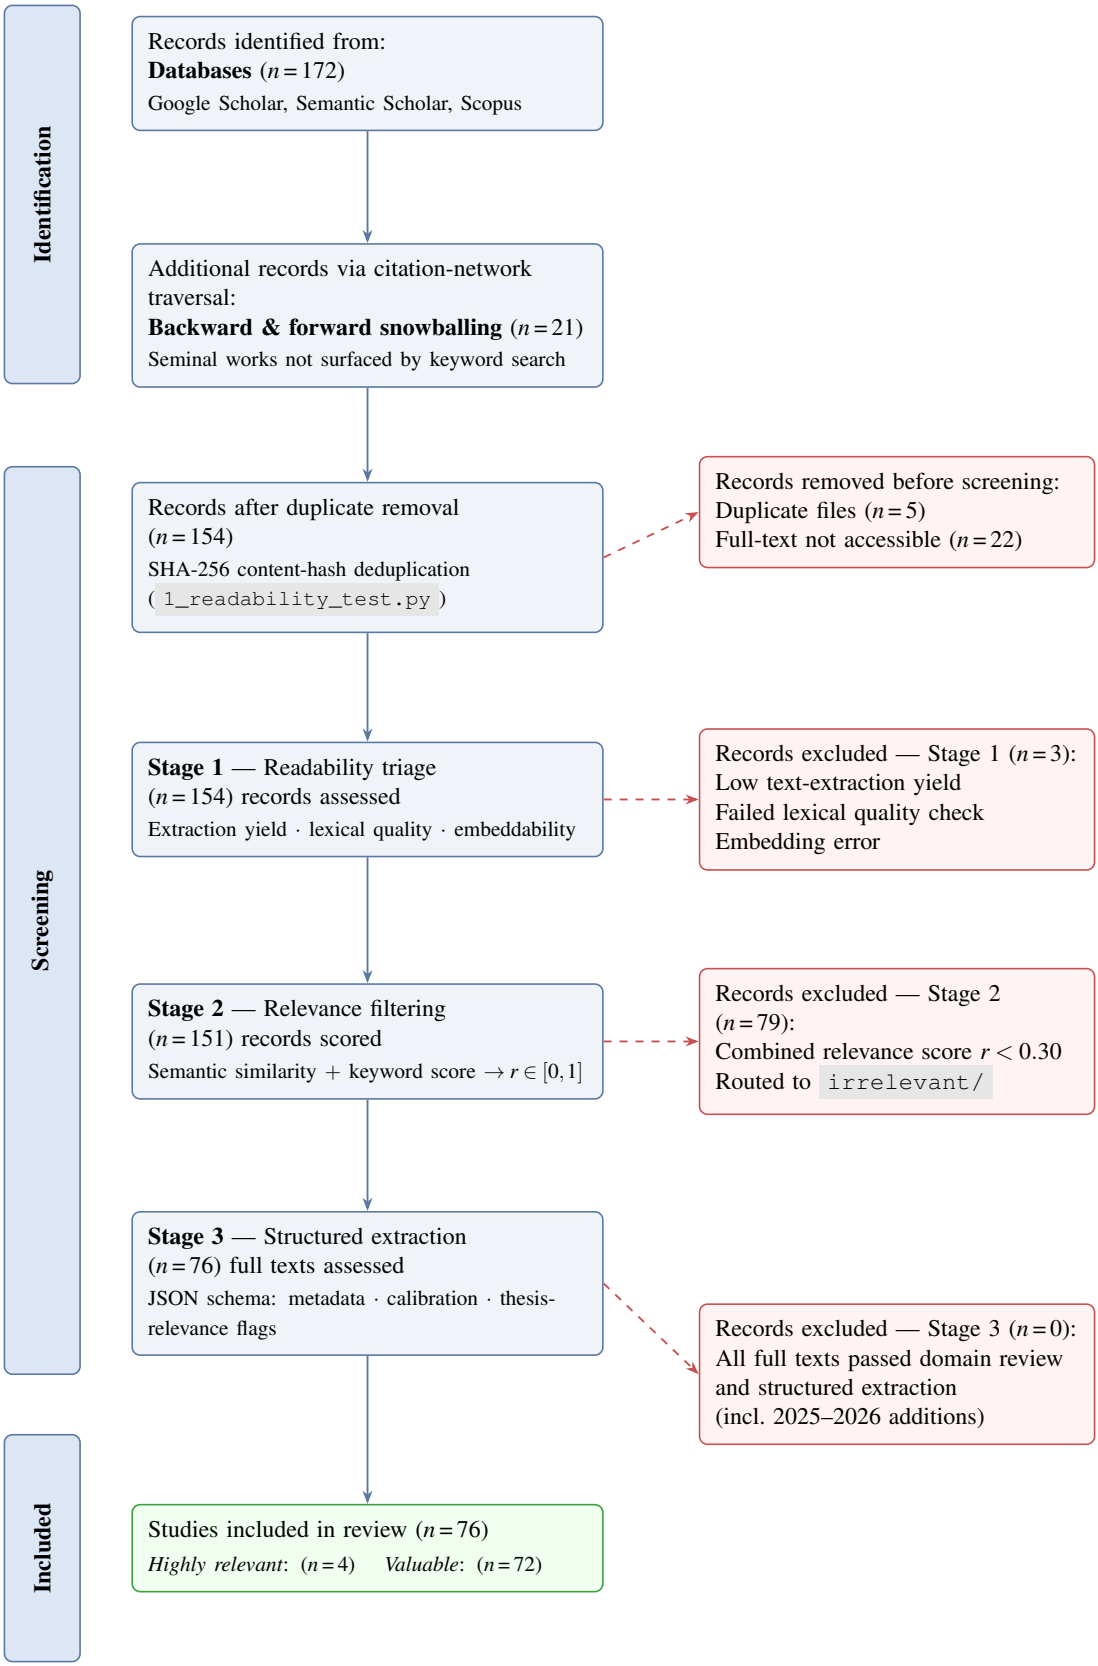
Identification
Records identified from:
Databases (n = 172)
Google Scholar, Semantic Scholar, Scopus
Additional records via citation-network traversal:
Backward & forward snowballing (n = 21)
Seminal works not surfaced by keyword search
Screening
Records after duplicate removal (n = 154)
SHA-256 content-hash deduplication (1_readability_test.py)
Records removed before screening:
Duplicate files (n = 5)
Full-text not accessible (n = 22)
Stage 1 — Readability triage (n = 154) records assessed
Extraction yield · lexical quality · embeddability
Records excluded — Stage 1 (n = 3):
Low text-extraction yield
Failed lexical quality check
Embedding error
Stage 2 — Relevance filtering (n = 151) records scored
Semantic similarity + keyword score → r ∈ [0, 1]
Records excluded — Stage 2 (n = 79):
Combined relevance score r < 0.30
Routed to irrelevant/
Stage 3 — Structured extraction (n = 76) full texts assessed
JSON schema: metadata · calibration · thesis-relevance flags
Records excluded — Stage 3 (n = 0):
All full texts passed domain review and structured extraction
(incl. 2025–2026 additions)
Included
Studies included in review (n = 76)
Highly relevant: (n = 4) Valuable: (n = 72)