Project Portfolio Budgeting as a Markov Decision Process: Environment Design, Baseline Benchmarks, and Optimality Gap Analysis

Mahriar Gharaghani*
Seyed Farid Ghannadpour

Department of Industrial Engineering, Iran University of Science and Technology
Tehran, Iran

Abstract

Allocating a shared capital reserve across a portfolio of simultaneously executing contracted projects is a high-frequency, high-stakes decision for engineering contractors. Classical approaches either assume full foresight or collapse the portfolio to a single project, leaving practitioners with no principled tool for the sequential, multi-project setting under real-world uncertainty. This paper closes that gap by delivering a complete, verified simulation environment that makes the problem tractable for data-driven methods.

We formalise the problem as a Constrained Markov Decision Process in which a portfolio manager allocates budget fractions across N concurrent projects at each discrete period. The state encodes Earned Value Management indicators, S-curve planned progress, milestone payment lookaheads, and a per-project breach-tolerance counter. The transition dynamics reproduce the full FIDIC contract cash-flow cycle: advance payment and recovery, progress-weighted milestone gross payments, retention withholding and release, interest charges on treasury draws, and a two-condition termination gate. The reward signal is the γ-discounted portfolio net cash flow, accumulating to an episode NPV that serves as the primary performance criterion.

The environment is released as a configurable, Gymnasium-compatible package with a SQLite audit ledger that records every allocation, cash-flow event, and breach flag — eliminating the black-box problem inherent to learned policies. To establish a performance baseline and motivate the need for learned agents, we run three deterministic greedy heuristics (Reach-Plan, Catchup-T, and Catchup-Next-T) and a manually computed full-information optimum across four benchmark configurations of increasing portfolio complexity. The results show a systematic and widening gap between the greedy baselines and the optimal solution as portfolio complexity grows, confirming that standard rule-based allocation is insufficient and that a well-defined improvement target exists for reinforcement learning agents. Training such agents and evaluating them against a formal Linear Programming bound are identified as the natural next steps, addressed in subsequent work.

Keywords: project portfolio management; budget allocation under performance uncertainty; constrained Markov decision process; reinforcement learning; FIDIC contracts; earned value management; Gymnasium environment; greedy baseline; optimality gap.

*Corresponding author: mahriargharaghani@gmail.com

1 Introduction

1.1 Motivation and Problem Context

Large infrastructure and engineering projects routinely overrun their budgets and miss their schedules (Flyvbjerg et al., 2003; Merrow, 2011). When a portfolio of such projects executes simultaneously under a shared capital commitment, the risks compound: a performance shortfall on one project propagates through the shared budget to all others, and the decision of how much to allocate to each project at each period can determine whether the portfolio succeeds or fails.

Despite this, the formal treatment of portfolio-level budget allocation as a sequential decision problem is absent from the project portfolio management literature. Research has concentrated on the *selection* stage — which projects to include — while the *operational governance* stage — how to manage a committed portfolio through to completion — has received comparatively little analytical treatment (Archer and Ghasemzadeh, 1999; Cooper et al., 2001; Turner, 2014). The present work fills this gap by delivering a verified simulation environment that makes the problem tractable for data-driven methods and establishing a performance baseline that quantifies the cost of naive allocation.

We consider a portfolio of n projects, each governed by a bilateral contract consistent with FIDIC forms (Fédération Internationale des Ingénieurs-Conseils, 2017), executing concurrently under a relaxed capital commitment. A portfolio manager exercises one recurring decision at each discrete planning period: how much budget to allocate to each active project. The formal problem is to find the allocation policy that maximises discounted net cash flow across the portfolio horizon while respecting contract mechanics.

1.2 The Project Management Decision Hierarchy

The present contribution operates at the portfolio governance layer of the three-tier project management hierarchy (Project Management Institute, 2017). Portfolio selection — which projects to commit to — is treated as an exogenous input, consistent with the two-stage governance model of Turner (2014). At the governance layer, the portfolio manager does not control project execution directly but allocates budget to each project and observes earned value outcomes. Termination is not a deliberate agent action but an endogenous environment mechanic. It is triggered by one of two independent conditions. The first is a hard deadline: if the current episode timestep reaches planned finish + finish delay cap, the project is terminated immediately regardless of progress or tolerance. The second is sustained underperformance: if a projects actual progress falls behind its planned progress by more than progress delay cap, a tolerance counter decrements by one each period the breach persists and resets to its initial value whenever the breach clears; termination fires when the counter reaches zero.

Project termination differs from episode termination a project terminates individually and triggers a settlement cash flow, while the episode ends only when every project has either completed or been terminated, or when the global horizon is reached, or when the portfolio budget is exhausted with no pending inflows.

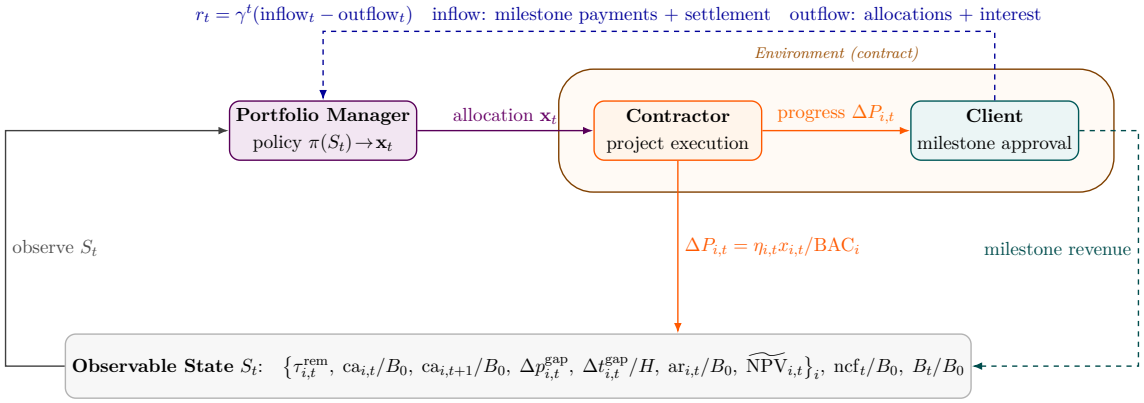


Figure 1: Portfolio governance as a reinforcement learning feedback loop. B_0: initial budget; H: horizon; τ^{rem}: tolerance remaining; ca: catchup allocation; Δp^{gap}: progress gap to next milestone; Δt^{gap}: timestep gap to milestone deadline; ar: required allocation to reach milestone; $\widetilde{\text{NPV}}$: normalised discounted milestone value; ncf: net cashflow.

1.3 Contract Mechanics and Cash-Flow Structure

The financial interaction between the portfolio manager and each contractor is governed by FIDIC contract forms (Fédération Internationale des Ingénieurs-Conseils, 2017). Three mechanisms determine the cash-flow structure.

Milestone-based payment. Payment is contingent on certified progress thresholds. The gross payment at milestone j is a contractually fixed fraction $\phi_{i,j}$ of the contract price $CP_i = (1 + \pi_i) \cdot \text{BAC}_i$. Because payment is contingent on certified progress, the timing of cash inflows depends on the contractor's actual performance and the portfolio manager's allocation decisions. Cash-flow uncertainty in the present model arises from a single source: the stochastic budget-to-progress efficiency $\eta_{i,t}$, which governs how much physical progress a given budget allocation yields in each period. Uncertain productivity induces uncertain milestone timing, which in turn induces uncertain payment arrival — without requiring any separate model of payment delay. This formulation is sufficient to justify the uncertainty framing of the title and admits a cleaner identification of the allocation problem structure than models with multiple independent uncertainty sources.

Advance payment and progressive recovery. At project initiation, an advance payment $A_i = \alpha_i \cdot CP_i$ is disbursed to the contractor. It is recovered through percentage deductions from subsequent interim payment certificates, with commencement threshold ψ_i and deduction rate δ_i^{rec} specified as project-specific contract parameters (Fédération Internationale des Ingénieurs-Conseils, 2017).

Retention. A fraction ρ_i of each milestone payment is withheld as a performance bond. The accumulated retention $R_i^{\text{ret}} = \rho_i \cdot CP_i$ is released in full upon project completion, creating a structural completion incentive embedded in the payment mechanics.

Earned value management as a performance signal. The EVM framework (Fleming and Koppelman, 2010) provides the diagnostic vocabulary for monitoring contractor performance. Define the Budgeted Cost of Work Performed as $\text{BCWP}_{i,t} = \text{EV}_{i,t} \cdot \text{BAC}_i$, where $\text{EV}_{i,t} \in [0, 1]$ is the fraction of work completed, and the Actual

Cost of Work Performed as $\mathrm{ACWP}_{i,t} = \sum_{\tau \le t} x_{i,\tau}$, the cumulative monetary outflow to project i. The Cost Performance Index $\mathrm{CPI}_{i,t} = \mathrm{BCWP}_{i,t}/\mathrm{ACWP}_{i,t}$ measures cost efficiency, while the Schedule Performance Index $\mathrm{SPI}_{i,t} = \mathrm{ES}_{i,t}/\tau_i$ measures schedule efficiency via earned schedule $\mathrm{ES}_{i,t}$—the project-time equivalent of the work accomplished, obtained by inverting the planned S-curve. The Estimate at Completion uses a composite denominator, $\mathrm{EAC}_{i,t} = \mathrm{ACWP}_{i,t} + (\mathrm{BAC}_i - \mathrm{BCWP}_{i,t})/(\mathrm{CPI}_{i,t} \cdot \mathrm{SPI}_{i,t})$, projecting total cost under the assumption that both cost and schedule inefficiency persist. The To-Complete Performance Index $\mathrm{TCPI}_{i,t} = (\mathrm{BAC}_i - \mathrm{BCWP}_{i,t})/(\mathrm{BAC}_i - \mathrm{ACWP}_{i,t})$ measures the cost efficiency required to finish on budget. These four indicators are tracked internally and logged to the audit record each period, but are not directly exposed in the observation vector o_t; instead they are intermediates from which the catchup-allocation and target-milestone features are derived.

Termination. Termination is an endogenous environment mechanic governed by two independent conditions. A *progress delay breach* fires when the lag between planned and actual progress exceeds the contractual tolerance $\delta_i^{\max}$, i.e. $\delta_{i,t} = \mathrm{PV}_i(\tau_i) - \mathrm{EV}_{i,t} > \delta_i^{\max}$. A *hard deadline breach* fires at episode period $t \ge f_i + \Delta_i^{\max}$, where $\Delta_i^{\max}$ is the contractually permitted finish overrun; this condition triggers immediate termination regardless of the cure-period counter. Each period a progress delay breach is active, a project-specific cure-period counter $\theta_{i,t}^{\mathrm{rem}}$ decrements by one; it resets to its initial value θ_i whenever the breach clears. When the counter reaches zero, the project terminates. Upon termination by either condition, a settlement cash flow $S_i = \mathrm{EV}_{i,t} \cdot P_i - \mathrm{CF}_{i,t}^{+}$ enters the portfolio cash balance, where $\mathrm{CF}_{i,t}^{+}$ denotes the cumulative milestone payments already received from the client for project i up to period t. The settlement is positive when earned value at the contract price exceeds payments received, and negative when the portfolio has already collected more than the work delivered warrants.

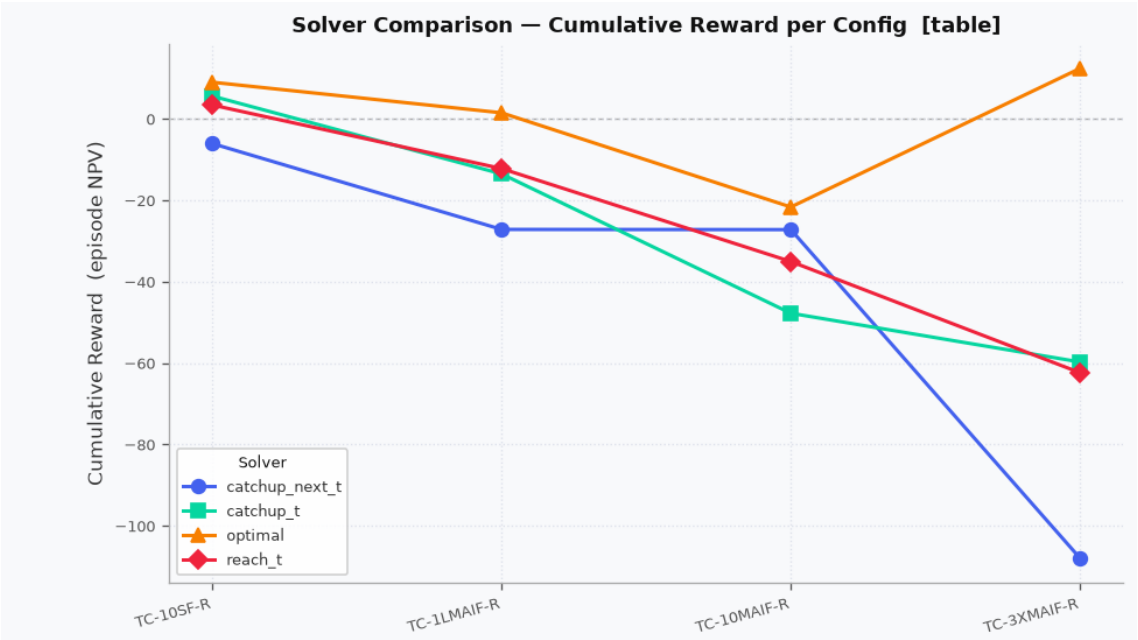
1.4 The Optimality Gap and the Case for Learned Policies

A central empirical question is whether the problem is hard enough to warrant learned policies at all. To answer this, we run three deterministic greedy heuristics against a manually computed full-information optimum across four benchmark portfolio configurations of increasing complexity.

The **Reach-Plan** heuristic (π^{rp}) allocates exactly enough budget each period to bring every active project to its S-curve planned progress, closing the gap entirely. The **Catchup-T** heuristic (π^{cu}) allocates only enough to stay within the contractual progress delay tolerance, minimising expenditure. The **Catchup-Next-T** heuristic ($\pi^{\mathrm{cu+1}}$) is the most risk-averse: it pre-emptively allocates enough this period so that the projected delay next period also remains within tolerance. The **full-information optimum** (π^{*}) is computed manually with all stochastic realisations known in advance, serving as an upper-bound reference.

Figure 7 reports cumulative episode NPV for each method across four configurations. On the simplest configuration (single project, fixed efficiency) the greedy baselines come within a small margin of the optimum. As portfolio complexity grows — more projects, staggered starts, interacting cash flows, and milestone sequencing constraints — the gap widens systematically. On the most complex configuration the best greedy baseline trails the optimum by more than 60 NPV units while one baseline degrades to -107.96, a loss-making outcome. This monotone widening confirms that the allocation problem becomes genuinely difficult at scale and that a well-defined improvement target exists

for reinforcement learning agents.



	10SF-R	1LMAIF-R	10MAIF-R	3XMAIF-R
catchup_next_t	-6.07	-27.17	-27.22	-107.96
catchup_t	5.60	-13.53	-47.77	-59.75
optimal	9.00	1.49	-21.66	12.40
reach_t	3.40	-12.21	-35.06	-62.38

Figure 2: Cumulative episode NPV (reward) for three greedy heuristics and the full-information optimum across four benchmark configurations of increasing portfolio complexity. The widening gap as complexity grows motivates the use of learned allocation policies.

1.5 Contributions and Paper Organisation

The contributions of this paper are threefold.

C1. Environment. A configurable, Gymnasium-compatible CMDP environment encoding the full FIDIC contract cash-flow cycle: advance payment and progressive recovery, progress-weighted milestone payments, retention withholding and release, treasury-draw interest on negative cash balances that persists until the portfolio returns to surplus, and a two-condition termination gate with a project-specific cure-period counter. All environment parameters are specified through a configuration dictionary whose entries are probability distributions, allowing practitioners in any project-based sector—construction, IT, R&D, or otherwise—to instantiate domain-specific environments by fitting those distributions to historical company data. Gymnasium API compliance is verified across 51 checks in 10 test sections; the full compatibility certificate is reproduced in Appendix B. The environment ships with a SQLite audit ledger that records every allocation, cash-flow event, and breach flag at each episode step, making learned policy behaviour fully interpretable.

C2. MDP Formulation. A formal specification of the portfolio budgeting problem as a CMDP: state space, action space, transition dynamics, reward signal, and the observation vector presented in closed form.

C3. Baseline Benchmark and Optimality Gap. Three deterministic greedy heuristics evaluated against a full-information optimum across four benchmark configurations, producing a quantified optimality gap that grows monotonically with portfolio complexity and establishes the performance target for future reinforcement learning agents.

Section 2 reviews the relevant literature across four streams: project management foundations, portfolio management and selection, PPM budgeting and cash-flow modelling, and reinforcement learning for sequential resource allocation. Section 3 develops the MDP formulation formally. Section 4 presents the benchmark evaluation: three deterministic greedy heuristics and a full-information optimum across four portfolio configurations, quantifying the optimality gap that motivates learned agents. Training RL agents on the environment and evaluating them against a formal Linear Programming bound are left as the natural next step, addressed in subsequent work.

2 Literature Review

2.1 Overview and Review Strategy

This section maps the intellectual lineage of the present work across four adjacent knowledge domains: project management (PM) foundations, project portfolio management (PPM) and selection, PPM budgeting and cash-flow modelling, and reinforcement learning (RL) applied to sequential resource-allocation problems. Together these streams define both the theoretical foundation and the gap that motivates the proposed framework. The review follows the PRISMA 2020 reporting guidelines (Page et al., 2021) to ensure transparency and reproducibility in source identification, screening, and inclusion decisions.

2.1.1 Search and Screening Protocol

Literature was retrieved from Google Scholar, Semantic Scholar, and Scopus using a progressively narrowing search strategy. Queries began at the broadest relevant domain (*project management*) and were successively refined through *project portfolio management*, *portfolio budgeting and cash-flow optimisation*, *earned value management under uncertainty*, and *reinforcement learning for project scheduling and resource allocation*. Citation-network traversal — backward and forward snowballing from high-citation anchor papers — was used to recover seminal works not surfaced by keyword search. For instance, a high-citation anchor paper on portfolio optimisation exposes its own foundational references through backward traversal, while forward traversal through its citing works surfaces more recent developments; both layers would have been invisible to keyword search alone.

Records were screened in three sequential automated stages consistent with the PRISMA 2020 flow diagram (Appendix A). The screening pipeline, together with its archived corpus and output artefacts, is implemented as a reproducible processing chain whose repository location and concept-level mapping are documented in Appendix C.

Stage 1 — Identification and Machine Readability. Candidate PDFs were subjected to automated bibliographic and readability triage. Each document was assessed

for text-extraction yield (minimum readable character density across pages), lexical quality (alphabetic-token ratio and mean word length), and sentence-transformer embeddability. Documents failing any criterion were quarantined in a dedicated rejection archive, with a content hash logged to prevent duplicate submissions from re-entering the pipeline.

Stage 2 — Relevance Filtering. Records passing Stage 1 were ranked by a composite relevance score. A semantic similarity component embedded each paper's abstract, introduction, and conclusion using a sentence-transformer model and computed the maximum cosine similarity against twelve thesis-derived query vectors spanning the target knowledge domains. A keyword-scoring component assigned weighted counts of domain-specific high-value and medium-value terms, with a penalty for domain-irrelevant vocabulary. The two components were combined with equal weight (0.5 each) to yield a combined relevance score $r \in [0, 1]$. Papers were then triaged into three cohorts: *highly relevant* ($r \geq 0.45$), *valuable* ($r \geq 0.30$), and *irrelevant* ($r < 0.30$), with the first two cohorts retained for structured extraction.

Stage 3 — Structured Extraction. Each retained paper in the *highly relevant* and *valuable* cohorts was subjected to structured knowledge extraction. The abstract, introduction, and conclusion of each PDF were parsed and submitted to a large language model with a fixed schema prompt, producing a per-paper record covering: problem framing, portfolio level, decision type, stochastic cash-flow treatment, contract mechanics, termination-settlement provisions, modelling paradigm, calibration parameters (CPI/SPI statistics, cost-overrun distributions, S-curve shape parameters), and thesis-relevance flags. The resulting structured corpus subsequently drove an automated synthesis dashboard and a table generator, producing the gap-analysis matrix, domain-distribution charts, and the literature map reproduced in Table 1.

The complete pipeline — processing scripts, the structured record corpus, dashboard outputs, and synthesis tables — is archived within the project repository and constitutes a reproducible, machine-auditable literature review corpus; its directory-level organisation is detailed in Appendix C. The PRISMA flow diagram quantifying records at each stage is provided in Appendix A.

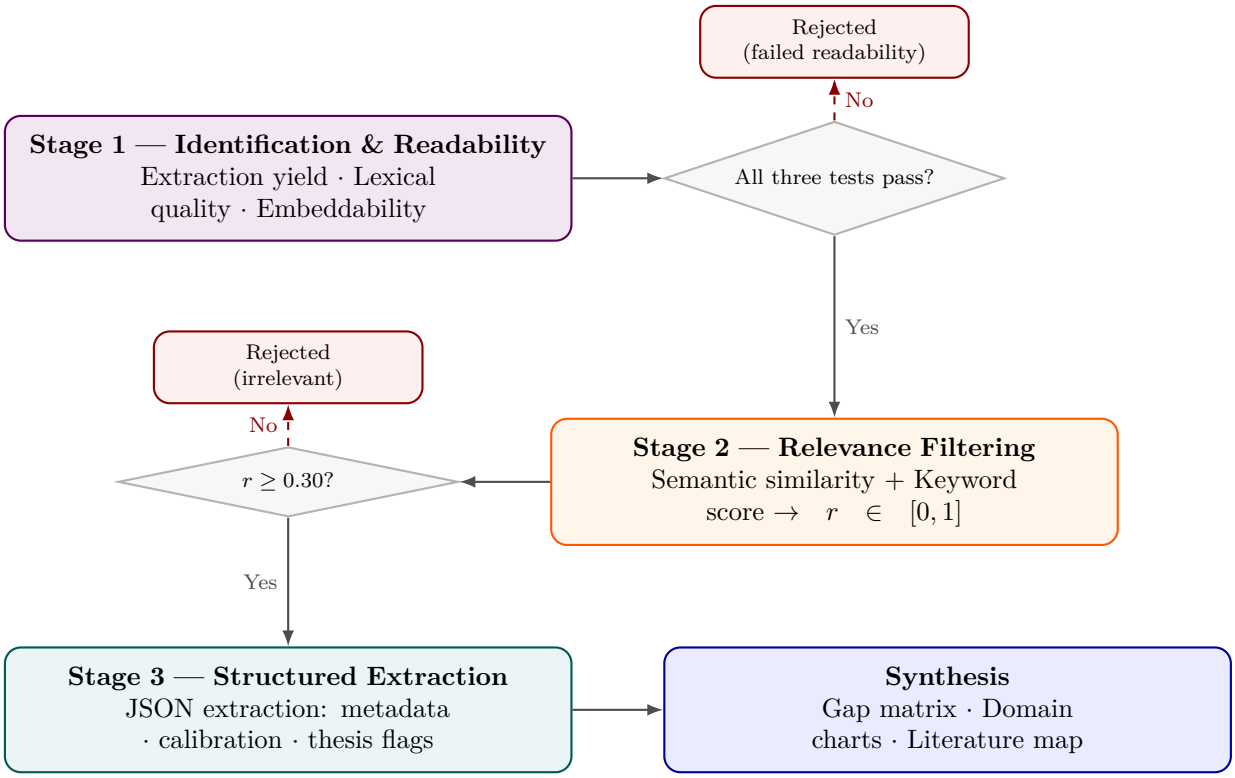


Figure 3: Three-stage PRISMA-aligned screening pipeline, arranged as a serpentine flow. Boxes denote automated processing stages; diamonds denote pass/fail gates; dashed arrows indicate rejection routing. Implementation details and repository locations are provided in Appendix C.

2.2 Project Management: Foundational Concepts

The academic study of project management draws on schedule control (CPM/PERT), earned value management (EVM), and risk modelling as its core quantitative pillars (Fleming and Koppelman, 2010; Project Management Institute, 2021). EVM introduced the Cost Performance Index (CPI) and Schedule Performance Index (SPI) as real-time diagnostic signals relating budgeted cost of work performed to actual expenditure and planned schedule — providing the performance vocabulary that the present model inherits directly. The present model adopts the Earned Schedule correction to SPI (Lipke, 2003), replacing the classical ratio BCWP/BCWS with SPI(t) = ES/AT, which continues to degrade past the planned finish date rather than freezing when BCWS reaches BAC — a known limitation of classical SPI in post-planned-finish recovery periods (Lipke, 2003; **?**). The To-Complete Performance Index (TCPI) is also included as a forward-looking complement to CPI, measuring the cost efficiency required on remaining work to finish within budget. Values of CPI < 1 signal cost overrun; SPI < 1 signals schedule lag. Stochastic extensions of EVM (Acebes et al., 2014; Batselier and Vanhoucke, 2017) and Monte Carlo simulation frameworks (Kwak and Ingall, 2007) have modelled performance uncertainty at the project level, establishing that CPI and SPI follow empirically characterisable distributions over project life. Empirical evidence on contractor performance variability from EVM datasets (Acebes et al., 2014; Batselier and Vanhoucke, 2017) provides the calibration basis for the stochastic efficiency parameter $\eta_{i,t}$ that parameterises the stochastic program at Level 2 and the CMDP environment at Level 3.

S-curve analysis, formalised by Kenley and Wilson (2009) for construction cash flows, established that budget consumption follows characteristic sigmoid profiles whose shape parameters are project-type- and contract-form-dependent. The regularized incomplete Beta function provides a flexible single-family parameterisation of this profile (Barraza et al., 2004), and is adopted directly in the planned expenditure model of the present formulation.

Large-scale empirical evidence on project performance is provided by Merrow (2011) through the Independent Project Analysis database, and by Flyvbjerg et al. (2003); Flyvbjerg (2017) through infrastructure cost-overrun datasets spanning decades and geographies. These datasets inform both the calibration of uncertainty distributions in the present model and the benchmarks against which synthetic environment parameters are validated. Complementary empirical benchmarks are drawn from AACE International and Construction Industry Institute cost-engineering databases (Christensen and Heise, 1993).

Contract governance structures — particularly FIDIC forms (Red, Yellow, and Silver Books) — prescribe the financial interaction between owner and contractor through milestone payment schedules, advance payment and recovery mechanisms, and retention clauses (Fédération Internationale des Ingénieurs-Conseils, 2017). The Estimate at Completion $\mathrm{EAC}_{i,t} = \mathrm{ACWP}_{i,t} + (\mathrm{BAC}_i - \mathrm{BCWP}_{i,t})/\mathrm{CPI}_{i,t}$ extends CPI into a forward-looking cost trajectory signal; Christensen and Heise (1993) established across 155 contracts that this formulation produces the most accurate and stable cost-at-completion forecasts, with reliability increasing beyond the 20% completion mark (Lipke, 2003). The present model adopts EAC as the second termination condition signal, directly parallel to the forecast finish deviation used in Condition 1. Under FIDIC Sub-Clause 14.2, the advance payment is an interest-free mobilisation loan recovered through percentage deductions from interim payment certificates: if not otherwise specified in the Appendix to Tender, deductions commence once cumulative certified interim payments exceed 10% of the Accepted Contract Amount and are applied at a rate of 25% of each subsequent payment certificate until the advance is fully recovered (Fédération Internationale des Ingénieurs-Conseils, 2017). The present model parameterises this mechanism with a project-specific commencement threshold ψ_i and deduction rate δ_i^{rec}, accommodating the full range of particular conditions encountered in practice. These mechanisms translate directly into the cash inflow and outflow dynamics formulated in Section **??**.

2.3 Project Portfolio Management: Selection and Governance

Project portfolio management formalises the problem of allocating an organisation's finite resources across a competing set of candidate projects to maximise strategic value (Archer and Ghasemzadeh, 1999; Cooper et al., 2001). The literature distinguishes two nested decision layers: *portfolio selection* — the binary or mixed-integer optimisation of which projects to include — and *portfolio governance* — the ongoing operational oversight of approved, active projects.

Portfolio selection models have been extensively developed. Markowitz-style mean–variance frameworks adapted from financial portfolio theory (Markowitz, 1952) have been applied to project baskets to trade off expected NPV against variance (Eilat et al., 2006). Multi-criteria decision analysis extensions incorporate strategic alignment, risk ratings, and resource interdependencies (Martinsuo and Lehtonen, 2007; Killen et al., 2008). Metaheuristic approaches — genetic algorithms, particle swarm optimisation — address large combinatorial instances (Zhang et al., 2020). Resource-constrained portfolio selection under capital rationing has been studied by Kavadias and Loch (2003), whose

no-borrowing constraint is adopted directly in the present formulation.

The present work *explicitly adopts the output of portfolio selection as an exogenous input.* The selection problem is treated as solved upstream; the focus is entirely on the operational governance layer that receives a committed portfolio and must manage it through execution. This decomposition is consistent with the two-stage governance model of Turner (2014) and the PMI Standard for Portfolio Management (Project Management Institute, 2017), which separates strategic selection from operational execution oversight. In doing so, this work builds directly on the selection literature while advancing the frontier into the governance domain that has received substantially less formal analytical treatment.

2.4 PPM Budgeting, Cash-Flow Modelling, and Uncertainty

Within portfolio governance, budget allocation across active projects over time represents a high-frequency, high-stakes operational decision. Despite its practical centrality, this specific problem has received limited formal treatment relative to portfolio selection. Existing work addresses project-level cash-flow modelling (Kenley and Wilson, 2009; Park and Peña-Mora, 2004), EVM-based forecasting of individual project cost-at-completion (Kim and Ballard, 2002; Acebes et al., 2014), and multi-project resource levelling (Lova et al., 2000; Van Peteghem and Vanhoucke, 2010). However, these contributions treat budgets as deterministic constraints rather than as dynamic decision variables subject to portfolio-level sequential optimisation.

Cash-flow uncertainty in the present model arises from a single source: the budget-to-progress efficiency $\eta_{i,t}$, which governs how much physical progress a given budget allocation yields in each period (Acebes et al., 2014; Batselier and Vanhoucke, 2017). Uncertain productivity determines whether and when each milestone threshold is crossed, which in turn determines the timing of the contractually fixed payment $R_{i,j}^{\text{net}}$. Payment delays between certification and receipt are not modelled as an independent uncertainty source. This relaxation is defensible: the performance-driven channel is the structurally primary source of cash-flow timing uncertainty in milestone-payment contracts, and isolating it yields a cleaner identification of the allocation problem without loss of the essential uncertainty framing. Other sources of cash-flow uncertainty — scope changes, external market shocks (Doloi, 2012) — are treated as exogenous and outside the scope of the present formulation. Stochastic programming (Birge and Louveaux, 2011) provides the canonical OR framework for optimisation under such uncertainty, and its two-stage variant has been applied to project scheduling under resource uncertainty (Lova et al., 2000). However, multi-stage stochastic programs for portfolio-level sequential budget allocation face scenario-tree explosion as the number of projects and periods increases. The single uncertainty source $\eta_{i,t}$ now branches the scenario tree at every period: if $\eta_{i,t}$ takes K discrete values at each of H periods for each of n projects, the tree produces K^{nH} leaf nodes. The effective planning horizon is itself portfolio-dependent—determined by the latest project finish—compounding the explosion, since H grows with the most delayed project in the committed set. Both effects are formalised in Section 3.5.

The gap in the literature is precise: *no established framework models the portfolio-level budget allocation decision as a sequential stochastic optimisation problem in which a central decision-maker allocates funds across multiple active projects whose cash inflow trajectories are uncertain through stochastic budget-to-progress efficiency, subject to contractual milestone certification windows that bound the speed at which progress converts to cash inflows.* The present work addresses this gap. The objective is pure discounted net cash flow; no deviation penalty enters the objective function. This is a deliberate design

choice: the consequences of sustained deviation — termination and loss of all future milestone payments and retention — are sufficient to make deviation undesirable without a separate penalty term, and including one would introduce an arbitrary scalarisation weight with no natural calibration from contract data (Ng et al., 1999). Symmetrically, the discount factor γ encodes not only time preference but the opportunity cost of tied-up capital in a competitive environment where future projects are foregone while current ones remain incomplete (Damodaran, 2012). This interpretation gives γ a concrete calibration anchor beyond abstract preference — it can be estimated from the firm's observed pipeline redeployment rates or from prevailing cost-of-capital benchmarks. The termination settlement further grounds the model in real asset values: the settlement payment $R_i^{\text{term}} = P_{i,t}^{\text{actual}} \cdot CP_i - (A_i + \sum_{j \in \mathcal{M}_i^{\text{cert}}} R_{i,j}^{\text{net}})$ computes the difference between the contractor's entitlement for work delivered and cumulative inflows already received, which can be positive or negative depending on progress at termination relative to certified payments (Dixit and Pindyck, 1994; Huchzermeier and Loch, 2001), rather than an arbitrarily weighted deviation penalty. A consequence is that the learned policy may produce allocation trajectories that concentrate expenditure around milestone certification thresholds, mirroring documented contractor financial strategy in milestone-payment contracts (Kenley and Wilson, 2009; Kaka and Price, 1993; Arditi and Chotibhongs, 2005).

2.5 Contractor Underperformance, Cure Periods, and Termination

When a contractor fails to meet schedule or cost performance thresholds, international EPC contracts prescribe a structured escalation sequence before termination may be exercised. Under FIDIC Sub-Clause 15.1, the client must issue a formal notice of default and allow a cure period — typically 28 days — during which the contractor is obliged to demonstrate measurable recovery before the client may proceed to termination (Fédération Internationale des Ingénieurs-Conseils, 2017; Ndekugri and Rycroft, 2008). This cure period is not a discretionary grace; it is a contractual obligation whose omission constitutes wrongful termination. Empirical studies of international EPC disputes confirm that failure to observe the cure period is among the most frequently litigated grounds in contractor termination claims (Ndekugri and Rycroft, 2008; Ennis and Burr, 2021).

The behavioural project management literature has studied the conditions under which organisations fail to terminate underperforming projects at the economically rational point. Escalation of commitment (Staw and Ross, 1987) — the tendency to continue investing in failing projects beyond what expected value justifies — has been documented across construction, IT, and infrastructure programmes. Real options theory (Dixit and Pindyck, 1994; Huchzermeier and Loch, 2001) provides the normative counterpart: termination is the exercise of an abandonment option whose value exceeds continuation NPV once performance has deteriorated past a recoverable threshold. Together these streams establish that the decision to persist with or terminate an underperforming project is neither purely contractual nor purely financial — it involves performance signals, contractual obligations, and timing all at once.

The present model integrates all three dimensions into a single environment mechanic governed by two conditions. A *progress delay breach* fires whenever the lag between planned and actual progress exceeds the contractual tolerance $\delta_i^{\max}$, i.e. $\delta_{i,t} = PV_i(\tau_i) - EV_{i,t} > \delta_i^{\max}$, grounded in the FIDIC engineer's authority to issue a notice when progress falls materially behind programme (Fédération Internationale des Ingénieurs-Conseils, 2017). A *hard deadline breach* fires unconditionally at episode period $t \geq f_i + \Delta_i^{\max}$, where $\Delta_i^{\max}$ is the contractually permitted finish overrun, and triggers immediate termination regardless of the cure-period counter state. When a progress delay breach is active,

the cure-period counter $\theta_{i,t}^{\text{rem}}$ decrements by one; it resets to θ_i only when no breach is present. Termination executes when the counter reaches zero. This mechanic is grounded simultaneously in FIDIC contract law (Fédération Internationale des Ingénieurs-Conseils, 2017; Ndekugri and Rycroft, 2008), the real options framing of the abandonment threshold (Dixit and Pindyck, 1994; Huchzermeier and Loch, 2001), the EVM literature's validation of EAC as a stable forward cost signal (Christensen, 1993; Lipke, 2003), and the escalation-of-commitment literature's evidence that observable performance signals must be sustained — not transient — before termination is justified (Staw and Ross, 1987). The counter $\theta_{i,t}^{\text{rem}}$ is part of the observable portfolio state, giving the RL agent an anticipatory signal that reflects contractual reality without granting it discretionary control over the termination outcome. From the portfolio manager's perspective, this signal serves a dual function: it flags projects approaching contractual failure while also marking the point at which continued allocation yields diminishing returns relative to the settlement value of work already delivered. A rational manager reading a counter approaching zero in a high-opportunity-cost environment would recognise that the real value of future milestone payments is eroding while the settlement value of work in place is fixed in nominal terms — creating a natural reallocation incentive that the model captures through the interaction of γ and $S_i = EV_{i,t} \cdot P_i - \text{CF}_i^+$ without requiring explicit inflation modelling (Dixit and Pindyck, 1994; Fédération Internationale des Ingénieurs-Conseils, 2017).

2.6 Reinforcement Learning for Sequential Resource Allocation

Reinforcement learning provides a principled framework for sequential decision-making under uncertainty (Sutton and Barto, 2018). The MDP formalisation — state, action, transition, reward — maps naturally onto the portfolio budget allocation problem, and the constrained MDP extension (Altman, 1999) accommodates the hard budget feasibility constraints that the portfolio manager must satisfy at every period.

RL has been applied to project scheduling (Peng and Huang, 2021; Liu et al., 2023), resource-constrained project scheduling problems (RCPSP) (Tassel et al., 2023), and supply chain inventory management (Giannoccaro and Pontrandolfo, 2002). In each domain, RL demonstrates advantages over classical OR methods when state spaces are large, transition dynamics are stochastic, the problem is non-stationary over the decision horizon, and the horizon length is itself uncertain at the outset of planning. These are precisely the characteristics of the portfolio budget allocation problem at Levels 2 and 3.

Proximal Policy Optimisation (Schulman et al., 2017) is adopted as the solution algorithm for its stability in continuous action spaces, its sample efficiency relative to earlier policy-gradient methods, and its established use in operations management domains. Domain randomisation (Tobin et al., 2017) — randomising environment parameters during training to improve out-of-sample policy robustness — is adopted to ensure the learned policy generalises across the range of project types, contract structures, and market conditions represented in the calibration data. Walk-forward validation (Cerqueira et al., 2020) is used as the empirical evaluation protocol to guard against data leakage given the time-series structure of project cash-flow data.

Direct application of RL to portfolio-level budget allocation remains absent from the literature. The present work provides the first formulation of this problem as a CMDP and the first empirical demonstration that a PPO agent can learn non-trivial allocation policies in a calibrated synthetic environment.

2.7 Synthesis: Positioning the Present Contribution

Table 1 maps representative works across the four literature streams against the key modelling dimensions of the present framework, as generated by the automated synthesis pipeline described in Section 2.1.1. Three structural gaps emerge from this mapping.

Gap 1 — Selection versus operational governance. The transition from portfolio *selection* to portfolio *operational governance* has not been modelled as a sequential decision problem. Selection literature assumes away post-selection management; project scheduling literature ignores portfolio interdependencies and the shared capital constraint. No existing work closes both simultaneously.

Gap 2 — Project-level versus portfolio-level uncertainty propagation. Existing stochastic models of project cash-flow uncertainty operate at the individual project level and do not propagate uncertainty upward into a portfolio-level decision architecture governed by a shared budget constraint. The feedback loop between a project's stochastic efficiency trajectory and its portfolio-level funding entitlement remains unmodelled.

Gap 3 — Absence of endogenous termination mechanics. The joint modelling of a progress-delay breach condition tracked against a contractual tolerance, a decrementing cure-period counter governing grace-period exhaustion, an independent hard-deadline breach condition that triggers immediate termination regardless of the cure-period state, and a termination settlement cash flow reconciling earned value against cumulative payments received, has not been embedded within a portfolio-level sequential decision framework in which budget allocation is the sole control variable.

These gaps are not isolated oversights; they reflect an incomplete progression along the path toward a deployable portfolio-level decision-support system. Figure 4 situates this progression as three sequential steps.

Step 1 — domain grounding — is well established by the literature reviewed above: contract mechanics, EVM performance signals, S-curve expenditure profiles, and portfolio selection models collectively provide the inputs and boundary conditions that any operational governance model requires.

Step 2 — tractable problem formulation — is the gap this paper closes. The contribution is a verified, configurable CMDP environment that encodes the full FIDIC cash-flow cycle and a two-condition termination mechanic, released as a Gymnasium-compatible package with a complete SQLite audit ledger. Baseline greedy heuristics are benchmarked against a full-information optimum, producing a quantified optimality gap that grows monotonically with portfolio complexity and establishes the performance target for future learned agents.

Step 3 — validated deployment within incumbent project management information systems such as Oracle Primavera P6 — remains future work, contingent on the trained agents that Step 2 now makes possible.

The present work addresses all three gaps and completes Step 2 of this progression. The Gymnasium-compatible CMDP environment — with parameterisable contract mechanics, stochastic budget-to-progress efficiency as the single uncertainty source, and an endogenous termination mechanic — constitutes the primary methodological contribution. Baseline heuristics are evaluated against a full-information optimum in Section 4, providing a quantified optimality gap that grows monotonically with portfolio complexity and defines the improvement target for future RL agents.

Table 1: Representative literature mapped against modelling dimensions of the present framework. *Decision type*: S = selection, Sc = scheduling, B = budgeting. ✓ = fully addressed; ○ = partially addressed; – = not addressed.

Work	Portfolio level	Decision type	Sequential decisions	Stochastic cash flows	Contract mechanics	Termination settlement
Yamakawa et al. (2019)	✓	S	–	○	–	–
Kosztyán et al. (2022)	✓	S	–	○	–	–
Touran (2006)	✓	Sc/B	–	○	–	–
Touran (2010)	✓	Sc/B	–	○	–	–
Ghapanchi et al. (2012)	✓	S	–	○	–	–
Pourahmadi et al. (2015)	✓	S	–	○	–	–
El-Sabek and McCabe (2018)	–	B	–	–	✓	–
Beseiso and Kumar (2021)	✓	S	–	○	–	–
Karimi et al. (2021)	✓	S	–	○	–	–
Bahroun et al. (2023)	–	Sc	✓	–	–	–
Ranjbar et al. (2022)	✓	S/Sc	–	–	○	–
Szwarcfiter et al. (2023)	✓	Sc	✓	○	–	–
Asghari et al. (2023)	–	–	✓	○	–	–
Mazhari and Tajik (2024)	✓	S/B	✓	–	–	–
Jiang et al. (2024)	✓	–	✓	○	–	–
Nabot et al. (2025)	✓	Sc	✓	–	✓	✓
Darvish and Sepehri (2025)	✓	B	✓	○	○	✓
Gholami et al. (2025)	✓	Sc	✓	○	–	✓
Kehinde (2025)	✓	B	✓	✓	–	–
Ajibade (2025)	✓	Sc	✓	○	–	–
Huddar and Joshi (2025)	✓	B	✓	○	–	–
Akbari and Shirazi (2026)	✓	B	✓	○	✓	✓
Ebnerasoul et al. (2023)	✓	S	–	○	–	–
Noroozi et al. (2016)	✓	B	–	○	–	–
Jiang et al. (2023)	–	B	✓	–	–	–
Present work	✓	**B**	✓	✓	✓	✓

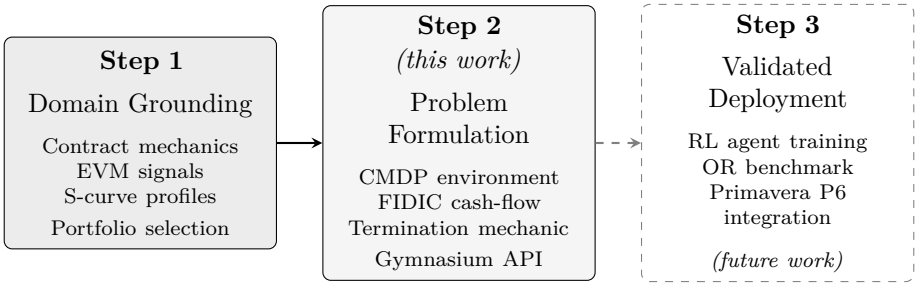


Figure 4: Maturity progression toward a deployable portfolio-level decision-support system. Step 2, completed by the present work, enables Step 3. Dashed border and arrow denote future work.

3 Problem Formulation

This section develops the formal problem specification in three parts. Section 3.1 and 3.2 define the environment primitives shared across all subsequent analysis: project dynamics, FIDIC contract cash-flow mechanics, EVM performance signals, and the termination mechanic. Section 3.3 consolidates all parameters. Sections 3.4 and 3.5 frame the deterministic upper bound (Level 1 MILP) and establish why the stochastic extension is computationally intractable (Level 2), motivating the CMDP reformulation. Section 3.6 specifies the CMDP and its Gymnasium-compatible implementation — the primary deliverable of this paper. Training a policy on this environment is the natural next step and is addressed in subsequent work.

3.1 Problem Statement and Scope

The problem is operational, not strategic. A portfolio of n contracted projects is fixed: commitments are signed, start dates are scheduled, and contracts are parameterised consistently with FIDIC Red, Yellow, and Silver Book forms (Fédération Internationale des Ingénieurs-Conseils, 2017). The decision-maker — a portfolio manager, board, or programme director — exercises a single recurring lever at each discrete planning period $t \in \mathcal{T} = \{1, \ldots, H\}$: the budget $x_{i,t} \geq 0$ directed to each active project i. The planning horizon H is not a fixed configuration parameter but is determined endogenously as $H = \max_i f_i - \min_i s_i + 1$, where s_i and f_i are the planned start and finish of project i; the index is shifted so that period $t = 1$ aligns with the earliest project start. All outcomes of interest — milestone achievement, cash inflows, schedule adherence, and project termination — flow through this allocation decision. Portfolio selection is treated as an exogenous upstream input (Archer and Ghasemzadeh, 1999).

Four structural assumptions underlie the entire formulation. (i) *Independent bilateral contracts*: projects interact only through the shared cash balance; no resource or workforce sharing occurs, consistent with standard EPC contracting practice. (ii) *Budget as the sole control*: the contractor's internal resource decisions are absorbed into a stochastic efficiency parameter $\eta_{i,t}$; the portfolio manager observes how much physical progress a budget allocation yields, not how it is achieved. $\eta_{i,t}$ is the sole stochastic primitive: it is unobservable at the time allocation $x_{i,t}$ is committed, and its realisation determines whether and when each milestone threshold is crossed, which in turn determines the timing of cash inflows. The timing of that allocation — specifically how quickly milestones are certified — determines the discounted value of inflows under $\gamma < 1$, making the scheduling of spend, not just its total magnitude, the operative decision variable for a

return-maximising portfolio manager (Damodaran, 2012). (iii) *Discrete monthly periods*: $\mathcal{T} = \{1, \ldots, H\}$, where H is determined dynamically from the portfolio (see Section 3.2.1), reflects the standard monthly governance cycle of portfolio management. A relaxed budget regime ($\kappa = B_1 / \sum_i \mathrm{BAC}_i \gg 1$) is adopted throughout, so the liquidity constraint is non-binding and the formulation isolates allocation timing and performance management as the operative decision dimensions. (iv) *FIDIC contract mechanics*: milestone payments, advance payment and recovery, retention, and termination settlement are all calibrated parameters, not fixed constants, accommodating the full range of particular conditions encountered in practice.

3.2 Environment Definition

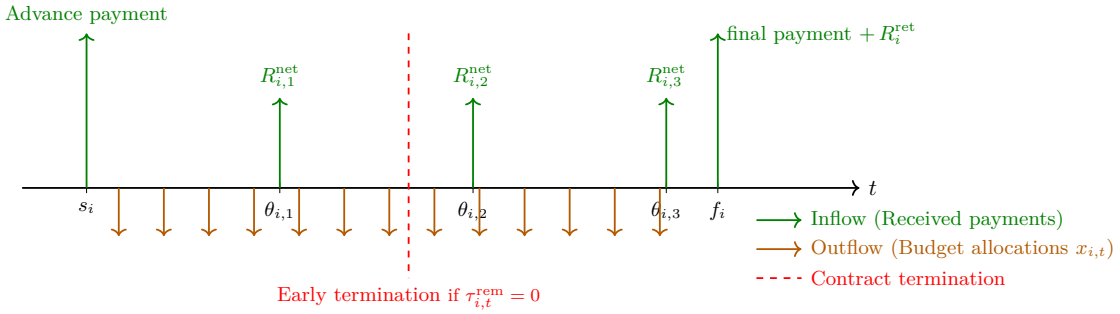


Figure 5: Cash-flow timeline for project i under FIDIC contract mechanics.

3.2.1 Sets, Indices, and the Decision Variable

$\mathcal{P}$	Projects $\{1, \ldots, n\}$, indexed by i; all committed at horizon start.
$\mathcal{T}$	Planning periods $\{1, \ldots, H\}$, indexed by t (calendar months), where $H = \max_{i \in \mathcal{P}} f_i - \min_{i \in \mathcal{P}} s_i + 1$ is derived endogenously from the committed portfolio; period $t = 1$ corresponds to the calendar month of the earliest project start $\min_i s_i$.
$\mathcal{M}_i$	Ordered milestones $\{1, \ldots, M_i\}$ for project i, indexed by j, with progress thresholds $\theta_{i,1} < \cdots < \theta_{i,M_i} = 1$.
$\mathcal{A}_i^{\text{plan}}$	Planned execution window $\{s_i, \ldots, f_i\}$; anchors S-curve computations.
$\mathcal{P}_t^{\text{active}}$	Projects active at period t; fixed at Level 1, endogenous at Levels 2–3.
$\mathcal{M}_i^{\text{cert}}(t)$	Milestones of project i certified by period t: $\{j : \theta_{i,j} \leq P_{i,t}^{\text{actual}}$ and $t \geq e_{i,j}\}$, where $e_{i,j} \in \mathcal{T}$ is the contractually specified earliest certification period for milestone j.
$x_{i,t} \geq 0$	**Decision variable**: budget allocated to project i at period t. The period-t action vector is $\mathbf{x}_t = (x_{i,t})_{i \in \mathcal{P}_t^{\text{active}}}$.
$B_t \geq 0$	Portfolio cash balance at the start of period t; B_1 is given.

3.2.2 Progress and Expenditure Dynamics

Budget $x_{i,t}$ converts to physical progress through the efficiency scalar $\eta_{i,t}$, which absorbs all sources of conversion uncertainty (contractor productivity, supply chain, weather):

$$P_{i,t}^{\text{actual}} = P_{i,t-1}^{\text{actual}} + \eta_{i,t} \cdot \frac{x_{i,t}}{\text{BAC}_i}, \qquad \text{ACWP}_{i,t} = \text{ACWP}_{i,t-1} + x_{i,t}, \tag{1}$$

with $P_{i,0}^{\text{actual}} = \text{ACWP}_{i,0} = 0$. Here BAC_i is the Budget at Completion — the total planned contractor expenditure for project i, which serves as the EVM cost basis and the normalisation denominator for all progress and performance calculations. At Level 1, $\eta_{i,t}$ is replaced by a known deterministic value $\bar{\eta}_{i,t}$; at Levels 2–3 it is the sole stochastic primitive, realised only after $x_{i,t}$ is committed.

The planned cumulative expenditure fraction follows a normalised S-curve (Kenley and Wilson, 2009; Barraza et al., 2004) parameterised by shape parameters (a_i, b_i):

$$P_{i,t}^{\text{plan}} = I_{\tau_{i,t}^c}(a_i, b_i), \qquad \tau_{i,t} = \frac{t - s_i + 1}{D_i^{\text{plan}}}, \quad \tau_{i,t}^c = \min\{1, \max\{0, \tau_{i,t}\}\}, \tag{2}$$

where $I_\tau(\cdot,\cdot)$ is the regularized incomplete Beta function and $D_i^{\text{plan}} = f_i - s_i + 1$ is the planned duration. This profile is a fixed parameter derived at commitment; it is never a decision variable.

3.2.3 Contractual Cash-Flow Mechanics

Payment inflows follow FIDIC mechanics (Fédération Internationale des Ingénieurs-Conseils, 2017). At project initiation ($t = s_i$), the portfolio manager receives the advance payment $A_i = \alpha_i \cdot CP_i$, where α_i is the advance ratio and $CP_i = \text{BAC}_i(1 + \pi_i)$ is the contract price with profit margin π_i. The contract price is the sole payment accumulation baseline: all milestone revenues, advance payment, retention, and final settlement sum to CP_i. All EVM performance signals and progress calculations are normalised against BAC_i.

Milestone j is certified at the first period t at which $P_{i,t}^{\text{actual}} \geq \theta_{i,j}$ *and* $t \geq e_{i,j}$, where $e_{i,j}$ is the contractually specified earliest certification period. Under FIDIC Sub-Clause 14.3, interim payment applications may not be submitted before the intervals stated in the Appendix to Tender (Fédération Internationale des Ingénieurs-Conseils, 2017); $e_{i,j}$ parameterises this constraint directly, preventing a portfolio manager from accelerating cash inflows beyond the contractual payment schedule regardless of how quickly physical progress is achieved. The net payment received, after retention withholding at rate ρ_i and advance recovery deduction at rate δ_i^{rec} (commencing once cumulative certified payments exceed threshold ψ_i), is:

$$R_{i,j}^{\text{net}} = \phi_{i,j} \cdot CP_i(1 - \rho_i) - \delta_i^{\text{rec}} \cdot \phi_{i,j} \cdot CP_i \cdot \mathbf{1}\left\{\textstyle\sum_{k \leq j} \phi_{i,k} \geq \psi_i\right\}, \tag{3}$$

where $\phi_{i,j}$ is the contractual weight of milestone j with $\sum_j \phi_{i,j} = 1$. This parameterisation is consistent with FIDIC Sub-Clause 14.2, under which advance recovery deductions commence once cumulative certified interim payments exceed a specified fraction of the contract price (Fédération Internationale des Ingénieurs-Conseils, 2017). Upon project completion ($P_{i,t}^{\text{actual}} = 1$), the accumulated retention $R_i^{\text{ret}} = \rho_i \cdot CP_i$ is released together with the final milestone payment. The payment identity $A_i + \sum_j R_{i,j}^{\text{net}} + R_i^{\text{ret}} = CP_i$ holds exactly.

Milestone payments are received in the same period they are certified (Relaxation B). The aggregate period inflow is:

$$V_t = \sum_{i,j} R^{\text{net}}_{i,j} \cdot \mathbf{1}\{\tau^{\text{cert}}_{i,j} = t\} + \sum_i A_i \cdot \mathbf{1}\{t = s_i\} + \sum_{i \in \mathcal{P}^{\text{term}}_t} R^{\text{term}}_i, \tag{4}$$

where $\mathcal{P}^{\text{term}}_t$ collects projects newly terminated at t, with settlement R^{term}_i defined in Section 3.2.5 below.

3.2.4 Performance Monitoring: EVM Signals

The Earned Value Management framework (Fleming and Koppelman, 2010) provides the diagnostic vocabulary used to monitor contractor performance and to drive the termination mechanic. The Earned Value at period t is $\text{EV}_{i,t} = P^{\text{actual}}_{i,t} \cdot \text{BAC}_i$ (equivalently, the Budgeted Cost of Work Performed, BCWP), expressed in cost units against the BAC_i baseline. Three derived indices are computed each period:

$$\text{SPI}_{i,t} = \frac{\text{ES}_{i,t}}{t - s_i + 1}, \tag{5}$$

$$\text{CPI}_{i,t} = \frac{\text{EV}_{i,t}}{\text{ACWP}_{i,t}} = \frac{P^{\text{actual}}_{i,t} \cdot \text{BAC}_i}{\text{ACWP}_{i,t}}, \tag{6}$$

$$\text{EAC}_{i,t} = \text{ACWP}_{i,t} + \frac{\text{BAC}_i - \text{EV}_{i,t}}{\text{CPI}_{i,t}}. \tag{7}$$

The Earned Schedule $\text{ES}_{i,t}$ is the planned time at which the s-curve (2) would have reached the current actual progress level $P^{\text{actual}}_{i,t}$; formally it is the value $\tau^* \in [0, D^{\text{plan}}_i]$ satisfying $I_{\tau^*/D^{\text{plan}}_i}(a_i, b_i) = P^{\text{actual}}_{i,t}$, computed by binary search over the Beta CDF. Unlike classical SPI = BCWP/BCWS, which freezes once BCWS reaches BAC_i at the planned finish, $\text{SPI}_{i,t} = \text{ES}_{i,t}/(t - s_i + 1)$ degrades continuously as real time passes without progress, correctly distinguishing a project one period late from one three periods late (Lipke, 2003). This correction is adopted throughout the environment; classical SPI is not used.

$\text{SPI}_{i,t} < 1$ signals schedule lag; $\text{CPI}_{i,t} < 1$ signals cost overrun. $\text{EAC}_{i,t}$ is the forward-looking Estimate at Completion: it projects the total cost of the project at the current cost performance rate by scaling remaining planned work by the inverse of $\text{CPI}_{i,t}$, directly analogous to the forecast finish date computation used in Condition 1 of the termination mechanic (Section 3.2.5). Christensen (1993) established that this formulation produces the most accurate and stable cost-at-completion forecasts, with reliability increasing beyond the 20% completion mark (Lipke, 2003). The To-Complete Performance Index $\text{TCPI}_{i,t}$ (**??**) complements $\text{CPI}_{i,t}$ with a forward-looking cost signal: it measures the cost efficiency that must be sustained on remaining work to finish within BAC_i. $\text{TCPI}_{i,t} > 1$ indicates that future performance must exceed historical performance; values above 1.10 are widely considered unrecoverable in practice (Fleming and Koppelman, 2010; Christensen, 1993). While $\text{CPI}_{i,t}$ reflects what has happened, $\text{TCPI}_{i,t}$ reflects what must happen — the two together give the agent a complete cost picture at each period. All four indices are observable state variables; their stochastic dynamics at Levels 2–3 are calibrated from empirical EVM datasets (Acebes et al., 2014; Batselier and Vanhoucke, 2017).

3.2.5 Termination Mechanic

Termination is an endogenous environment mechanic, not an agent action. At each period t, the environment evaluates two conditions for every active project i. Both are

forward-looking trajectory estimates, symmetric in structure to one another.

Condition 1 — Progress delay breach. At each period t the planned cumulative progress is evaluated from the S-curve (2) at project-time $\tau_i = t - s_i$:

$$PV_{i,t} = I_{\tau_i/D_i^{\text{plan}}}(a_i, b_i). \tag{8}$$

The progress lag is $\delta_{i,t} = PV_{i,t} - EV_{i,t}$, where $EV_{i,t} = P_{i,t}^{\text{actual}}$ is actual progress. Condition 1 activates when this lag exceeds the contractual progress delay cap:

$$\delta_{i,t} \;>\; \delta_i^{\max}, \tag{9}$$

where $\delta_i^{\max} \in [0, 1]$ is the project-specific tolerance fixed at commitment. This condition is grounded in the FIDIC engineer's authority to issue a notice when progress falls materially behind programme (Fédération Internationale des Ingénieurs-Conseils, 2017).

Condition 2 — Hard deadline breach. A hard deadline fires unconditionally at:

$$t \;\geq\; f_i + \Delta_i^{\max}, \tag{10}$$

where $\Delta_i^{\max} \in \mathbb{Z}^+$ is the maximum permitted finish overrun in periods. This condition triggers immediate termination regardless of the cure-period counter state, reflecting the FIDIC Sub-Clause 8.7 entitlement to terminate when time for completion has definitively expired (Fédération Internationale des Ingénieurs-Conseils, 2017).

Cure-period counter and termination execution. When Condition 1 is active the environment decrements a project-specific cure-period counter, modelling the FIDIC Sub-Clause 15.1 obligation to allow a recovery interval before termination may be executed (Fédération Internationale des Ingénieurs-Conseils, 2017; Ndekugri and Rycroft, 2008):

$$\theta_{i,t}^{\text{rem}} = \begin{cases} \theta_{i,t-1}^{\text{rem}} - 1 & \text{if (9) holds,} \\ \theta_i & \text{if (9) lapses.} \end{cases} \tag{11}$$

The counter initialises at $\theta_{i,s_i}^{\text{rem}} = \theta_i$, where $\theta_i \in \mathbb{Z}^+$ is the project-specific tolerance fixed at commitment. Soft termination executes when the counter reaches zero: $t_i^{\text{term}} = \min\{t : \theta_{i,t}^{\text{rem}} = 0\}$. Hard termination (10) overrides the counter and executes immediately. The counter is part of the observable portfolio state; the agent influences it indirectly through allocations that drive progress relative to $PV_{i,t}$.

Upon termination at t_i^{term}, the net settlement payment is:

$$R_i^{\text{term}} = P_{i,t_i^{\text{term}}}^{\text{actual}} \cdot CP_i \;-\; \left(A_i + \sum_{j \in \mathcal{M}_i^{\text{cert}}(t_i^{\text{term}})} R_{i,j}^{\text{net}} \right). \tag{12}$$

The first term is the contractor's entitlement for work physically delivered: the fraction of contract price proportional to actual progress at termination. The second term is the cumulative inflows already received: the advance payment A_i plus all net milestone payments $R_{i,j}^{\text{net}}$ certified up to and including t_i^{term}, where $R_{i,j}^{\text{net}}$ is as defined in (3). The settlement is the difference between what the contractor is owed for delivered scope and what they have already collected.

Positive R_i^{term} arises when delivered work exceeds cumulative receipts — typically when termination occurs just below a milestone threshold, so the contractor has performed

work not yet captured in a certified payment. Negative R_i^{term} constitutes a net contractor obligation — typically when a large advance has been disbursed against work that was not subsequently delivered, or when certified milestone payments have run ahead of physical progress. In either case the settlement is fully determined by contract parameters and the progress level at termination; no penalty term or external calibration is required. This grounds the model's financial signals in the real asset value of work in place (Dixit and Pindyck, 1994) rather than in arbitrarily weighted deviation terms.

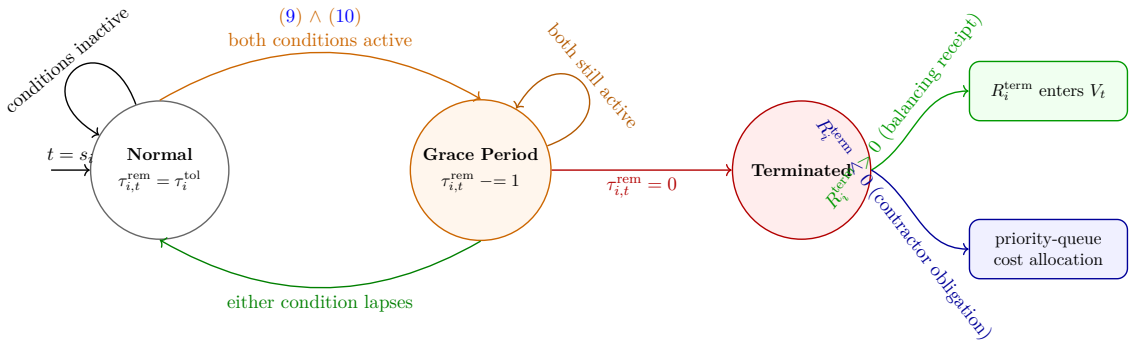


Figure 6: Finite-state representation of the termination mechanic for project i.

3.2.6 Portfolio Cash Balance

The portfolio cash balance evolves as:

$$B_{t+1} = B_t - \sum_{i \in \mathcal{P}_t^{\text{active}}} x_{i,t} + V_t. \tag{13}$$

Under the relaxed budget regime ($\kappa \gg 1$), the liquidity constraint $B_t \geq 0$ is non-binding by construction throughout all experiments. The shared balance B_t is the sole coupling across otherwise independent project dynamics and is retained in the formulation for completeness; it becomes binding in the tight and scarce regimes ($\kappa \approx 1$ and $\kappa \ll 1$) studied in sensitivity analysis.

3.2.7 Observable State and Shared Objective

The observation vector returned to the policy at period t is a flattened concatenation of seven per-project signals followed by two portfolio-level scalars, giving dimension $7n + 2$:

$$o_t = \Big(\big\{ \tilde{\theta}^{\text{rem}}_{i,t},\ \tilde{x}^{\text{cu}}_{i,t},\ \tilde{x}^{\text{cu}+1}_{i,t},\ \Delta\tilde{\psi}^{*}_{i,t},\ \Delta\tilde{\tau}^{*}_{i,t},\ \tilde{x}^{*}_{i,t},\ \widetilde{\text{NPV}}^{*}_{i,t} \big\}_{i \in \mathcal{P}},\ \widetilde{\Delta B}_t,\ \tilde{B}_t \Big), \tag{14}$$

where all components are normalised to $[-1, 1]$ or $[0, 1]$ as detailed in Table 2. The per-project signals encode: the remaining cure-period tolerance $\tilde{\theta}^{\text{rem}}_{i,t}$; the current and next-period catchup allocations $\tilde{x}^{\text{cu}}_{i,t}$, $\tilde{x}^{\text{cu}+1}_{i,t}$ (minimum spend to avoid a breach this period and next); the target milestone progress gap $\Delta\tilde{\psi}^{*}_{i,t}$, timestep gap $\Delta\tilde{\tau}^{*}_{i,t}$, required allocation $\tilde{x}^{*}_{i,t}$, and discounted NPV $\widetilde{\text{NPV}}^{*}_{i,t}$ for the next uncertified milestone. The portfolio-level scalars are the last-period net cash-flow $\widetilde{\Delta B}_t$ and the available budget $\tilde{B}_t$, both normalised by B_0. All components are observable at decision time. The sole unobservable primitive is $\eta_{i,t}$, realised only after $x_{i,t}$ is committed; it drives the stochastic transition but is not part of the observation.

Index	Symbol	Clip	Description
$7i$	$\tilde{\theta}^{\text{rem}}_{i,t}$	$[0, \theta_i]$	Tolerance remaining
$7i+1$	$\tilde{x}^{\text{cu}}_{i,t}$	$[0,1]$	Catchup alloc. (current), normalised by B_0
$7i+2$	$\tilde{x}^{\text{cu}+1}_{i,t}$	$[0,1]$	Catchup alloc. (next period), normalised by B_0
$7i+3$	$\Delta\tilde{\psi}^{*}_{i,t}$	$[0,1]$	Target milestone progress gap
$7i+4$	$\Delta\tilde{\tau}^{*}_{i,t}$	$[-1,1]$	Target milestone timestep gap, normalised by T
$7i+5$	$\tilde{x}^{*}_{i,t}$	$[0,1]$	Target required allocation, normalised by B_0
$7i+6$	$\widetilde{\text{NPV}}^{*}_{i,t}$	$[-1,1]$	Target milestone NPV, normalised by $\max_i \lvert\text{NPV}^{*}_{i,t}\rvert$
$7n$	$\widetilde{\Delta B}_t$	$[-2,2]$	Last-period net cash-flow, normalised by B_0
$7n+1$	$\tilde{B}_t$	$[0,2]$	Available budget, normalised by B_0

Table 2: Observation vector components. Per-project block i occupies indices $7i$ through $7i+6$. All normalisation denominators use $B_0 = \kappa \sum_i \text{BAC}_i$.

The decision-maker maximises discounted net cash flow over the horizon:

$$\max_{\{x_{i,t}\}} \quad \mathbb{E}\left[\sum_{t=1}^{H} \gamma^{t-1}\left(V_t - \sum_{i \in \mathcal{P}_t^{\text{active}}} x_{i,t}\right)\right], \tag{15}$$

subject to, for all $i \in \mathcal{P}$, $t \in \mathcal{T}$:

$$x_{i,t} \geq 0, \tag{16}$$

$$\sum_{i \in \mathcal{P}_t^{\text{active}}} x_{i,t} \leq B_t, \tag{17}$$

$$x_{i,t} = 0 \quad \text{if status}_{i,t} \neq \texttt{active}. \tag{18}$$

Constraint (17) is retained for generality; under the relaxed budget regime it is non-binding. It becomes the operative constraint in the tight and scarce budget regimes.

There is no terminal sum and no deviation penalty in the objective. Completion incentives are embedded structurally: R_i^{ret} and R_{i,M_i}^{net} enter V_t at the final certification period; R_i^{term} (possibly negative) enters at termination, creating a natural deterrent against allowing project performance to deteriorate to the termination threshold (Ng et al., 1999). Symmetrically, the discounting of inflows by γ^{t-1} in (15) creates a completion incentive from above — earlier certification of milestone j yields $R_{i,j}^{\text{net}}$ at a higher discount weight, so a return-maximising agent will front-load allocations to projects where earlier certification is feasible, without any additional objective term. In competitive project-based markets where π_i is fixed at contract signing, this discounting effect is the formal expression of the capital velocity argument: the annualised return on committed capital increases with completion speed, and γ encodes the opportunity cost of capital that makes this concrete (Damodaran, 2012). The expectation $\mathbb{E}[\cdot]$ collapses to a deterministic evaluation at Level 1 and is retained at Levels 2 and 3.

3.3 Consolidated Parameter Reference

Table 3 lists all project parameters fixed at commitment, grouped by function. The planned S-curve and advance recovery parameters introduced in Sections 3.2.2–3.2.3 are included for completeness. The discount factor γ and initial balance B_1 are portfolio-level parameters shared across all projects.

Parameter	Symbol	Domain	Definition
Schedule			
Planned start & finish	$s_i,\ f_i$	$\mathbb{Z}^+$	Anchor of execution window and S-curve.
Planned duration	D_i^{plan}	$\mathbb{Z}^+$	$f_i - s_i + 1$; normalises progress against plan.
S-curve shape	$a_i,\ b_i$	$\mathbb{R}^+$	Beta CDF parameters governing planned expenditure profile (2).
Milestone threshold	$\theta_{i,j}$	$(0, 1]$	Cumulative progress fraction required to certify milestone j.
Earliest certification	$e_{i,j}$	$\mathbb{Z}^+$	Earliest period at which milestone j may be certified, regardless of progress achieved; enforces FIDIC Sub-Clause 14.3 payment interval constraints (Fédération Internationale des Ingénieurs-Conseils, 2017).
Schedule tolerance	Ω_i	$\mathbb{Z}^+$	Maximum admissible forecast finish overrun in periods before Condition 1 (9) activates.
To-Complete Performance Index	$\text{TCPI}_{i,t}$	$\mathbb{R}^+$	Forward-looking cost efficiency required on remaining work to finish within BAC_i; computed as $(\text{BAC}_i - \text{EV}_{i,t})/(\text{BAC}_i - \text{ACWP}_{i,t})$. Complements $\text{CPI}_{i,t}$ (historical) with a prospective signal; values above 1.10 are widely considered unrecoverable (Fleming and Koppelman, 2010).
Contract			
Contract price	CP_i	$\mathbb{R}^+$	Total client payment; $CP_i = \text{BAC}_i(1 + \pi_i)$; payment accumulation baseline only.
Budget at completion	BAC_i	$\mathbb{R}^+$	Total planned contractor expenditure; EVM cost basis and progress normalisation denominator.
Profit margin	π_i	$\mathbb{R}_0^+$	Contractual markup on BAC_i.
Advance payment ratio	α_i	$[0, 1]$	$A_i = \alpha_i \cdot CP_i$ received at $t = s_i$.
Recovery deduction rate	δ_i^{rec}	$(0, 1]$	Fraction of each interim payment deducted toward advance recovery once threshold ψ_i is exceeded; default 0.25 per FIDIC Sub-Clause 14.2 (Fédération Internationale des Ingénieurs-Conseils, 2017).
Recovery threshold	ψ_i	$[0, 1]$	Cumulative certified payments as fraction of CP_i at which advance recovery begins; default 0.10 per FIDIC Sub-Clause 14.2 (Fédération Internationale des Ingénieurs-Conseils, 2017).
Retention ratio	ρ_i	$[0, 1]$	Fraction withheld per milestone; released as $\rho_i \cdot CP_i$ at completion.
Milestone weight	$\phi_{i,j}$	$\mathbb{R}^+$	Fraction of CP_i at milestone j; $\sum_j \phi_{i,j} = 1$.
Termination			
Max cost overrun ratio	μ_i	$\mathbb{R}_0^+$	Contractual ceiling on projected total cost expressed as fractional overrun on BAC_i; activates Condition 2 (10).
Termination tolerance	τ_i^{tol}	$\mathbb{Z}^+$	Cure-period length in periods; models FIDIC Sub-Clause 15.1 (Fédération Internationale des Ingénieurs-Conseils, 2017; Ndekugri and Rycroft, 2008).
Portfolio			

3.4 Level 1: Deterministic Full-Foresight MILP

Setting $\eta_{i,t} = \bar{\eta}_{i,t}$ (known) renders S_t a deterministic function of $\mathbf{x}_t$, so $\mathcal{P}_t^{\text{active}}$, milestone certification times, and termination periods t_i^{term} are all computable ex ante. Under deterministic efficiency, the trajectories of $\text{SPI}_{i,t}$, $\text{CPI}_{i,t}$, $\text{EAC}_{i,t}$, and the counter $\tau_{i,t}^{\text{rem}}$ are fully determined before solving; both termination conditions (9)–(10) and the counter dynamics (11) are parameter-verifiable, so termination periods t_i^{term} and their settlements R_i^{term} enter the MILP as fixed parameters rather than optimisation constraints.

Introducing the certification binary $u_{i,j,t} \in \{0,1\}$ (= 1 iff milestone j of project i certified by period t) and activity indicator $\text{act}_{i,t} \in \{0,1\}$ (fixed parameter), the stochastic objective (15) collapses to:

$$\max \quad \sum_{t=1}^{H} \gamma^{t-1} \left(V_t^{\text{det}} - \sum_i \text{act}_{i,t}\, x_{i,t} \right) \tag{19a}$$

$$\text{s.t.} \quad B_{t+1} = B_t - \sum_i \text{act}_{i,t}\, x_{i,t} + V_t^{\text{det}}, \qquad \forall\, t \tag{19b}$$

$$\sum_i \text{act}_{i,t}\, x_{i,t} \le B_t, \qquad \forall\, t \tag{19c}$$

$$P_{i,t}^{\text{actual}} = P_{i,t-1}^{\text{actual}} + \bar{\eta}_{i,t}\, \tfrac{x_{i,t}}{\text{BAC}_i}, \quad P_{i,t}^{\text{actual}} \le 1, \qquad \forall\, i,t \tag{19d}$$

$$P_{i,t}^{\text{actual}} \ge \theta_{i,j}\, u_{i,j,t} - \mathbf{M}(1 - u_{i,j,t}), \qquad \forall\, i,j,t \tag{19e}$$

$$P_{i,t}^{\text{actual}} \le \theta_{i,j} + \mathbf{M}\, u_{i,j,t}, \quad u_{i,j,t} \ge u_{i,j,t-1}, \quad u_{i,j+1,t} \le u_{i,j,t}, \qquad \forall\, i,j,t \tag{19f}$$

$$u_{i,j,t} = 0 \quad \text{if } t < e_{i,j}, \qquad \forall\, i,j,t \tag{19g}$$

$$V_t^{\text{det}} = \sum_{i,j} R_{i,j}^{\text{net}} \left(u_{i,j,t} - u_{i,j,t-1}\right) + \sum_i R_i^{\text{ret}} \left(u_{i,M_i,t} - u_{i,M_i,t-1}\right) + \sum_i A_i \cdot \mathbf{1}\{t = s_i\} + \sum_{i \in \mathcal{P}^{\text{term}}} R_i^{\text{term}} \cdot \mathbf{1}\{t = t_i^{\text{term}}\}, \qquad \forall\, t \tag{19h}$$

$$x_{i,t} = 0 \quad \text{if } \text{act}_{i,t} = 0, \qquad \forall\, i,t \tag{19i}$$

$$x_{i,t} \ge 0, \quad u_{i,j,t} \in \{0,1\}, \quad u_{i,j,0} = P_{i,0}^{\text{actual}} = 0. \tag{19j}$$

The big-M constant is $\mathbf{M} = \max_{i,j} \theta_{i,j}$. Constraints (19e)–(19f) implement exact if-and-only-if certification logic with monotone, ordered binary activation. Constraint (19g) enforces the contractual earliest certification date: $u_{i,j,t}$ is fixed to zero for all periods prior to $e_{i,j}$, preventing the MILP from certifying a milestone before its contractual window opens regardless of how much budget has been allocated. This makes the optimal allocation timing uniquely determined under $\gamma < 1$: the agent will front-load spend to reach $\theta_{i,j}$ exactly at $e_{i,j}$, then hold until the window opens. The difference $(u_{i,j,t} - u_{i,j,t-1})$ in (19h) equals one in the sole period a milestone is first certified, linearising the payment trigger without additional binaries.

Upper bound. Let Z_{L1}^* denote the Level 1 optimum. For any causal policy π feasible under uncertainty, $\sup_\pi \mathbb{E}[Z^\pi] \le Z_{\text{L1}}^*$: every feasible causal allocation sequence is also feasible in the Level 1 program solved on the realised trajectory, so the full-foresight optimum dominates in expectation. The gap $Z_{\text{L1}}^* - Z^\pi$ quantifies the cost of uncertainty and anchors the empirical comparison in Section 5.

3.5 Level 2: Stochastic Extension and Intractability

3.5.1 Stochastic Primitive and Distributional Assumption

At Level 2, the deterministic efficiency scalar $\bar{\eta}_{i,t}$ is replaced by a random variable $\eta_{i,t}$ defined on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$. The efficiency realisations are assumed to satisfy:

Assumption 1 (Efficiency structure)**.** For each project $i \in \mathcal{P}$ and period $t \in \mathcal{T}$:

(i) $\eta_{i,t} > 0$ almost surely;

(ii) $\{\eta_{i,t}\}_{t \in \mathcal{T}}$ are independent across periods given project i;

(iii) $\eta_{i,t}$ are independent across projects $i \in \mathcal{P}$;

(iv) the marginal distribution of $\eta_{i,t}$ is project-specific and calibrated from empirical EVM datasets (Acebes et al., 2014; Batselier and Vanhoucke, 2017).

The realisation $\eta_{i,t}(\omega)$ is revealed only after allocation $x_{i,t}$ is committed. This single timing constraint defines the one decision point at which the portfolio manager faces irreducible uncertainty: the allocation decision under unknown productivity. Under Relaxation B, payments are received in the period of certification, so no second stochastic primitive is required.

3.5.2 Information Structure and Non-Anticipativity

Let $\mathcal{F}_t$ denote the σ-algebra generated by all efficiency realisations and the resulting state transitions, up to and including period t:

$$\mathcal{F}_t = \sigma\big(\eta_{i,\tau}(\omega) : i \in \mathcal{P},\ \tau \leq t\big), \qquad \mathcal{F}_0 = \{\emptyset, \Omega\}. \tag{20}$$

The history available to the portfolio manager at the start of period t, prior to committing $\mathbf{x}_t$, is:

$$\mathcal{H}_t = \big(S_1,\, \mathbf{x}_1,\, S_2,\, \mathbf{x}_2,\, \ldots,\, S_{t-1},\, \mathbf{x}_{t-1},\, S_t\big), \tag{21}$$

where S_t is the portfolio state defined in (14), which is $\mathcal{F}_{t-1}$-measurable since it reflects only realisations up to period $t-1$. The *non-anticipativity constraint* requires that the allocation at period t depends only on the available history:

$$\mathbf{x}_t \in \sigma(\mathcal{H}_t) \qquad \forall\, t \in \mathcal{T}. \tag{22}$$

This constraint distinguishes a causal policy from the full-foresight MILP of Level 1, which optimises over the entire realised trajectory $(\eta_{i,t})_{i,t}$ simultaneously.

3.5.3 The Multi-Stage Stochastic Program

Under Assumption 1 and constraint (22), the Level 1 MILP (19) becomes a multi-stage stochastic program (Birge and Louveaux, 2011). The decision at each stage is a *policy* $\pi_t : \mathcal{H}_t \to \mathbb{R}_{\geq 0}^{|\mathcal{P}_t^{\text{active}}|}$ rather than a fixed scalar, and the objective is an expectation over all

scenario realisations. The program is:

$$\max_{\{\pi_t\}} \quad \mathbb{E}\left[\sum_{t=1}^{H} \gamma^{t-1}\left(V_t - \sum_{i\in\mathcal{P}_t^{\text{active}}} x_{i,t}\right)\right] \tag{23a}$$

$$\text{s.t.} \quad x_{i,t} = \pi_t(\mathcal{H}_t)_i, \qquad \mathbf{x}_t \in \sigma(\mathcal{H}_t), \qquad \forall\, i,\, t \tag{23b}$$

$$P_{i,t}^{\text{actual}} = P_{i,t-1}^{\text{actual}} + \eta_{i,t} \cdot \frac{x_{i,t}}{\text{BAC}_i}, \qquad \forall\, i,\, t \tag{23c}$$

$$B_{t+1} = B_t - \sum_{i\in\mathcal{P}_t^{\text{active}}} x_{i,t} + V_t, \qquad \forall\, t \tag{23d}$$

$$\sum_{i\in\mathcal{P}_t^{\text{active}}} x_{i,t} \le B_t, \qquad \forall\, t \tag{23e}$$

$$x_{i,t} = 0 \quad \text{if } \text{status}_{i,t} \neq \texttt{active}, \qquad \forall\, i,\, t \tag{23f}$$

$$x_{i,t} \ge 0, \qquad \forall\, i,\, t \tag{23g}$$

where V_t is as defined in (4), now a random variable whose realisation depends on $(\eta_{i,\tau})_{\tau\le t}$ through the progress recursion (23c) and the milestone certification events it generates. The termination counter (11), and therefore $\mathcal{P}_t^{\text{active}}$ and $\text{status}_{i,t}$, are likewise random variables endogenous to the realised efficiency trajectory, in contrast to Level 1 where they were fixed parameters.

3.5.4 Scenario-Tree Representation and Exponential Growth

The canonical solution method for multi-stage stochastic programs is scenario-tree enumeration (Birge and Louveaux, 2011). The scenario tree for (23) has the following structure. Discretise $\eta_{i,t}$ to K values per project per period. At each period t, the tree branches by K^n (one branching factor of K per project, with n projects independent by Assumption 1(iii)). Over a horizon of H periods the tree has:

$$|\mathcal{S}| = K^{nH} \tag{24}$$

leaf nodes, each corresponding to a complete realisation of $(\eta_{i,t})_{i\in\mathcal{P},\, t\in\mathcal{T}}$.

For a representative infrastructure portfolio with $n = 10$ projects, a horizon H derived endogenously from the committed portfolio (typically $H \in [18, 36]$ months in practice), and $K = 5$ discrete efficiency scenarios per project per period, the leaf-node count is:

$$|\mathcal{S}| = 5^{10\times 24} = 5^{240} \approx 10^{167}, \tag{25}$$

a number that exceeds the estimated number of atoms in the observable universe by more than 80 orders of magnitude. This is not a consequence of the binary milestone variables inherited from the MILP structure: replacing them with continuous relaxations leaves the exponent nH unchanged, since the exponential growth is intrinsic to the multi-stage stochastic formulation and independent of the within-stage problem structure.

Endogenous horizon compounding. A further compounding factor distinguishes this problem from standard multi-stage stochastic programs with fixed horizons. The planning horizon $H = \max_i f_i - \min_i s_i + 1$ is itself a function of the committed portfolio; it is not an external configuration parameter. Under stochastic efficiency, the effective execution window of each project is uncertain — a project with $\text{SPI}_{i,t} < 1$ throughout will overrun

f_i, extending the realised horizon beyond H. In the worst case, every project overruns simultaneously, and the actual number of periods over which decisions must be made grows with the portfolio. Each additional period multiplies the scenario tree by K^n, so horizon uncertainty compounds (24) multiplicatively.

Termination as structural bifurcation. At Level 1, termination periods t_i^{term} were fixed parameters precomputed from the known efficiency trajectory. At Level 2, the period at which the cure-period counter (11) reaches zero depends on the realised sequence of $\text{SPI}_{i,t}$ and $\text{CPI}_{i,t}$ values, which are in turn determined by $(\eta_{i,\tau})_{\tau \leq t}$. Each termination event removes project i from $\mathcal{P}_t^{\text{active}}$ and delivers R_i^{term} into V_t — a structural change in the active project set that bifurcates the scenario tree at every period in which a project's counter is positive. Across n projects each with counter depth τ_i^{tol}, the number of structurally distinct subtrees grows combinatorially in the termination event space, adding a second source of exponential scaling beyond (24).

3.5.5 Intractability and the Motivation for Level 3

The exponential scaling established in (24) constitutes a fundamental computational barrier. It is not a function of solver technology or implementation efficiency; it is intrinsic to the multi-stage stochastic formulation of any sequential decision problem with independent per-stage uncertainty across n components over H stages (Birge and Louveaux, 2011). Approximation methods for multi-stage stochastic programs — Sample Average Approximation (Kleywegt et al., 2002), Progressive Hedging (Rockafellar and Wets, 1991), and Stochastic Dual Dynamic Programming (Pereira and Pinto, 1991) — reduce the computational burden but do not escape the exponential scaling in n and H at portfolio scales of practical interest. Sample Average Approximation, for instance, requires a sample size that grows with the variance of the objective and the number of stages to achieve a given optimality gap (Kleywegt et al., 2002); for $n = 10$ and $H = 24$ this remains intractable without structural decomposition that the shared cash balance (23d) prevents.

Proposition 1 (Level 1 as upper bound)**.** *Let Z_{L1}^* denote the optimal value of the Level 1 MILP (19) evaluated on a realised efficiency trajectory $(\bar{\eta}_{i,t})_{i,t}$. Let Z^π denote the objective value achieved by any causal policy π satisfying (22) on the same trajectory. Then:*

$$\sup_{\pi \in \Pi^{\text{causal}}} \mathbb{E}[Z^\pi] \leq Z_{\text{L1}}^*, \tag{26}$$

where Π^{causal} denotes the set of all policies satisfying (22).

Proof. Fix any realised trajectory $\omega \in \Omega$. Any allocation sequence $\{x_{i,t}\}$ produced by a causal policy π under ω is a feasible solution to the Level 1 MILP solved on ω (since feasibility constraints (19b)– (19i) are identical to (23d)–(23f) for fixed $\bar{\eta}_{i,t} = \eta_{i,t}(\omega)$; the single uncertainty source means no separate payment delay trajectory need be fixed). Therefore $Z^\pi(\omega) \leq Z_{\text{L1}}^*(\omega)$ for every ω. Taking expectations on both sides yields (26). □

The gap $Z_{\text{L1}}^* - \mathbb{E}[Z^\pi]$ quantifies the cost of uncertainty: the value destroyed by operating under incomplete information relative to a hypothetical full-foresight oracle. This gap is the primary empirical quantity reported in Section 5. The intractability of computing the exact solution to (23) motivates the Level 3 reformulation: rather than enumerating the scenario tree, the problem is recast as a sequential interaction between a policy and an environment, solved through sampled trajectory experience rather than exhaustive enumeration.

3.6 Level 3: Constrained MDP and RL Agent

The intractability established in Section 3.5 rules out exact solution of (23) at realistic portfolio scales. The Level 3 reformulation replaces exhaustive scenario-tree enumeration with a sequential interaction framework: the portfolio manager is modelled as a *policy* that observes the current state, commits an allocation, and receives a reward signal, with the environment transitioning stochastically to the next state. Optimisation proceeds through repeated sampled interaction rather than explicit enumeration, making the computational cost per training step independent of $|\mathcal{S}| = K^{nH}$.

3.6.1 Constrained Markov Decision Process Formulation

Definition 1 (Portfolio Governance CMDP)**.** The portfolio budget allocation problem at Level 3 is formalised as a finite-horizon Constrained Markov Decision Process (Altman, 1999):

$$\mathcal{M} = \big(\mathcal{S},\ \mathcal{A},\ \mathcal{T},\ R,\ \mathcal{C},\ \gamma,\ H\big), \tag{27}$$

with components defined as follows.

State space $\mathcal{S}$. The state at period t is the portfolio state S_t defined in (14):

$$S_t \in \mathcal{S} \subseteq [0,1]^n \times \mathbb{R}^n_{\geq 0} \times \mathbb{R}^{4n} \times \{0,1,2\}^n \times \mathbb{Z}^n_{\geq 0} \times \mathbb{R}_{\geq 0}, \tag{28}$$

collecting, for each project i: actual progress $P^{\text{actual}}_{i,t} \in [0,1]$, cumulative actual cost $\text{ACWP}_{i,t} \in \mathbb{R}_{\geq 0}$, performance indices $\text{SPI}_{i,t}$, $\text{CPI}_{i,t}$, $\text{TCPI}_{i,t}$, $\text{EAC}_{i,t}$, schedule deviation $\Delta f_{i,t}$, project status $\text{status}_{i,t} \in \{0,1,2\}$ encoding $\{\texttt{active}, \texttt{completed}, \texttt{terminated}\}$, and cure-period counter $\tau^{\text{rem}}_{i,t} \in \mathbb{Z}_{\geq 0}$; together with the portfolio cash balance $B_t \in \mathbb{R}_{\geq 0}$. The state dimension is $9n + 1$ per period. All components are observable. One primitive is unobservable at decision time: $\eta_{i,t}$, realised only after allocation $x_{i,t}$ is committed. It is not part of the state vector; it is an environment-internal draw that the policy must act without knowing in advance.

Action space $\mathcal{A}$ and feasible set. The raw action space is:

$$\mathcal{A} = \mathbb{R}^n_{\geq 0}, \tag{29}$$

with the state-dependent feasible action set:

$$\mathcal{A}(S_t) = \left\{ \mathbf{x}_t \in \mathbb{R}^n_{\geq 0} \;\middle|\; \sum_{i \in \mathcal{P}^{\text{active}}_t} x_{i,t} \leq B_t, \quad x_{i,t} = 0 \text{ if } \text{status}_{i,t} \neq \texttt{active} \right\}, \tag{30}$$

encoding constraints (17) and (18) directly. The hard liquidity constraint $B_t \geq 0$ is enforced by construction through the projection operator defined below.

Projection operator. To guarantee $\mathbf{x}_t \in \mathcal{A}(S_t)$ for any raw network output $\tilde{\mathbf{x}}_t \in \mathbb{R}^n$, the environment applies a two-step projection before executing the transition. First, allocations to inactive projects are zeroed: $\tilde{x}_{i,t} \leftarrow 0$ for all i with $\text{status}_{i,t} \neq \texttt{active}$. Second, the remaining vector is projected onto the budget simplex:

$$\mathbf{x}_t = B_t \cdot \frac{\max(\tilde{\mathbf{x}}^{\text{active}}_t,\ \mathbf{0})}{\left\|\max(\tilde{\mathbf{x}}^{\text{active}}_t,\ \mathbf{0})\right\|_1 \vee 1}, \tag{31}$$

where $\tilde{\mathbf{x}}_t^{\text{active}}$ denotes the sub-vector restricted to active projects, and $\vee$ denotes the maximum operator ensuring the denominator is at least one when the agent outputs all-zero or all-negative values. This projection is differentiable almost everywhere and preserves gradient flow through the policy network during training (Schulman et al., 2017).

Transition kernel $\mathcal{T}$. The stochastic transition $\mathcal{T}(S_{t+1} \mid S_t, \mathbf{x}_t)$ is defined by the composition of environment mechanics applied sequentially within each period:

(i) Draw $\eta_{i,t} \sim F_i$ independently for each active project i, where F_i is the project-specific efficiency distribution calibrated from Acebes et al. (2014); Batselier and Vanhoucke (2017).

(ii) Update progress and actual cost via (1).

(iii) Compute EVM signals via (5)–(7).

(iv) Evaluate termination conditions (9)–(10) and update counter via (11).

(v) Certify milestones for which $P_{i,t}^{\text{actual}} \geq \theta_{i,j}$ *and* $t \geq e_{i,j}$. Each newly certified milestone (i,j) contributes $R_{i,j}^{\text{net}}$ to V_t immediately in the certification period (Relaxation B). Compute V_t via (3)–(4).

(vi) Update portfolio cash balance via (13).

(vii) Execute termination for any project with $\tau_{i,t}^{\text{rem}} = 0$ or $t \geq f_i + \Omega_i$, compute R_i^{term} via (12).

The transition kernel is Markovian: S_{t+1} depends on $(S_t, \mathbf{x}_t, \boldsymbol{\eta}_t)$ only, not on $\mathcal{H}_{t-1}$. This is the key structural property that makes the CMDP formulation exact rather than approximate: the full portfolio state S_t defined in (14) is a sufficient statistic for the history $\mathcal{H}_t$ defined in (21).

Reward signal R. The per-period reward is the discounted net cash flow term from (15):

$$r_t(S_t, \mathbf{x}_t, S_{t+1}) = \gamma^{t-1} \left(V_t - \sum_{i \in \mathcal{P}_t^{\text{active}}} x_{i,t} \right), \tag{32}$$

where V_t is realised after the transition. The reward contains no deviation penalty term; completion incentives are embedded structurally through R_i^{ret} and R_{i,M_i}^{net} at project completion and through R_i^{term} (possibly negative) at termination, as established in Section 3.2.7.

Constraint set $\mathcal{C}$. The hard liquidity constraint $B_t \geq 0$ is the sole constraint promoted to the CMDP constraint set:

$$\mathcal{C} = \{ \mathbf{x}_t \in \mathcal{A} \mid B_{t+1} \geq 0 \}. \tag{33}$$

This constraint is enforced by the projection (31) rather than through a Lagrangian penalty, ensuring hard satisfaction at every step rather than satisfaction only in expectation (Altman, 1999).

CMDP objective. The policy optimisation problem is:

$$\pi^* = \text{argmax}_{\pi \in \Pi} \; J(\pi) \;=\; \text{argmax}_{\pi \in \Pi} \; \mathbb{E}_\pi \left[\sum_{t=1}^{H} r_t(S_t, \mathbf{x}_t, S_{t+1}) \right], \tag{34}$$

subject to $\mathbf{x}_t \in \mathcal{C}$ for all t, where Π denotes the set of all measurable policies satisfying (22) and $\mathbb{E}_\pi$ denotes expectation under the distribution over trajectories induced by π and $\mathcal{T}$. The objective $J(\pi)$ in (34) is identical to the intractable stochastic program (23a); the CMDP reformulation does not change the objective but changes the solution method from scenario-tree enumeration to policy-gradient optimisation.

Remark 1 (Relationship to Level 1 upper bound). By Proposition 1, $J(\pi^*) \leq Z^*_{\text{L1}}$ for any causal policy π^*. The CMDP does not close this gap; it provides the tractable path to finding the best achievable $J(\pi)$ within Π^{causal}. The empirical gap $Z^*_{\text{L1}} - J(\hat{\pi})$, where $\hat{\pi}$ is the trained PPO agent, quantifies the combined cost of uncertainty and approximation error in Section 5.

3.6.2 Gymnasium-Compatible Environment

The CMDP $\mathcal{M}$ is implemented as a configurable Gymnasium-compatible environment (Towers et al., 2024). The environment exposes the standard `reset()` / `step()` interface, with `observation_space` $= \mathcal{S}$ and `action_space` $= \mathcal{A}$. All project and portfolio parameters (Table 3) are passed as a configuration dictionary at initialisation; the planning horizon H is derived automatically as $\max_i f_i - \min_i s_i + 1$ from the committed project set, making the environment fully parameterisable across contract types and portfolio sizes without code modification.

The observation vector returned at each step is a flattened concatenation of S_t as defined in (28), normalised to $[0, 1]$ per component to facilitate stable policy gradient estimation. The environment enforces the projection (31) on the raw action vector before applying the transition kernel, guaranteeing $\mathbf{x}_t \in \mathcal{A}(S_t)$ regardless of the policy output. Episodes terminate when $t = H$ or when all projects have reached `completed` or `terminated` status, whichever occurs first.

The stochastic efficiency $\eta_{i,t}$ is drawn each period from a project-specific Beta distribution calibrated from empirical EVM datasets (Acebes et al., 2014; Batselier and Vanhoucke, 2017); the Beta family accommodates the right-skewed, bounded support structure of empirical CPI and SPI distributions documented in Lipke (2003) and Batselier and Vanhoucke (2017). Under Relaxation B, milestone payments are received in the certification period; no payment delay distribution is sampled.

3.6.3 Outlook: Policy Training

The CMDP $\mathcal{M}$ is designed to be solved by any policy-gradient algorithm operating on the standard `reset()` / `step()` interface. Proximal Policy Optimisation (Schulman et al., 2017) is the natural candidate: its clipped surrogate objective provides stable updates in continuous action spaces, its on-policy sample collection is compatible with the non-stationary episode structure of $\mathcal{M}$ in which $\mathcal{P}_t^{\text{active}}$ changes endogenously, and it has demonstrated consistent performance in operations management domains with similar state-action structure (Giannoccaro and Pontrandolfo, 2002; Peng and Huang, 2021). The observation vector (14) maps directly to a fully connected actor-critic network; the projection (31) enforces feasibility without a penalty term in the surrogate objective. Domain randomisation (Tobin et al., 2017) over the configuration space can be used to

train a policy that generalises across contract types and portfolio sizes without retraining. Training, hyperparameter selection, and evaluation against the Level 1 upper bound are addressed in subsequent work.

3.6.4 Greedy Baseline Policies

Three deterministic greedy policies are defined over the observation vector (14) to establish a performance baseline and quantify the optimality gap. Each policy reads directly from the catchup and reach-plan fields embedded in the observation.

Reach-Plan (π^{rp}). At each period, allocate exactly enough to bring every active project to its S-curve planned progress:

$$x_{i,t}^{\text{rp}} = x_{i,t}^{\text{rp}} = \max\bigl(0,\ PV_i(\tau_i + 1) - EV_{i,t}\bigr) \cdot \text{BAC}_i. \tag{35}$$

This policy closes the progress gap entirely each period but does not account for milestone timing, contract profitability, or budget scarcity.

Catchup-T (π^{cu}). Allocate only enough to stay within the contractual progress delay cap $\delta_i^{\max}$ this period:

$$x_{i,t}^{\text{cu}} = \max\bigl(0,\ PV_i(\tau_i + 1) - \delta_i^{\max} - EV_{i,t}\bigr) \cdot \text{BAC}_i. \tag{36}$$

This is the minimum spend that avoids a breach flag this period. Beyond $\tau_i \geq D_i^{\text{plan}}$ the policy switches to $x_{i,t}^{\text{rp}}$ to prevent hard deadline termination.

Catchup-Next-T ($\pi^{\text{cu+1}}$). The most risk-averse policy: pre-emptively allocate enough this period so that the projected delay *next* period also remains within $\delta_i^{\max}$:

$$x_{i,t}^{\text{cu+1}} = \max\bigl(0,\ PV_i(\tau_i + 2) - \delta_i^{\max} - EV_{i,t}\bigr) \cdot \text{BAC}_i. \tag{37}$$

This one-period lookahead avoids both current and next-period breaches at the cost of higher early expenditure and correspondingly lower discounted returns.

All three policies are normalised by the available budget B_t before application, using the same projection (31) as the learned agent. Their empirical performance across four benchmark configurations is reported in Section 5.

4 Evaluation

5 Benchmark Evaluation and Optimality Gap

This section reports the baseline benchmark study conducted to establish a performance floor, motivate the need for learned allocation policies, and quantify the improvement target available to a future reinforcement learning agent. Three deterministic greedy heuristics and a manually computed full-information optimum are evaluated across four benchmark portfolio configurations of increasing structural complexity. All episodes are executed through the Gymnasium-compatible environment described in Section 3, with the SQLite audit ledger providing a complete per-period record of allocations, cash-flow events, and breach flags for every episode.

5.1 Shared Parameters and Duration Convention

All four configurations share the contract and environment parameters listed in Table 4. These are fixed across every episode and every solver to ensure that observed performance differences are attributable solely to configuration structure and solver policy.

Parameter	Symbol	Value
Retention rate	ν_i	0.05
Progress delay cap	$\delta_i^{\max}$	0.15
Finish delay cap	$\Delta_i^{\max}$	2 periods
Termination tolerance	θ_i	2 breach-periods
Annual interest rate	r	0.12
Per-period deficit rate	$r/12$	0.01
Budget tightness	κ	3.0
Profit margin	π_i	0.20
Discount factor	γ	0.95
S-curve shape	(a_i, b_i)	(1.5, 3.0)

Table 4: Parameters shared across all benchmark configurations. Retention is deducted from each milestone gross payment after advance recovery has been applied; the accumulated retention is released in full at the final milestone.

Duration convention. A project with `planned_duration` $= D$ executes over project-timesteps $\tau_i \in \{0, 1, \ldots, D\}$, giving $D + 1$ active periods. The final planned timestep is $\tau_i = D$, at which the last planned milestone is due. With finish delay cap $\Delta_i^{\max} = 2$, the hard deadline falls at episode period $f_i + 2$, and the total efficiency sequence supplied to the environment has length $D + 1 + 2 = D + 3$ entries covering all possible active periods.

Deduction order. Within each milestone payment, advance recovery is applied first to the gross payment; the retention deduction is then applied to the same gross payment independently. At the final milestone, any advance not yet recovered is deducted in full, and all accumulated retention is released. This ordering — advance recovery before retention, both applied to gross — is validated consistently across all configurations.

5.2 Benchmark Configurations

Table 5 summarises the four benchmark configurations. Each is designed to isolate a specific structural mechanic of the environment, so that the failure mode of each greedy heuristic can be attributed to a concrete feature of the problem.

TC-10SF-R — Isolated atom. One project, one non-advance milestone, no advance payment. With $D_i = 0$, the project has a single planned action at $\tau_i = 0$, and the hard deadline is at episode period 2. The efficiency sequence is [0.91, 1.17, 0.83]. Contract price $P = 100 \times 1.20 = 120$. The single milestone carries weight $w = 1.0$, so gross payment $= 120$. Retention of $0.05 \times 120 = 6$ is withheld and immediately released at the same (final) milestone, giving net payment $= 120$. There is no advance payment

Config	N	BAC_i	D_i	s_i	M_i	ϕ_i	ψ_i^{trig}	**Primary mechanic**
TC-10SF-R	1	100	0	0	1	0.00	—	Isolated atom: single action, single milestone
TC-1LMAIF-R	1	100	3	2	3	0.20	0.80	Delayed start, progress-triggered advance, high reco
TC-10MAIF-R	1	100	3	0	3	0.20	0.00	Diminishing efficiency, termination risk
TC-3XMAIF-R	3	100/200/300	3	0/6/7	3	0.20	0.00	Multi-project, gap then overlap, interacting cash flo

Table 5: Benchmark configuration summary. N: number of projects; D_i: planned duration (so contractual active periods = D_i+1); s_i: planned start (episode period); M_i: number of non-advance milestones; ϕ_i: advance rate; ψ_i^{trig}: progress threshold at which the advance is disbursed.

or recovery. This configuration provides the cleanest possible baseline: it tests the fundamental progress-to-cash-flow loop with no confounding payment mechanics.

TC-1LMAIF-R — Delayed start and progress-triggered advance. One project, planned start $s_i = 2$, planned duration $D_i = 3$ (four active periods: $\tau_i \in \{0, 1, 2, 3\}$), three non-advance milestones, and an advance rate $\phi_i = 0.20$ with advance trigger $\psi_i^{\text{trig}} = 0.80$. The trigger means the advance payment of $0.20 \times 120 = 24$ is disbursed only once cumulative progress $EV_{i,t} \geq 0.80$ is reached — *not* at project start. Consequently the advance is absent from the first interim milestone and the associated deductions begin only after the trigger fires.

Advance recovery rate is $\rho_i = 0.50$, the highest of all configurations. With milestone gross payments of 18, 42, and 60 (weights 0.15, 0.35, 0.50), the recovery trace is: $j = 1$: $\min(0.50 \times 18,\ 24) = 9.0$ recovered, advance remaining = 15.0; $j = 2$: $\min(0.50 \times 42,\ 15) = 15.0$ recovered, advance fully cleared; $j = 3$: no further recovery. Retention deductions (0.9, 2.1, 3.0) accumulate across the three milestones and are released in full (6.0) at $j = 3$. Net payments after deductions are 8.1, 24.9, and 63.0. The efficiency sequence alternates: [0.91, 1.17, 0.83, 0.91, 1.17, 0.83]. This configuration validates: (i) delayed project start and its interaction with the episode clock; (ii) progress-triggered advance disbursement; (iii) high advance recovery rate and the liquidity trough it creates after the advance is received; and (iv) the ordering invariant that advance recovery is deducted before retention on every gross payment.

TC-10MAIF-R — Diminishing efficiency and termination risk. One project, planned start $s_i = 0$, planned duration $D_i = 3$, three milestones, advance rate $\phi_i = 0.20$ (trigger at progress = 0, i.e. disbursed at project start), advance recovery rate $\rho_i = 0.20$. The efficiency sequence is monotonically *decreasing*: [0.95, 0.91, 0.55, 0.40, 0.30, 0.20]. Early periods are productive; late periods are near-inoperable.

The advance of 24 is disbursed at $t = 0$. Milestone payment trace (advance recovery before retention): $j = 1$: $\min(0.20 \times 18,\ 24) = 3.6$ recovered, advance remaining = 20.4, retention = 0.9, net = 13.5; $j = 2$: $\min(0.20 \times 42,\ 20.4) = 8.4$ recovered, advance remaining = 12.0, retention = 2.1, net = 31.5; $j = 3$: $\min(0.20 \times 60,\ 12.0) = 12.0$ recovered (advance fully cleared at the final milestone via the standard rate), retention = 3.0, released = 6.0, net = 51.0. This configuration provides a balanced three-milestone structure while stress-testing the termination mechanic: the sharp efficiency drop beyond $\tau_i = 1$ makes it progressively more expensive to maintain progress above $\delta_i^{\max} = 0.15$, and any policy that does not front-load work risks triggering the cure-period counter in the later periods.

TC-3XMAIF-R — Multi-project with gap and overlap. Three simultaneous projects with heterogeneous scale (BAC = [100, 200, 300], prices [120, 240, 360]) and staggered starts ($s_i = [0, 6, 7]$). Each project has planned duration $D_i = 3$ (four active periods), three milestones, and advance rate $\phi_i = 0.20$ triggered at project start. Advances at episode periods 0, 6, and 7 are 24, 48, and 72 respectively, creating three distinct outflow spikes that the policy must navigate against a shared budget.

The episode timeline has two structurally distinct phases. *Gap phase* ($t = 4, 5$): Project 0 completes at $t = 3$; Projects 1 and 2 have not yet started. No allocations are required and the budget accumulates interest on any surplus. *Overlap phase* ($t = 7, 8, 9$): both Projects 1 and 2 are simultaneously active, creating competing allocation demands against the shared reserve B_t. The concurrent advance outflows at $t = 6$ (-48) and $t = 7$ (-72) produce the largest single-period deficit in the benchmark suite.

Per-project efficiency sequences are: Project 0: [0.91, 1.17, 0.83, 0.91, 1.17, 0.83]; Project 1: [1.09, 0.84, 1.18, 1.09, 0.84, 1.18]; Project 2: [0.97, 1.13, 0.86, 0.97, 1.13, 0.86]. None is monotonically degrading; the challenge here is purely one of coordination and cash-flow timing, not individual project difficulty.

5.3 Solvers

Four methods are evaluated, as specified in Section 3.6.4.

Reach-Plan (π^{rp}) allocates $x_{i,t}^{\text{rp}}$ to bring every active project exactly to its S-curve planned progress each period, regardless of budget scarcity or milestone proximity.

Catchup-T (π^{cu}) allocates the minimum spend $x_{i,t}^{\text{cu}}$ required to keep the progress delay within $\delta_i^{\max} = 0.15$ this period, preserving budget for later periods and other projects.

Catchup-Next-T ($\pi^{\text{cu+1}}$) pre-emptively allocates $x_{i,t}^{\text{cu+1}}$ so that the projected delay next period also remains within $\delta_i^{\max}$, at the cost of higher early expenditure.

Full-information optimum (π^*) is computed manually with complete prior knowledge of all efficiency realisations ($\eta_{i,t}$). It serves as an upper-bound reference for each configuration. Implementing a tractable algorithmic solver — whether Linear Program or Dynamic Program — is left as future work; the manual computation suffices to establish the performance ceiling across the four configurations studied here.

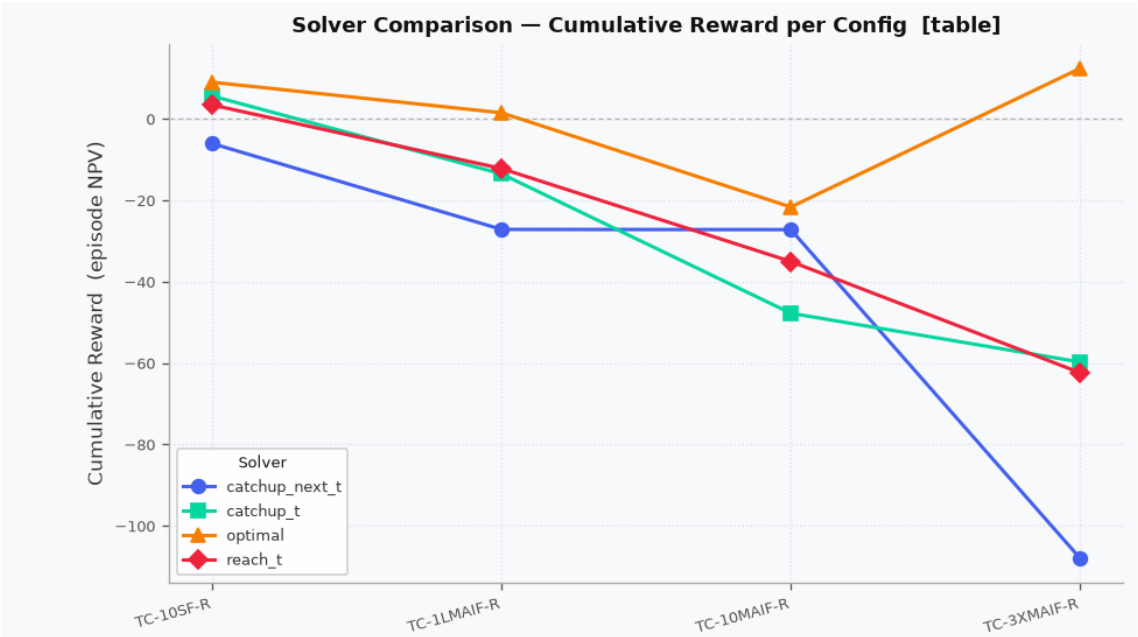
All methods are evaluated under seed = 0, with the fixed custom ordinal efficiency sequences described above, ensuring exact reproducibility.

5.4 Results

Figure 7 reports cumulative episode NPV for each method across the four configurations, ordered left to right by increasing portfolio complexity. Table 6 presents the numerical values and the best-heuristic gap $G(\pi^*) - \max_\pi G(\pi)$ for each configuration.

Config	$G(\pi^*)$	$G(\pi^{\text{rp}})$	$G(\pi^{\text{cu}})$	$G(\pi^{\text{cu+1}})$	**Best-heuristic gap**
TC-10SF-R	9.00	5.60	3.40	−6.07	3.40
TC-1LMAIF-R	1.49	−13.21	−13.21	−27.17	14.70
TC-10MAIF-R	−21.66	−35.06	−47.77	−27.22	5.56
TC-3XMAIF-R	12.40	−59.75	−62.38	−107.96	72.15

Table 6: Episode NPV G for all methods across four configurations. Best-heuristic gap $= G(\pi^*) - \max_{\pi \in \{\pi^{\text{rp}}, \pi^{\text{cu}}, \pi^{\text{cu+1}}\}} G(\pi)$.



	10SF-R	1LMAIF-R	10MAIF-R	3XMAIF-R
catchup_next_t	-6.07	-27.17	-27.22	-107.96
catchup_t	5.60	-13.53	-47.77	-59.75
optimal	9.00	1.49	-21.66	12.40
reach_t	3.40	-12.21	-35.06	-62.38

Figure 7: Cumulative episode NPV ($G = \sum_{t=0}^{T} \gamma^t \Delta B_t$) for three deterministic greedy heuristics and the full-information optimum across four benchmark configurations ordered by increasing portfolio complexity. The widening gap confirms that rule-based allocation deteriorates systematically as structural complexity grows.

TC-10SF-R. On the isolated single-period configuration all three heuristics produce positive NPV and the gap is modest. Reach-Plan achieves 5.60 against the optimal 9.00 (gap 3.40); Catchup-T follows at 3.40. Catchup-Next-T finishes at −6.07: with only one active period and one milestone, its one-period lookahead over-spends relative to what the single milestone requires, incurring interest costs without a compensating return. The result confirms that the environment correctly prices conservative over-allocation as a loss.

TC-1LMAIF-R. The progress-triggered advance and high recovery rate ($\rho_i = 0.50$) create a liquidity trough immediately after the advance is disbursed. All three heuristics produce negative NPV while the optimum remains marginally positive at 1.49; the best-heuristic gap is 14.70. The optimal allocation exploits the known efficiency sequence to time the milestone certifications so that the advance trigger fires at the most advantageous moment and the large recovery deductions at $j = 1$ (−9.0) and $j = 2$ (−15.0) are offset by subsequent cash inflows. The greedy policies, which allocate based on current-period signals only, cannot anticipate the trigger timing and over-allocate during the low-return period before the advance fires. The fact that both π^{rp} and π^{cu} achieve identical NPV (−13.21) indicates that the binding constraint here is not the delay cap but the payment timing: both rules happen to satisfy the tolerance in the same periods, but neither correctly sequences allocation to capture the advance inflow early.

TC-10MAIF-R. The monotonically diminishing efficiency sequence makes the allocation timing problem acute. Even the full-information optimum achieves only −21.66: the late-stage efficiency collapse ($\eta = 0.30$, 0.20 in the delay window) makes full project recovery uneconomical. All heuristics perform materially worse. Catchup-Next-T (−27.22) outperforms both Reach-Plan (−35.06) and Catchup-T (−47.77): its risk-averse lookahead front-loads work into the high-efficiency early periods, reducing expenditure in the near-inoperable late periods where each unit of allocation buys very little progress. Catchup-T performs worst because its minimum-spend rule continues allocating into low-efficiency periods to maintain the tolerance floor, accumulating cost without meaningful progress gain. The best-heuristic gap of 5.56 is the smallest in the benchmark suite — but even here the optimum is materially better and the gap arises from information the greedy rules cannot access.

TC-3XMAIF-R. The multi-project configuration produces the starkest result. The full-information optimum achieves +12.40: coordinating across three projects with known efficiency sequences, the oracle times advance disbursements, milestone certifications, and the gap-phase budget accumulation to convert the complex portfolio into a profitable episode. Every greedy heuristic produces a loss: Reach-Plan at −59.75, Catchup-T at −62.38, and Catchup-Next-T at −107.96. The best-heuristic gap is 72.15, a 21× increase over TC-10SF-R. The sign reversal — profitable optimum, loss-making heuristics — is the critical observation. The concurrent advance outflows at $t = 6$ and $t = 7$ (−48 and −72) create a large shared deficit that each single-project greedy rule handles independently, without accounting for how its own allocation depletes the reserve available to the concurrent project. Catchup-Next-T collapses to −107.96 because its aggressive pre-emptive spend across all three simultaneously-active projects exhausts the budget during the overlap phase, leaving insufficient reserve to certify the high-value late milestones of Projects 1 and 2.

5.5 Discussion

Three findings emerge from the benchmark study.

The gap is monotonically increasing with structural complexity. The best-heuristic gap grows from 3.40 (TC-10SF-R) to 14.70 (TC-1LMAIF-R) to 72.15 (TC-3XMAIF-R), a 21× increase from simplest to most complex. TC-10MAIF-R is the only exception to the monotone trend: its gap of 5.56 is smaller than TC-1LMAIF-R's 14.70 because the challenge there is project-level efficiency timing, not multi-project coordination. The gap in TC-3XMAIF-R is not marginal degradation; it reflects a structural failure of greedy allocation under shared-budget coordination requirements that no single-project rule can resolve.

No single heuristic dominates. Each heuristic is best on at least one configuration. Reach-Plan is best on TC-10SF-R; Catchup-Next-T is best on TC-10MAIF-R; Catchup-T is best on TC-3XMAIF-R. This confirms that the optimal policy is adaptive: the correct allocation depends jointly on the efficiency trajectory, the milestone payment structure, the advance trigger timing, and the number and phasing of concurrent projects — a combination no fixed rule can capture.

A well-defined improvement target exists for learned agents. The full-information optimum $G(\pi^*)$ provides a concrete upper bound. On TC-3XMAIF-R the swing between the best achievable (12.40) and the best greedy (-59.75) is 72.15 NPV units, corresponding to the difference between a profitable and a deeply loss-making portfolio outcome. This gap is the value of intelligent sequential allocation — and it is precisely the target that a policy trained on the CMDP environment of Section 3 is positioned to close. Training such a policy and evaluating it against a formal Linear Programming bound constitute the natural next step, addressed in subsequent work.

6 Conclusion

6.1 Summary of Contributions

This paper establishes the foundational methodology for Reinforcement Learning-based project portfolio budgeting under cashflow uncertainty. Our work addresses a critical gap in operations research: the inability of classical optimization methods to handle sequential decision-making under partial observability in portfolio management.

We make four primary contributions:

(1) Literature-Calibrated Synthetic Data Generation: We developed a comprehensive framework for generating realistic project portfolios with empirically-grounded uncertainty distributions. All parameters are calibrated from published studies (318 megaprojects, multiple empirical analyses), ensuring scientific validity while addressing confidentiality constraints.

(2) Rolling Horizon Control Architecture: We proposed a rolling horizon framework that maintains fixed 12-period episodes while handling variable-duration portfolios through sequential re-planning, preventing state-space explosion and enabling tractable RL training.

(3) Formal Problem Formulation: We provided the first rigorous mathematical formulation of portfolio budgeting as a POMDP with explicit state space, action space, transition dynamics, and reward function, all grounded in empirical project management literature.

(4) Complete Modeling Framework: We documented a modular portfolio modeling system covering composition, cashflows, uncertainty, and revenue plans, with

all assumptions explicitly stated and future extensions identified.

6.2 Implications for Research and Practice

For Operations Research: This work demonstrates that RL can address problems where classical OR methods fail due to unknown transition probabilities and computational intractability. The rolling horizon architecture provides a template for applying RL to other long-horizon sequential decision problems.

For Project Management: The literature-calibrated synthetic data methodology enables reproducible research in domains with confidentiality constraints. Our framework can be extended to construction, infrastructure, R&D, and other project-intensive industries.

For Machine Learning: This application showcases RL's potential in operations research domains, complementing recent successes in robotics, game playing, and control. The integration of domain knowledge (S-curves, milestone payments, SPI/CPI dynamics) with data-driven learning exemplifies hybrid AI approaches.

6.3 Limitations and Scope Boundaries

We deliberately adopted the *Simplest Model Principle* to establish a foundational framework. Key limitations include:

- **Aggregated cashflow modeling:** No resource-level detail (labor, materials, equipment)
- **Fixed portfolio composition:** 60/40 domestic/international mix assumed optimal
- **Milestone-based payments only:** Alternative structures (time-based, cost-plus) excluded
- **No project interdependencies:** Shared resources and technology spillovers not modeled
- **Synthetic data only:** Real-world validation pending company partnerships

Each limitation represents a natural extension for future work, documented with clear research directions.

6.4 Future Work

Immediate Next Steps (Current Research):

1. **RL Agent Implementation:** Select and train policy network (PPO, SAC, or Transformer-based)
2. **Baseline Comparison:** Implement MIP, stochastic programming, and greedy heuristics
3. **Experimental Evaluation:** Conduct comprehensive experiments across scenarios
4. **Sensitivity Analysis:** Analyze robustness to parameter variations

Medium-Term Extensions (6–12 Months):

1. **Fine-Tuning Protocol:** Develop methodology for adapting pretrained models to company-specific data
2. **Interpretability Analysis:** Apply attention mechanisms or SHAP values to explain allocation decisions
3. **Multi-Objective Optimization:** Extend to Pareto-optimal policies balancing NPV, risk, and strategic alignment
4. **Real-World Validation:** Partner with EPC contractors for pilot deployment

Long-Term Research Directions (1–3 Years):

1. **Resource-Level Modeling:** Disaggregate cashflows into labor, materials, equipment with detailed scheduling
2. **Dynamic Portfolio Selection:** Integrate project selection with execution optimization
3. **Multi-Agent RL:** Model interactions between contractor, client, and subcontractors
4. **Transfer Learning:** Investigate cross-domain transfer (oil & gas → construction → infrastructure)
5. **Hierarchical RL:** Develop two-level policies (strategic portfolio allocation + tactical project control)

6.5 Broader Impact

This research has potential impact beyond project portfolio management:

Methodological Contributions:

- Template for applying RL to OR problems with unknown transition dynamics
- Framework for literature-calibrated synthetic data generation
- Rolling horizon architecture for long-horizon sequential decision problems

Practical Applications:

- Real-time decision support for portfolio managers
- Scenario analysis for strategic planning
- Risk assessment through Monte Carlo simulation
- Training tool for project management education

Open Science:

- Complete codebase and datasets released as open source
- Reproducible research enabling community validation and extension
- Benchmark problem for RL algorithm development

6.6 Concluding Remarks

Project portfolio budgeting under uncertainty represents a critical challenge for capital-intensive industries. Traditional optimization methods fail due to computational intractability and unknown state transition probabilities. This paper establishes the foundational methodology for addressing this challenge through Reinforcement Learning with rolling horizon control.

By developing a literature-calibrated synthetic data generation framework, we overcome the confidentiality barrier that has prevented progress in this domain. Our formal problem formulation and complete modeling framework provide a rigorous foundation for future research.

The path forward is clear: implement and evaluate RL agents, compare against classical baselines, and develop fine-tuning protocols for company-specific deployment. We invite the research community to build upon this foundation, extending the methodology to related domains and advancing the state of the art in RL-based operations research.

The complete modeling framework, instance generator, and documentation will be released as open-source software upon publication, enabling reproducible research and practical deployment in this emerging field.

References

Acebes, F., Pajares, J., Galán, J. M., and López-Paredes, A. (2014). A new approach for project control under uncertainty: Going back to the basics. *International Journal of Project Management*, 32(3):423–434.

Ajibade, O. M. (2025). Ai enhanced project management: Leveraging predictive analytics and intelligent automation. *Int. J. Res. Publ. Rev*, 6:6828–6844.

Akbari, M. and Shirazi, M. A. (2026). A self-adjusting decision support system for strategic project portfolio prioritization using integrated ahp and reinforcement learning. *Applied Operations and Analytics*, 2(1):1–19.

Altman, E. (1999). *Constrained Markov Decision Processes*. Routledge, New York, 1st edition.

Archer, N. P. and Ghasemzadeh, F. (1999). An integrated framework for project portfolio selection. *International Journal of Project Management*, 17(4):207–216.

Arditi, D. and Chotibhongs, R. (2005). Issues in subcontracting practice. *Journal of Construction Engineering and Management*, 131(8):866–876.

Asghari, V., Biglari, A. J., and Hsu, S.-C. (2023). Multiagent Reinforcement Learning for Project-Level Intervention Planning under Multiple Uncertainties. *Journal of Management in Engineering*, 39(2):04022075.

Bahroun, Z., Tanash, M., As'ad, R., and Alnajar, M. (2023). Artificial Intelligence Applications in Project Scheduling: A Systematic Review, Bibliometric Analysis, and Prospects for Future Research. *Management Systems in Production Engineering*, 31:144–161.

Barraza, G. A., Back, W. E., and Mata, F. (2004). Probabilistic monitoring of project performance using S-curves. *Journal of Construction Engineering and Management*, 130(1):142–148.

Batselier, J. and Vanhoucke, M. (2017). Improving project forecast accuracy by integrating earned value management with exponential smoothing and reference class forecasting. *International Journal of Project Management*, 35(1):28–43.

Beseiso, M. and Kumar, G. (2021). A fuzzy computational approach for selecting interdependent projects using prioritized criteria. *Journal of Intelligent & Fuzzy Systems*, 40(6):11341–11354.

Birge, J. R. and Louveaux, F. (2011). *Introduction to Stochastic Programming*. Springer, New York, 2nd edition.

Cerqueira, V., Torgo, L., and Mozetič, I. (2020). Evaluating time series forecasting models: An empirical study on performance estimation methods. *Machine Learning*, 109:1997–2028.

Christensen, D. S. (1993). The estimate at completion problem: A review of three studies. *Project Management Journal*, 24(1):37–42.

Christensen, D. S. and Heise, S. R. (1993). Cost performance index stability. *National Contract Management Journal*, 25:7–15.

Cooper, R. G., Edgett, S. J., and Kleinschmidt, E. J. (2001). *Portfolio Management for New Products*. Perseus Publishing, Cambridge, MA, 2nd edition.

Damodaran, A. (2012). *Investment Valuation: Tools and Techniques for Determining the Value of Any Asset*. Wiley Finance. John Wiley & Sons, Hoboken, NJ, 3rd edition.

Darvish, A. and Sepehri, M. (2025). Artificial intelligence-driven project portfolio optimization under deep uncertainty using adaptive reinforcement learning. *Applied Sciences*, 15(23):12713.

Dixit, A. K. and Pindyck, R. S. (1994). *Investment under Uncertainty*. Princeton University Press, Princeton, NJ.

Doloi, H. (2012). Cost overruns and failure in project management: Understanding the roles of key stakeholders in construction projects. *Journal of Construction Engineering and Management*, 138(3):267–279.

Ebnerasoul, M., Ghannadpour, S. F., and Haeri, A. (2023). A collective efficacy-based approach for bi-objective sustainable project portfolio selection using interdependency network model between projects. *Environment, Development and Sustainability*, 25:13981–14001.

Eilat, H., Golany, B., and Shtub, A. (2006). Constructing and evaluating balanced portfolios of R&D projects with interactions: A DEA based methodology. *European Journal of Operational Research*, 172(3):1018–1039.

El-Sabek, L. M. and McCabe, B. Y. (2018). Coordination Challenges of Production Planning in the Construction of International Mega-Projects in The Middle East. *International Journal of Construction Education and Research*, 14(2):118–140.

Ennis, C. and Burr, A. (2021). Termination under FIDIC: Conditions, consequences and disputes. *International Construction Law Review*, 38(2):112–134.

Fédération Internationale des Ingénieurs-Conseils (2017). *Conditions of Contract for EPC/Turnkey Projects (Silver Book)*. FIDIC, Geneva, 2nd edition.

Fleming, Q. W. and Koppelman, J. M. (2010). *Earned Value Project Management*. Project Management Institute, Newtown Square, PA, 4th edition.

Flyvbjerg, B. (2017). *The Oxford Handbook of Megaproject Management*. Oxford University Press, Oxford.

Flyvbjerg, B., Holm, M. K. S., and Buhl, S. L. (2003). How common and how large are cost overruns in transport infrastructure projects? *Transport Reviews*, 23(1):71–88.

Ghapanchi, A. H., Tavana, M., Khakbaz, M. H., and Low, G. (2012). A methodology for selecting portfolios of projects with interactions and under uncertainty. *International Journal of Project Management*, 30(7):791–803.

Gholami, H., Azizi, A., Sabzehparvar, M., and Jafari, D. (2025). A fuzzy multi-agent model of project portfolio scheduling and selection taking into account environmental resilience. *Results in Control and Optimization*, 19:100544.

Giannoccaro, I. and Pontrandolfo, P. (2002). Inventory management in supply chains: A reinforcement learning approach. *International Journal of Production Economics*, 78(2):153–161.

Huchzermeier, A. and Loch, C. H. (2001). Project management under risk: Using the real options approach to evaluate flexibility in R&D. *Management Science*, 47(1):85–101.

Huddar, V. and Joshi, N. (2025). Ai-driven project portfolio optimization for sustainable development: A decision support framework. In *2025 First International Conference on Advances in Computer Science, Electrical, Electronics, and Communication Technologies (CE2CT)*, pages 1035–1038. IEEE.

Jiang, C., Li, X., Lin, J.-R., Liu, M., and Ma, Z. (2023). Adaptive control of resource flow to optimize construction work and cash flow via online deep reinforcement learning. *Automation in Construction*, 150:104817.

Jiang, Y., Olmo, J., and Atwi, M. (2024). Deep reinforcement learning for portfolio selection. *Global Finance Journal*, 62:101016.

Kaka, A. P. and Price, A. D. F. (1993). Modelling standard cost commitment curves for contractors' cash flow forecasting. *Construction Management and Economics*, 11(4):271–283.

Karimi, S., Saeed, M., and Pishvaee, M. (2021). A scenario-based mathematical approach to a robust project portfolio selection problem under fuzzy uncertainty. *Journal of Intelligent & Fuzzy Systems*, 42:1–14.

Kavadias, S. and Loch, C. H. (2003). *Project Selection under Uncertainty: Dynamically Allocating Resources to Maximize Value*. Kluwer Academic Publishers, Boston, MA.

Kehinde, O. (2025). Leveraging data-driven decision-making for enhanced risk management and resource allocation in projects. *International Journal of Computer Applications Technology and Research*, 14(2):1–17.

Kenley, R. and Wilson, O. (2009). *A Construction Management Approach to S-Curves and Cash Flow Management*. Spon Press, London.

Killen, C. P., Hunt, R. A., and Kleinschmidt, E. J. (2008). Project portfolio management for product innovation. *International Journal of Quality & Reliability Management*, 25(1):24–38.

Kim, E.-B. and Ballard, G. (2002). Case study: Using earned value method in field management of construction project. *Journal of Construction Engineering and Management*, 128(5):436–445.

Kleywegt, A. J., Shapiro, A., and Homem-de Mello, T. (2002). The sample average approximation method for stochastic discrete optimization. *SIAM Journal on Optimization*, 12(2):479–502.

Kosztyán, Z. T., Katona, A. I., Kuppens, K., Kisgyörgy-Pál, M., Nachbagauer, A., and Csizmadia, T. (2022). Exploring the structures and design effects of eu-funded r&d&i project portfolios. *Technological Forecasting and Social Change*, 180:121687.

Kwak, Y. H. and Ingall, L. (2007). Exploring Monte Carlo simulation applications for project management. *Risk Management*, 9(1):44–57.

Lipke, W. (2003). Schedule is different. *The Measurable News*, pages 10–15.

Liu, Z., Liang, Y., and Zheng, H. (2023). Multi-project scheduling based on deep reinforcement learning. *Automation in Construction*, 148:104721.

Lova, A., Maroto, C., and Tormos, P. (2000). A multicriteria heuristic method to improve resource allocation in multiproject scheduling. *European Journal of Operational Research*, 127(2):408–424.

Markowitz, H. (1952). Portfolio selection. *The Journal of Finance*, 7(1):77–91.

Martinsuo, M. and Lehtonen, P. (2007). Role of single-project management in achieving portfolio management efficiency. *International Journal of Project Management*, 25(1):56–65.

Mazhari, N. and Tajik, S. (2024). Evaluating the impact of artificial intelligence and machine learning in project portfolio management. Master's thesis, Stockholm University.

Merrow, E. W. (2011). *Industrial Megaprojects: Concepts, Strategies, and Practices for Success*. Wiley, Hoboken, NJ.

Nabot, A., Alnaimat, F., Jebreen, I., Al-Qerem, A., and Ali, A. M. (2025). Ai for software project management: Predicting delivery times & resource allocation. In *2025 12th International Conference on Information Technology (ICIT)*, pages 277–284. IEEE.

Ndekugri, I. and Rycroft, M. (2008). The FIDIC engineering contracts: A commentary. *Construction Law Journal*, 24(3):181–198. Analysis of Sub-Clause 15.1 cure period obligations and wrongful termination jurisprudence.

Ng, A. Y., Harada, D., and Russell, S. (1999). Policy invariance under reward transformations: Theory and application to reward shaping. In *Proceedings of the 16th International Conference on Machine Learning (ICML)*, pages 278–287.

Noroozi, A., Mokhtari, H., Bagherpour, M., and Sadjadi, S. J. (2016). A probabilistic portfolio budget allocation problem with CPI index under risk. *Journal of Industrial and Production Engineering*, 33(4):236–246.

Page, M. J., McKenzie, J. E., Bossuyt, P. M., Boutron, I., Hoffmann, T. C., Mulrow, C. D., Shamseer, L., Tetzlaff, J. M., Akl, E. A., Brennan, S. E., Chou, R., Glanville, J., Grimshaw, J. M., Hróbjartsson, A., Lalu, M. M., Li, T., Loder, E. W., Mayo-Wilson, E., McDonald, S., McGuinness, L. A., Stewart, L. A., Thomas, J., Tricco, A. C., Welch, V. A., Whiting, P., and Moher, D. (2021). The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. *BMJ*, 372:n71.

Park, M. and Peña-Mora, F. (2004). Reliability buffering for construction projects. *Journal of Construction Engineering and Management*, 130(5):626–637.

Peng, W. and Huang, M. (2021). A deep reinforcement learning approach for construction project scheduling. *Applied Sciences*, 11(3):1044.

Pereira, M. V. F. and Pinto, L. M. V. G. (1991). Multi-stage stochastic optimization applied to energy planning. *Mathematical Programming*, 52(1):359–375.

Pourahmadi, K., Nouri, S., and Yaghoubi, S. (2015). A scenario based project portfolio selection. *Management Science Letters*, pages 883–888.

Project Management Institute (2017). *The Standard for Portfolio Management*. Project Management Institute, Newtown Square, PA, 4th edition.

Project Management Institute (2021). *A Guide to the Project Management Body of Knowledge (PMBOK Guide)*. Project Management Institute, Newtown Square, PA, 7th edition.

Ranjbar, M., Nasiri, M. M., and Torabi, S. A. (2022). Multi-mode project portfolio selection and scheduling in a build-operate-transfer environment. *Expert Systems with Applications*, 189:116134.

Rockafellar, R. T. and Wets, R. J.-B. (1991). Scenarios and policy aggregation in optimization under uncertainty. *Mathematics of Operations Research*, 16(1):119–147.

Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.

Staw, B. M. and Ross, J. (1987). Knowing when to pull the plug. *Harvard Business Review*, 65(2):68–74.

Sutton, R. S. and Barto, A. G. (2018). *Reinforcement Learning: An Introduction*. MIT Press, Cambridge, MA, 2nd edition.

Szwarcfiter, C., Herer, Y. T., and Shtub, A. (2023). Balancing Project Schedule, Cost, and Value under Uncertainty: A Reinforcement Learning Approach. *Algorithms*, 16(8):395.

Tassel, P., Gebser, M., and Schekotihin, K. (2023). An end-to-end reinforcement learning approach to job-shop scheduling problems. *Proceedings of the International Conference on Automated Planning and Scheduling*, 33:338–347.

Tobin, J., Fong, R., Ray, A., Schneider, J., Zaremba, W., and Abbeel, P. (2017). Domain randomization for transferring deep neural networks from simulation to the real world. In *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 23–30.

Touran, A. (2006). A mathematical model for establishing budget levels for a portfolio of transit projects. *Proc., ICCCBE XI*.

Touran, A. (2010). Probabilistic Approach for Budgeting in Portfolio of Projects. *Journal of Construction Engineering and Management*, 136(3):361–370.

Towers, M., Terry, J. K., Kwiatkowski, A., et al. (2024). Gymnasium: A standard interface for reinforcement learning environments. *arXiv preprint arXiv:2407.17032*.

Turner, J. R. (2014). *The Handbook of Project-Based Management*. McGraw-Hill, New York, 4th edition.

Van Peteghem, V. and Vanhoucke, M. (2010). A genetic algorithm for the preemptive and non-preemptive multi-mode resource-constrained project scheduling problem. *European Journal of Operational Research*, 201(2):409–418.

Yamakawa, E. K., Cauchick-Miguel, P. A., Sousa-Zomer, T. T., and Killen, C. P. (2019). Project portfolio management: A landscape of the literature. *International Journal of Business Excellence*, 18(4):450–487.

Zhang, X., Qiu, Y., and Pardalos, P. M. (2020). Project portfolio selection and scheduling under uncertainty using multi-objective evolutionary algorithm. *Annals of Operations Research*, 285:3–24.

A Appendix A : PRISMA 2020 Flow Diagram

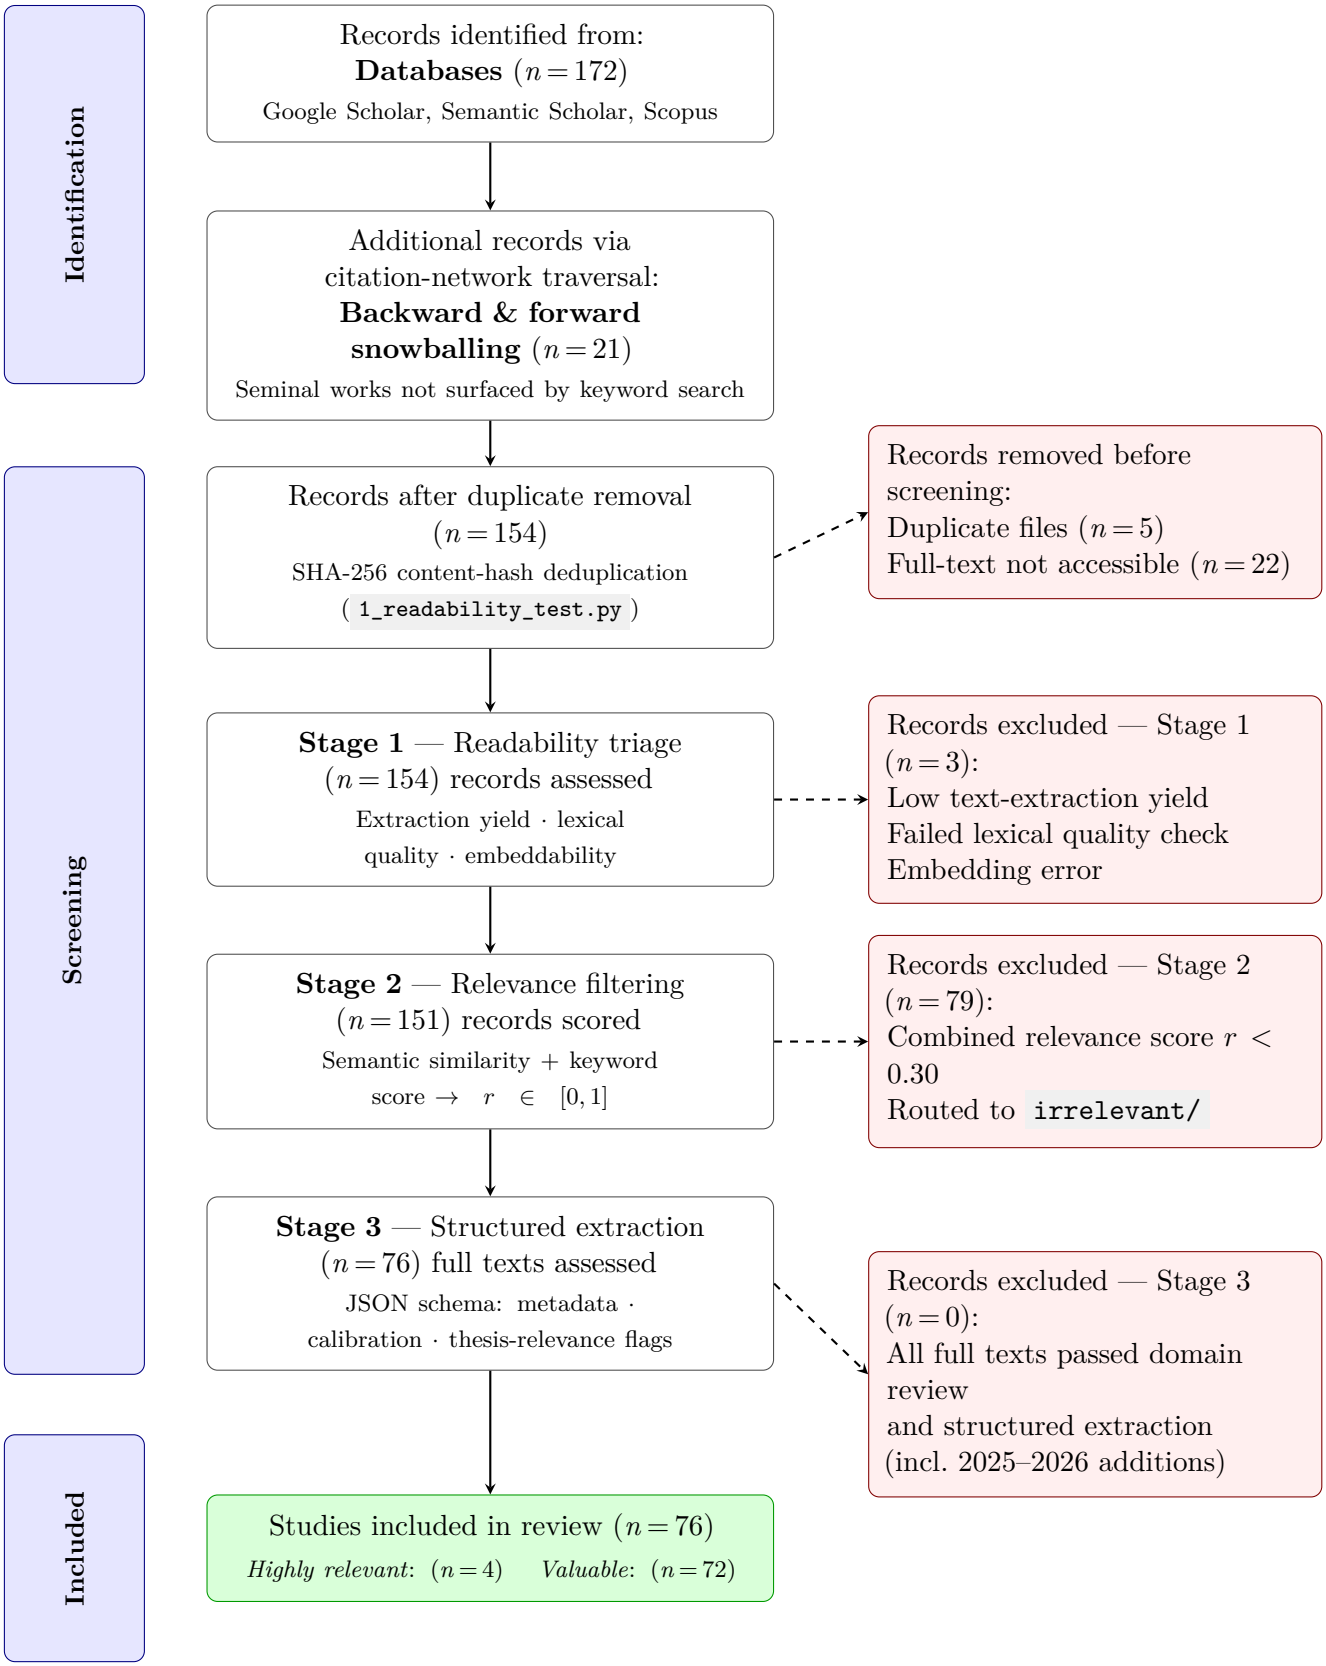


Figure 8: PRISMA 2020 flow diagram of the automated literature search and screening process. Stage labels correspond to scripts in `docs/lit_archive/`.

B Appendix B : Gymnasium Compatibility Certificate

`tests/gym_test.py` exercises `PortfolioBudgetingEnv` across ten test sections and 51 individual checks, covering import and instantiation, space declarations, the `reset()` and `step()` API contracts, determinism under fixed and varying seeds, random-policy and degenerate (zero-action, greedy) rollouts, multi-episode stability, and the official `gymnasium.utils.check_env` checker. The suite is run against benchmark configuration `TC-3XMAIF-R` (3 projects, mixed start times, fixed efficiency sequences). Table 7 summarises the ten sections; all 51 checks pass.

Table 7: Gymnasium compatibility test summary — `PortfolioBudgetingEnv`, config `TC-3XMAIF-R`, generated 2026-10-20.

§	Section	Passed
1	Import & Instantiation	8 / 8
2	Space Declarations	7 / 7
3	`reset()` Contract	13 / 13
4	`step()` Contract	10 / 10
5	Determinism	2 / 2
6	Random Policy Rollout	6 / 6
7	Zero-Action Rollout	1 / 1
8	Greedy Rollout	1 / 1
9	Multi-Episode Stability	2 / 2
10	`gymnasium.utils.env_checker`	1 / 1
Total		**51 / 51**

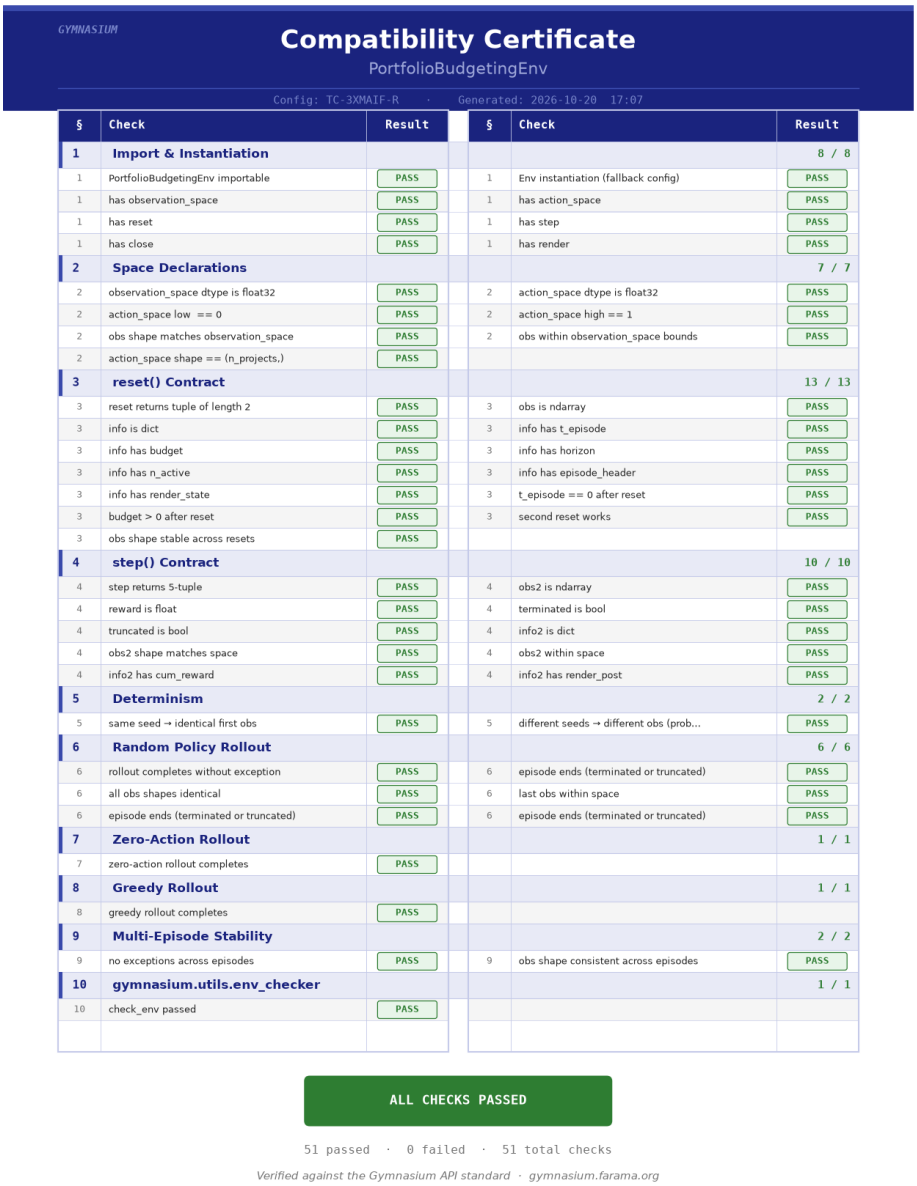


Figure 9: Compatibility certificate generated by `tests/gym_test.py`. All 51 checks pass; 0 failures.

C Reopsitory Architecture and Concept Map

C.1 Repository Architecture Overview

The complete implementation accompanying this paper — the environment, the baseline solvers, the verification suite, and the literature-review pipeline described in Section 2 — is organised as a single version-controlled repository. The repository separates *documentation* (this manuscript, the thesis, and the literature corpus) from *implementation* (the environment, ledger, baselines, and tests), as summarised below.

```
ppm_env/
|-- notation.yaml        parameter dictionary (academic <-> code)
|-- docs/
|   |-- article/          LaTeX source of this manuscript
|   |-- lit_archive/      literature screening pipeline and corpus
|   |-- mscn/             supplementary thesis documentation
|   |-- pp_presentation/  thesis defense slides (university deliverable)
|   `-- thesis_farsi/     Persian thesis manuscript (university deliverable)
`-- src/
    |-- db/               ledger schema, configs, run records
    |-- env/              core CMDP environment implementation
    `-- tests/            verification suite, baselines, gap analysis
```

The `mscn/`, `pp_presentation/`, and `thesis_farsi/` directories satisfy institutional degree-completion requirements and fall outside the reproducibility scope of this paper; they are listed here only for completeness. The `notation.yaml` file at the repository root provides an explicit, machine-readable correspondence between every academic symbol introduced in Section 3 and its counterpart variable name in the source code, and is intended as the primary cross-reference for readers wishing to connect the formal model to the implementation.

The literature-review pipeline described in Section 2.1.1 resides in `docs/lit_archive/` and retains, alongside the processing scripts, the full PDF archive (`lit_archive_passed/`, `lit_archive_failed/`) and the structured JSON extraction records that drove the synthesis table and gap-analysis figures reported in Section 2.

C.2 Core Contribution Modules

The modules implementing the formal model of Section 3 and the evaluation of Section 4 are contained entirely within `src/`. The environment dynamics — state construction, EVM/S-curve computation, the FIDIC cash-flow cycle, and the termination gate — are implemented in `src/env/`, with the environment's public Gymnasium-compatible interface and its internal support routines maintained as separate modules within that directory. During training or evaluation, the environment class is instantiated per episode through the standard Gymnasium `reset()` / `step()` interface; no additional driver component is required.

The SQLite audit ledger introduced in Section 3 is schema-defined and persisted under `src/db/`, which also holds the benchmark portfolio configurations used in Section 4 and the run-record databases generated by each executed episode.

The verification and evaluation layer resides in `src/tests/`, which bundles four functionally distinct components: the Gymnasium compliance suite, the baseline and optimal solvers, the benchmark configuration generator together with its archived reference episodes, and the post-hoc analysis scripts used to produce the figures and tables of Section 4.

```
src/tests/
|-- gym_test.py                 Gymnasium API compliance suite
|-- gym_certificate.png         compliance certificate (Appendix B)
|-- baseline_solvers.py         greedy heuristics and optimum solver
|-- reference_episodes_gen/     benchmark configuration generator
|-- reference_episodes_record/  archived configs and solutions
`-- analysis/                   gap-analysis and figure scripts
```

The `reference_episodes_record/` directory retains the configuration set and recorded solutions used to validate the environment and to produce the four benchmark configurations evaluated in Section 4; it is intended as a working template for practitioners who wish to instantiate the environment against their own company data.

C.3 Concept-to-Implementation Traceability Matrix

Table 8 maps every concept and promised contribution named in the main text to its corresponding location in the repository, and is provided to support independent verification of the claims made throughout this paper.

Table 8: Traceability matrix mapping concepts and promised contributions to their implementation in the repository.

Concept / Contribution	Ref.	Repository Location	Role
Academic–code notation dictionary	§3	`notation.yaml`	Maps formal symbols to code identifiers
EVM/S-curve state representation	§3	`src/env/`	State vector construction
FIDIC cash-flow transition dynamics	§3	`src/env/`	Advance, retention, interest, and milestone logic
Two-condition termination gate	§3	`src/env/`	Episode termination evaluation
SQLite audit ledger	§3	`src/db/`	Schema definition and per-run event logging
Benchmark portfolio configurations	§4	`src/db/`	Stored configuration instances for evaluation runs
Gymnasium compliance verification	App. B	`src/tests/`	API conformance tests and certificate generation
Greedy heuristic baselines (Reach-Plan, Catchup-T, Catchup-Next-T)	§4	`src/tests/`	Deterministic baseline policies

continued on next page

Table 8 – continued

Concept / Contribution	Ref.	Repository Location	Role
Manually computed full-information optimum	§4	`src/tests/`	Upper-bound benchmark solver
Benchmark configuration generation and reference episodes	§4	`src/tests/`	Configuration matrix, episode archive, and company-data templates
Gap-analysis figures and dashboards	§4	`src/tests/`	Post-hoc analysis and figure generation
Literature screening, Stages 1–3	§2.1.1	`docs/lit_archive/`	Readability triage, relevance filtering, structured extraction
Synthesis dashboard and literature map	§2.1.1	`docs/lit_archive/`	Gap matrix, domain charts, synthesis table generation
Manuscript source	—	`docs/article/`	LaTeX source, figures, and tables for this paper

C.4 Reproducibility Notes

The mapping in Table 8 reflects a specific snapshot of the repository at the time of writing; subsequent development may reorganise individual modules without affecting the directory-level structure described in this appendix. Access to the repository, including the full commit history, the benchmark configurations, and the Gymnasium compliance certificate reproduced in Appendix B, is available from the corresponding author upon request.